

Virtual memory

Acknowledgements:

CO and design- Chapter 5 - Henessey and Patterson

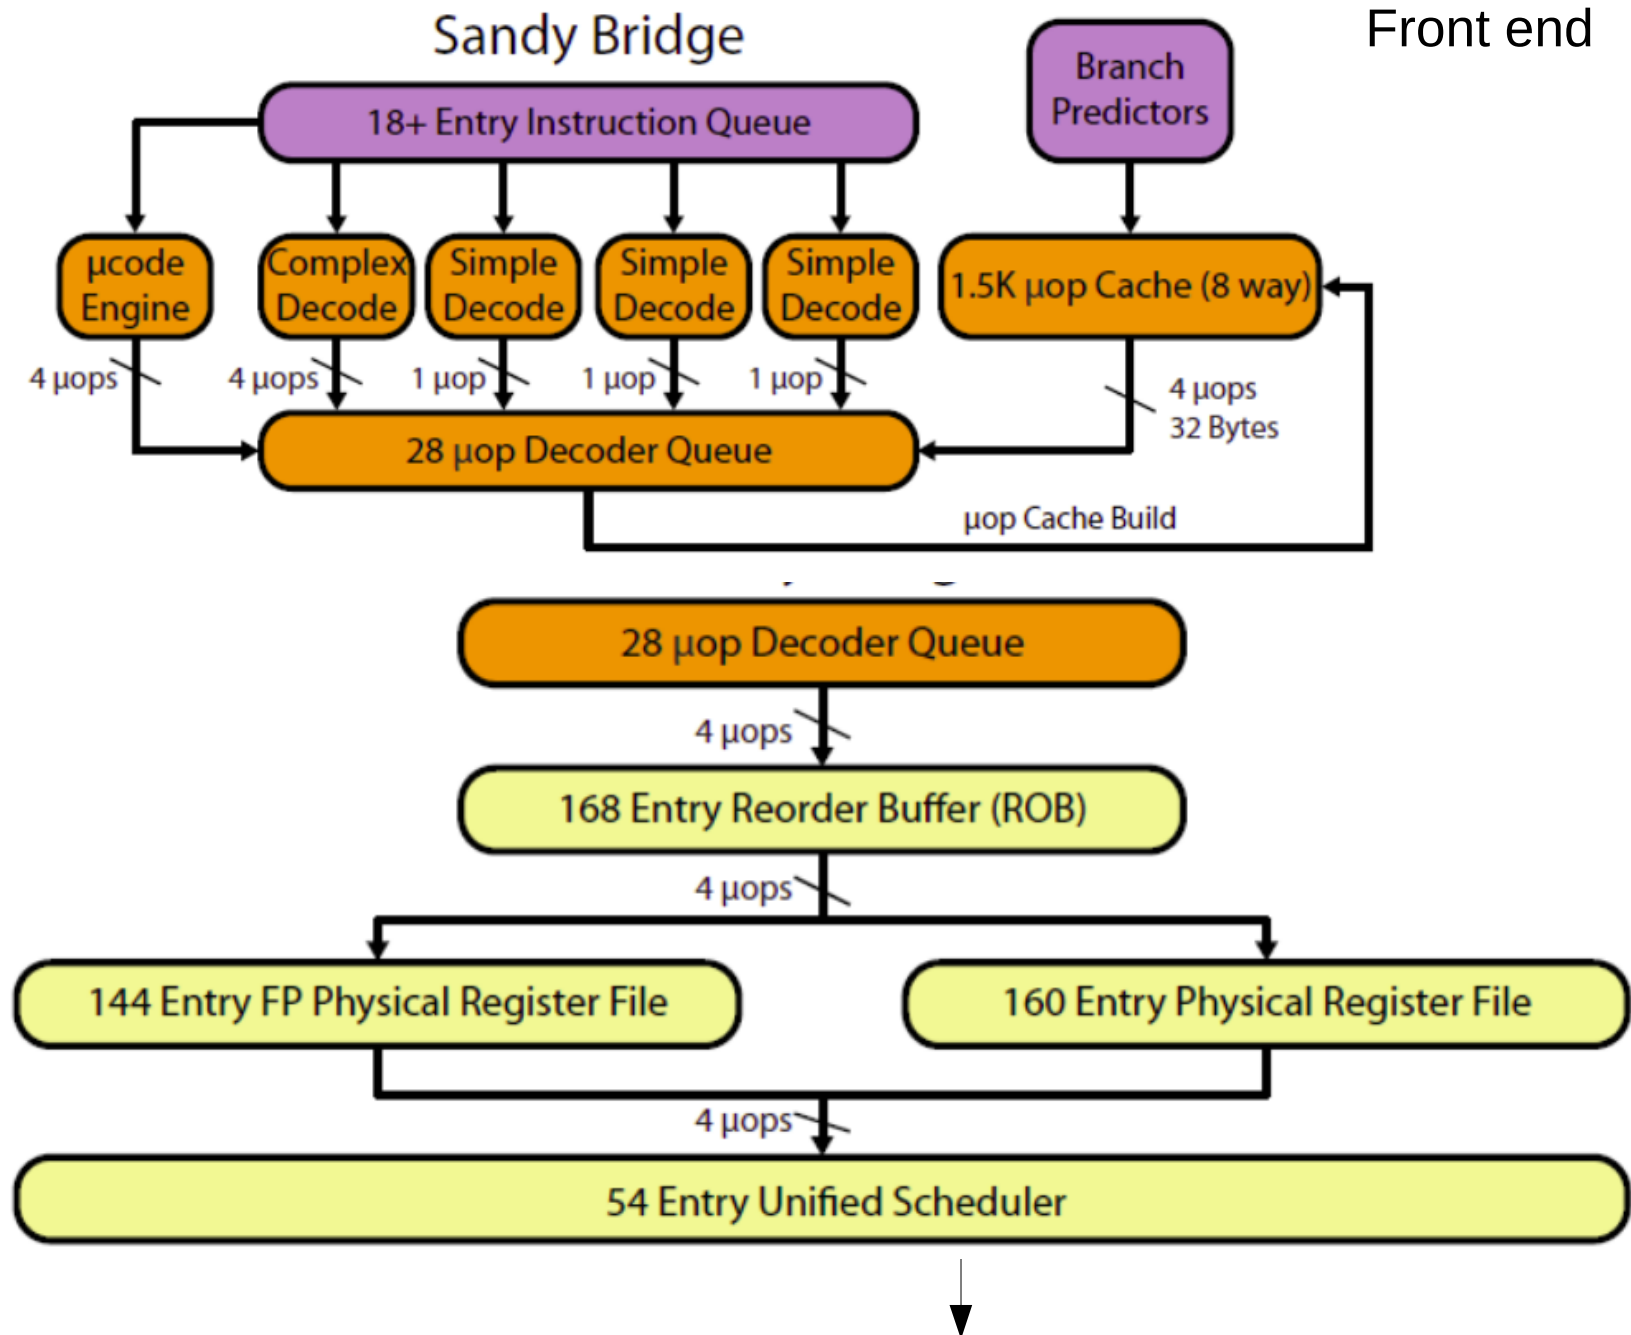
CA- Quant approach- H & P

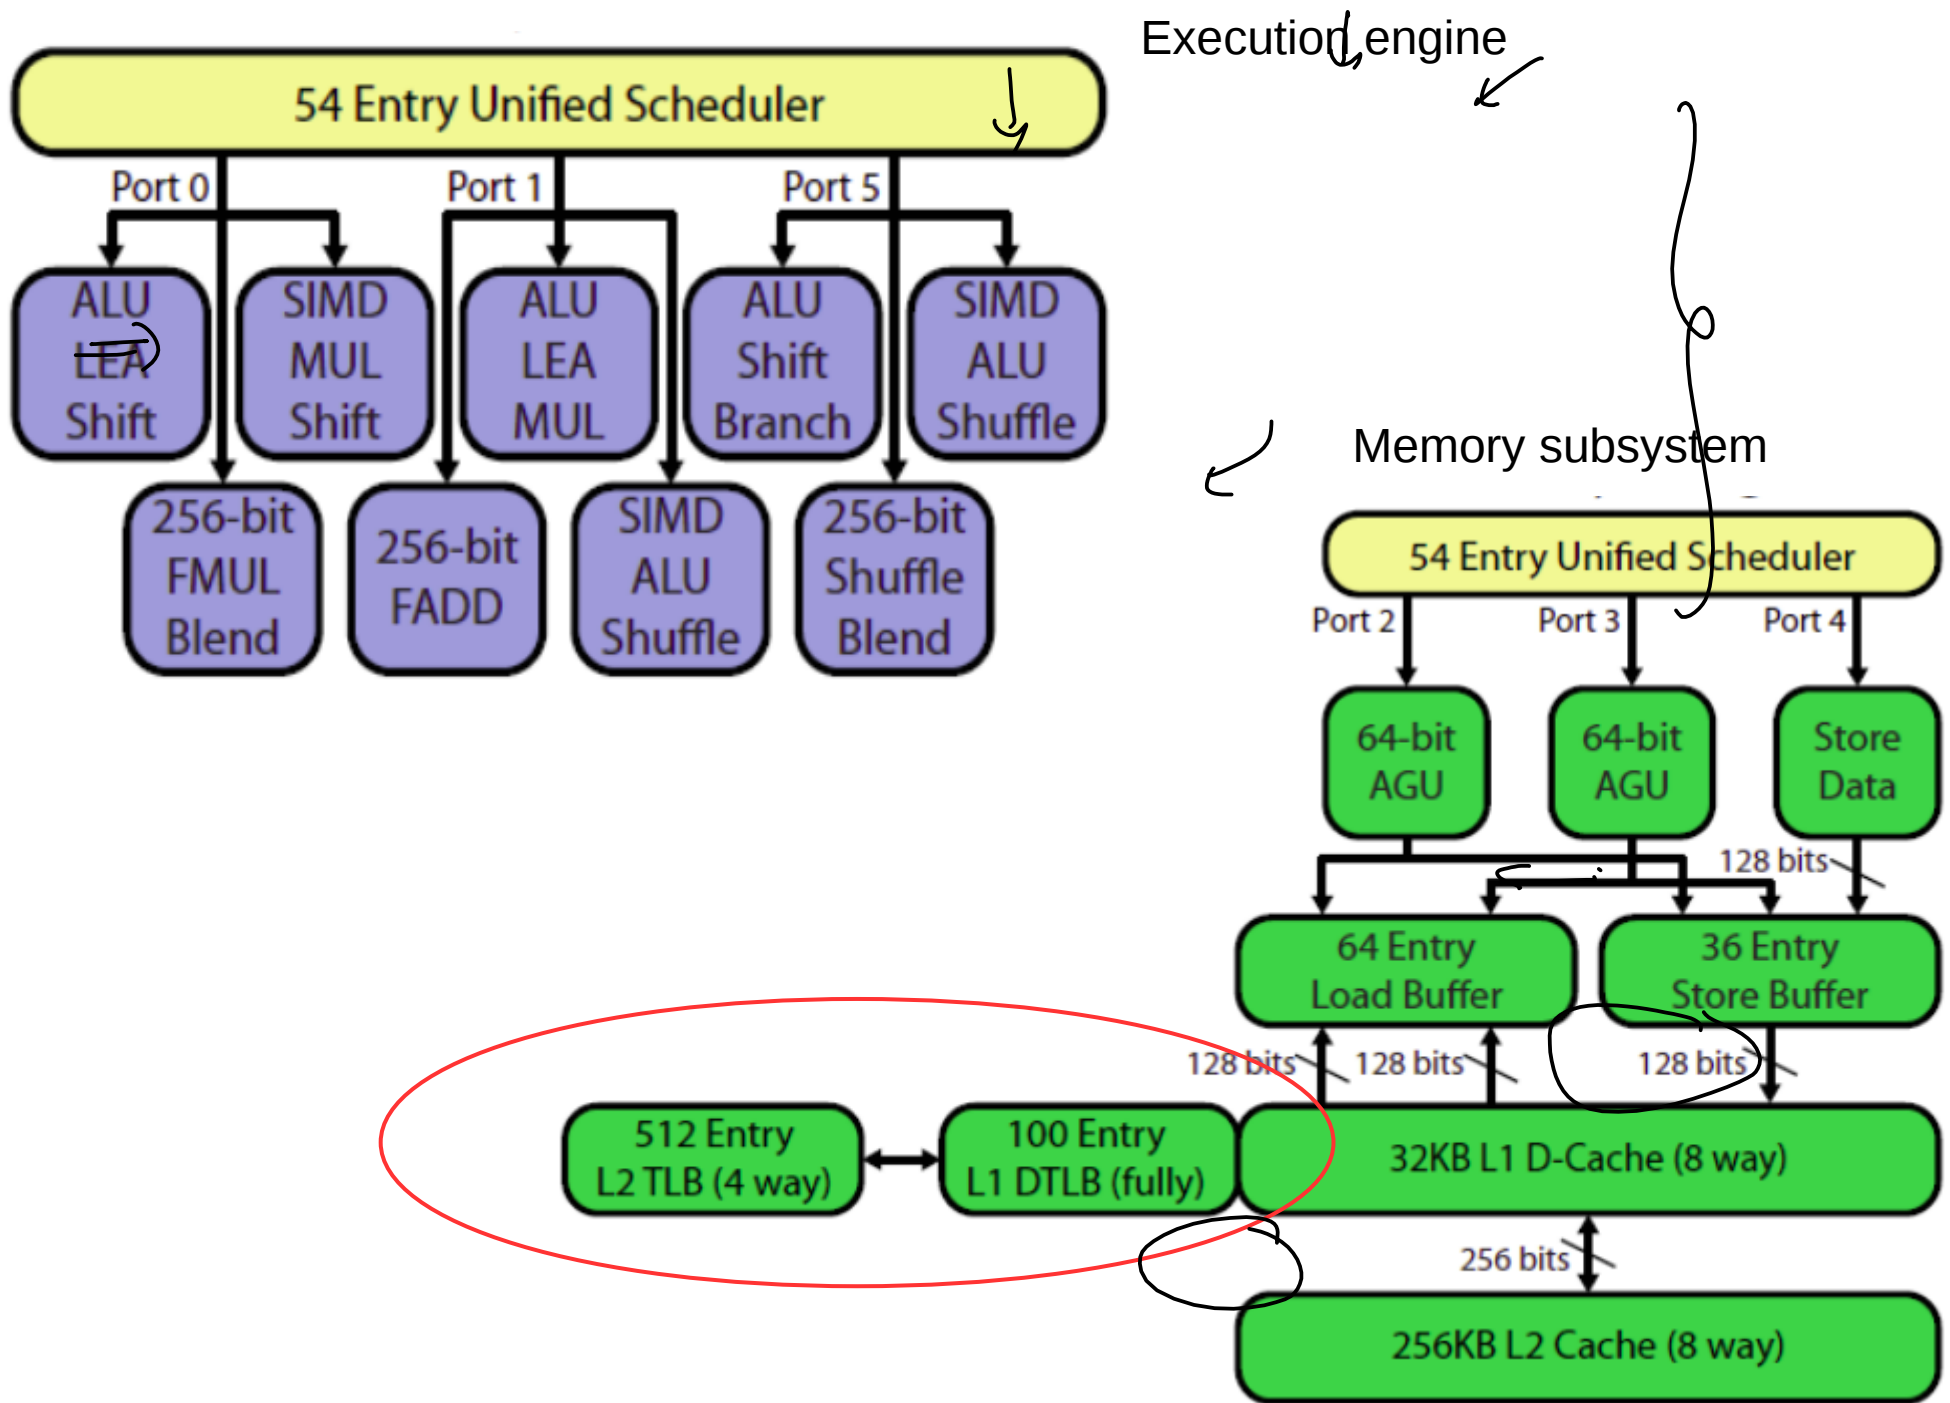
HPCA- Georgia Tech- Virtual memory

[https://www.youtube.com/watch?](https://www.youtube.com/watch?v=Dz9Hgq65iJw&list=PLAwxTw4SYaPn79fsplluZG34KwbkYSedj)

[v=Dz9Hgq65iJw&list=PLAwxTw4SYaPn79fsplluZG34KwbkYSedj](https://www.youtube.com/watch?v=Dz9Hgq65iJw&list=PLAwxTw4SYaPn79fsplluZG34KwbkYSedj)

Sandybridge architecture






Memory address space

- Processor's view of memory: 4GB
RAM – physical memory
- Programmers view of memory: MIPS promises each program a 32 bit address space – 2^{32} bytes – 4GB
- Multiple programs- Each program sees/thinks it has 4GB memory:
 2^{32} address space
- 64 bit address space: 2^{64} bytes: >> 4GB
- **Not enough RAM**


LW S1, 0(S2)
(4GB) 0(S2)
0+12



PA. \Rightarrow L1-Cache - L2-D \rightarrow RAM


Holes in physical memory

- Quit Program 1 --> Free 1GB space $\bigcirc \times 10 = 40GB$
- Assume program 3 needs 2GB space- cannot run
- Memory fragmentation



Security

- 2 programs access same address location unless sharing data
- Corrupt data/Crash
- Problem: Same memory space
- Solution: Give each program its own “virtual memory space”

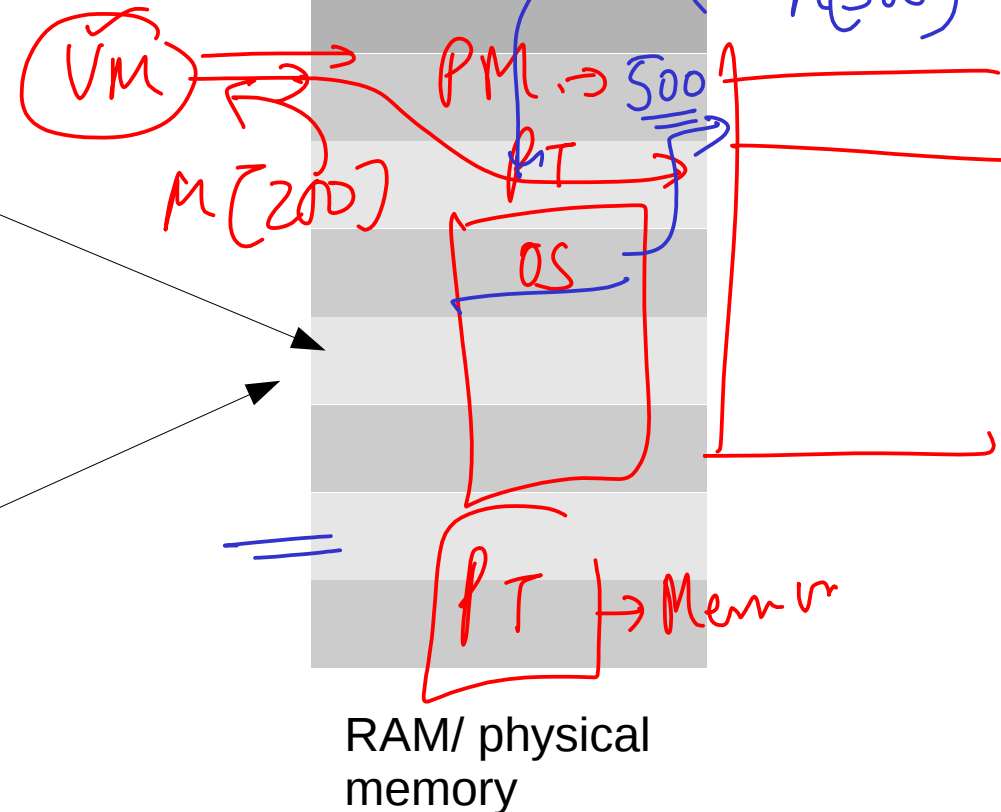




Virtual memory

- What is VM?
- Page tables and translation
- Page faults, Multi-level page tables
- Translation look aside buffer (TLB)
- TLB and Caches

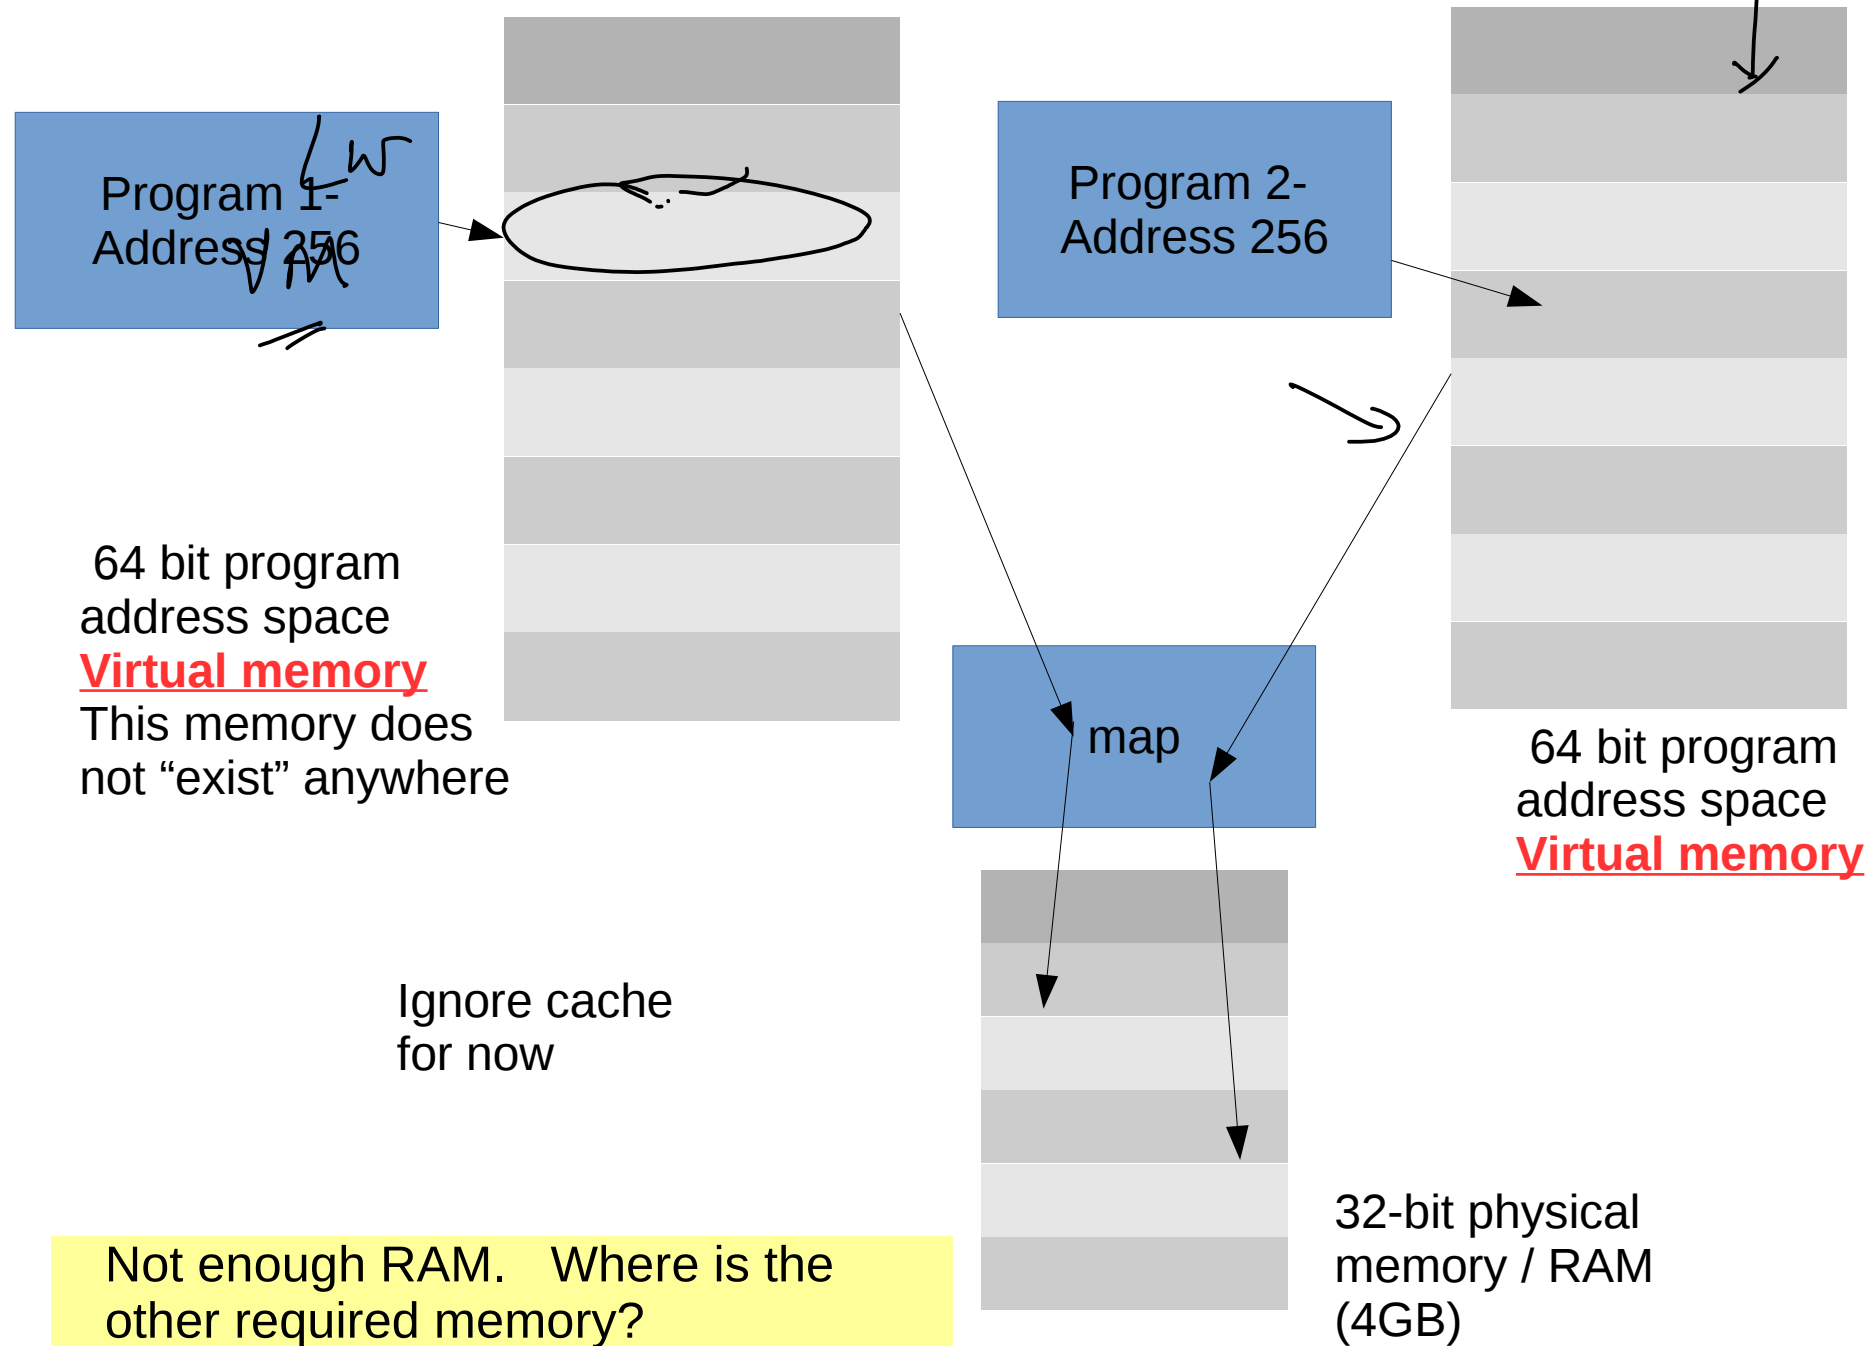
VM → 4GB
PM.

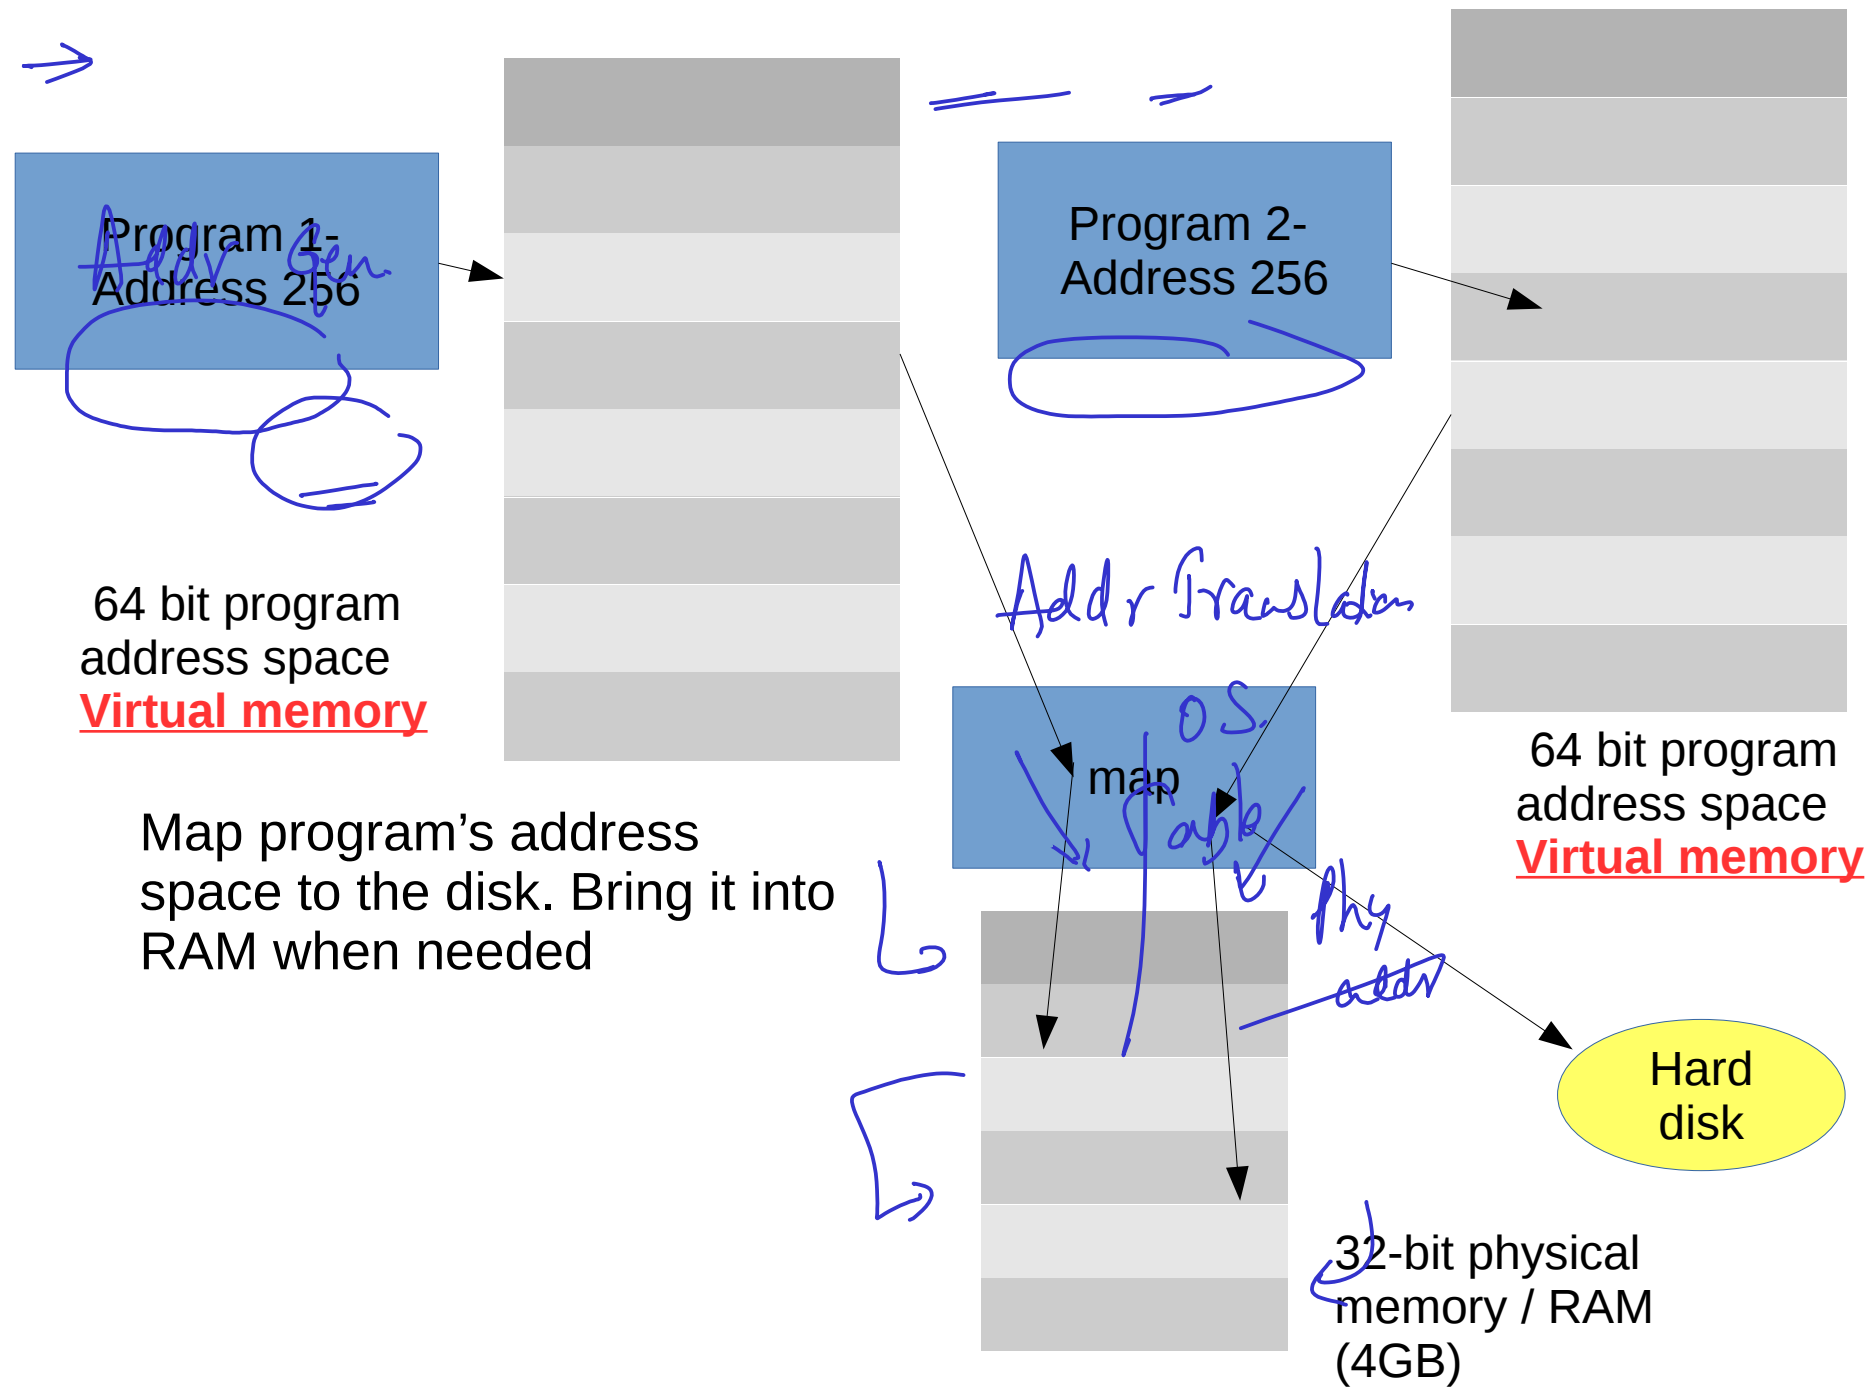
Physical address space



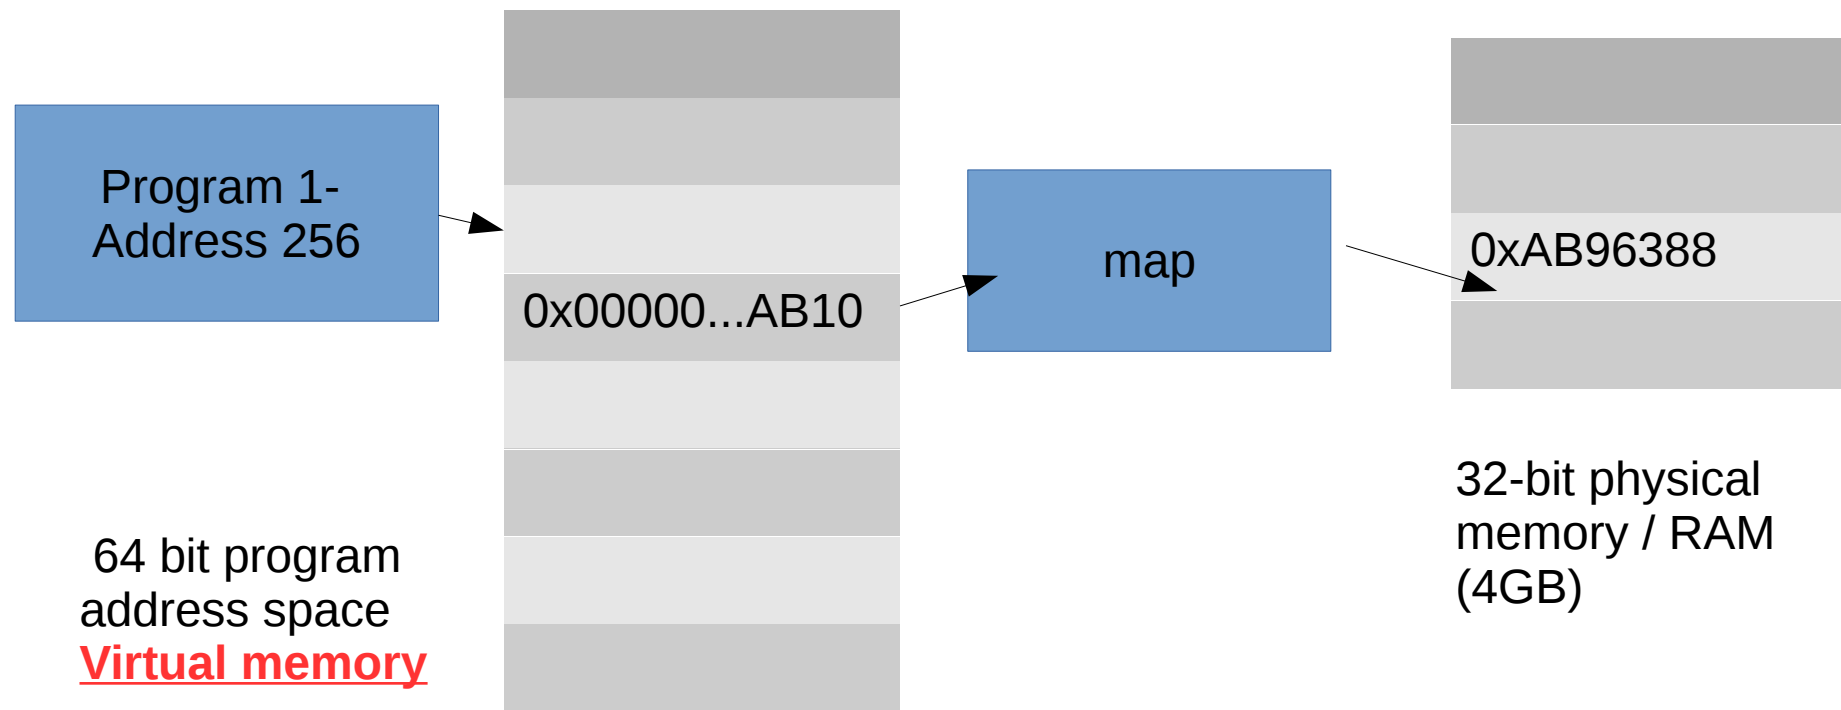
- Problem: Same memory space
 - Solution: Give each program its own “virtual memory space”
- > Map to the physical RAM memory space

- Without VM: Program address = RAM/physical address
- With VM: Program address *maps to* physical address





Mapping gives the illusion of an unlimited memory as long as Disk has 4GB space --> Every program can get 4GB of space. Data will be brought from disk to RAM



Problem of mapping:

- If we have a word aligned memory, how many mappings do you need?
 - 2^{30} --> 1 billion --> x 32 bits per entry (PA=32 bits) --> 1GB for mappings entries
- Mapping each virtual address to a physical address would result in a large number of mapping entries

Divide the program memory/virtual memory into equal sized "pages": Typically 4kB

2²⁰ virtual address
Program 1-
Address 256

64 bit program address space
Virtual memory

0 to 4kB	Page 0
4kB to 8kB	Page 1
To 12kB	Page 2
To 16kB etc	Page 3 etc

VA PA
0 → 2000
4 → 8000
2²⁰

map

0 to 4kB	Frame 0
4kB to 8kB	Frame 1 etc

4GB
32-bit physical memory / RAM (4GB)

Frame is of the same size as that of a page to hold a page

Program (LW) generates a Virtual address. Processor needs a physical address to access

VA.

Page-Frame

RAM phy mem

Program 1 VM

The OS does the mapping

0 to 4kB	Page 0
4kB to 8kB	Page 1
To 12kB	Page 2
.	.
.	.
To 32kB	Page 7 etc

Page table- translates virtual to physical address

page	frame
0	2
1	0
2	
3	1

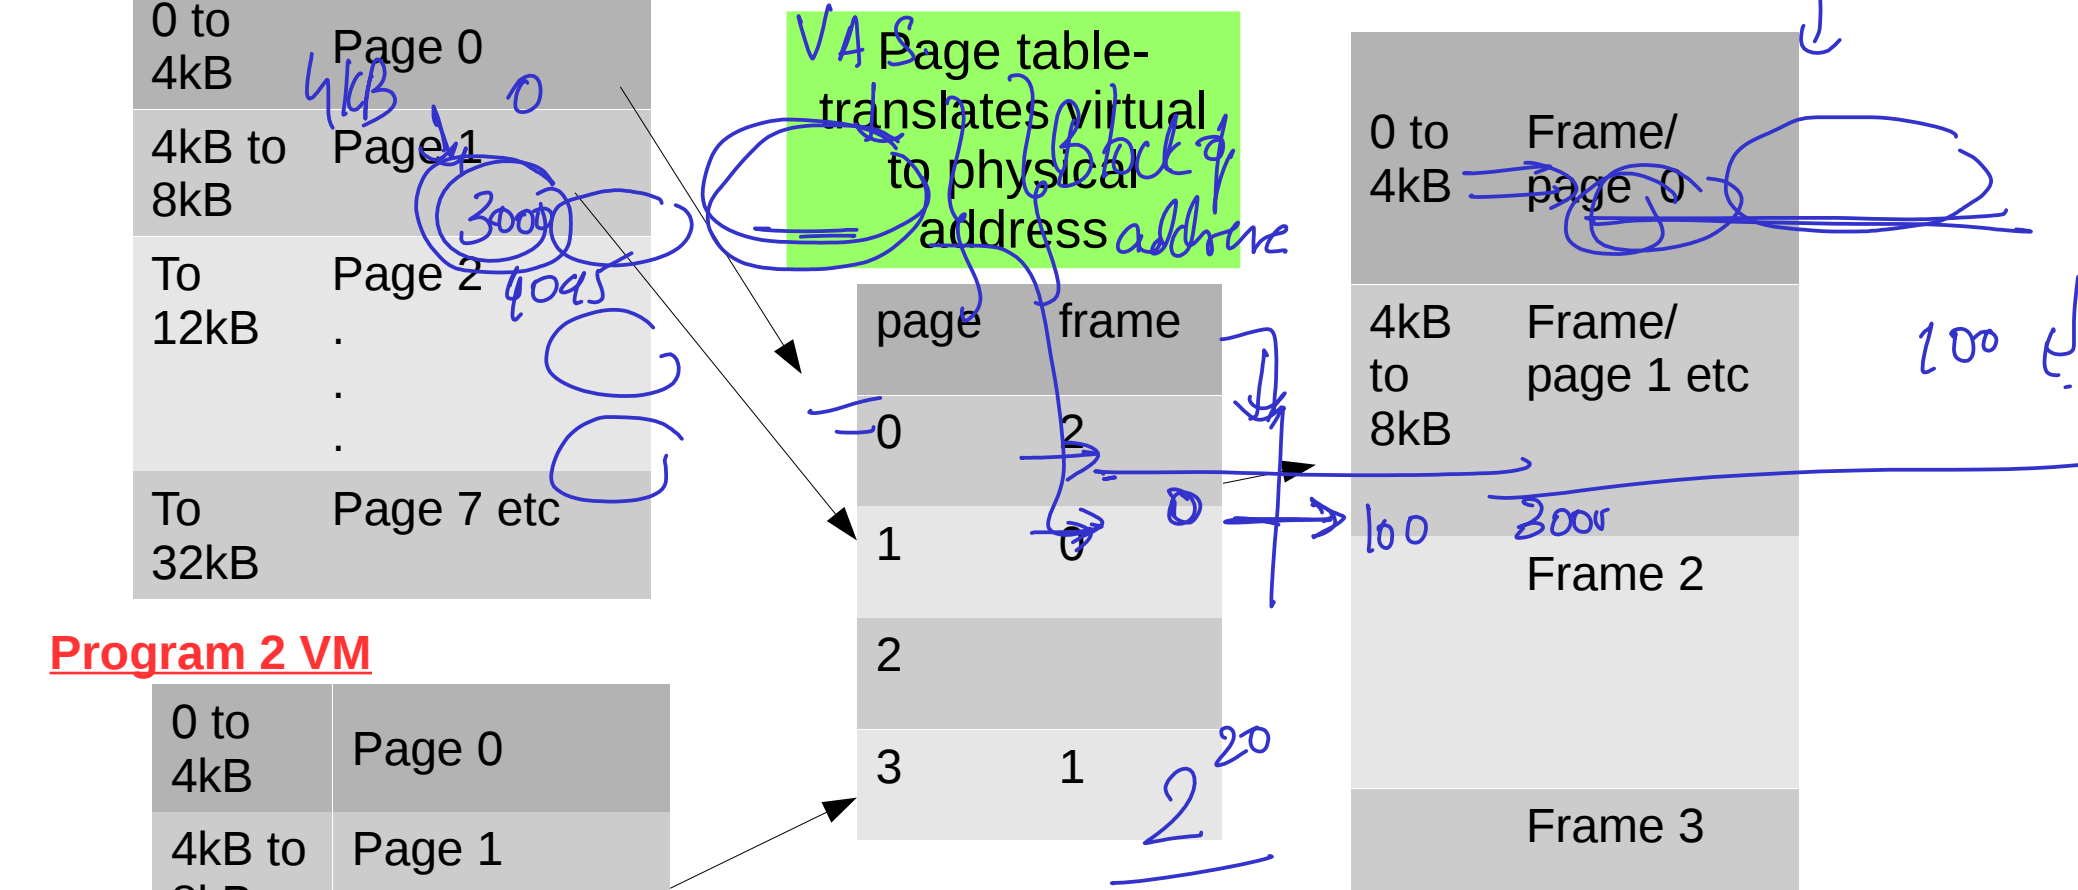
0 to 4kB	Frame/ page 0
4kB to 8kB	Frame/ page 1 etc
	Frame 2
	Frame 3

Program 2 VM

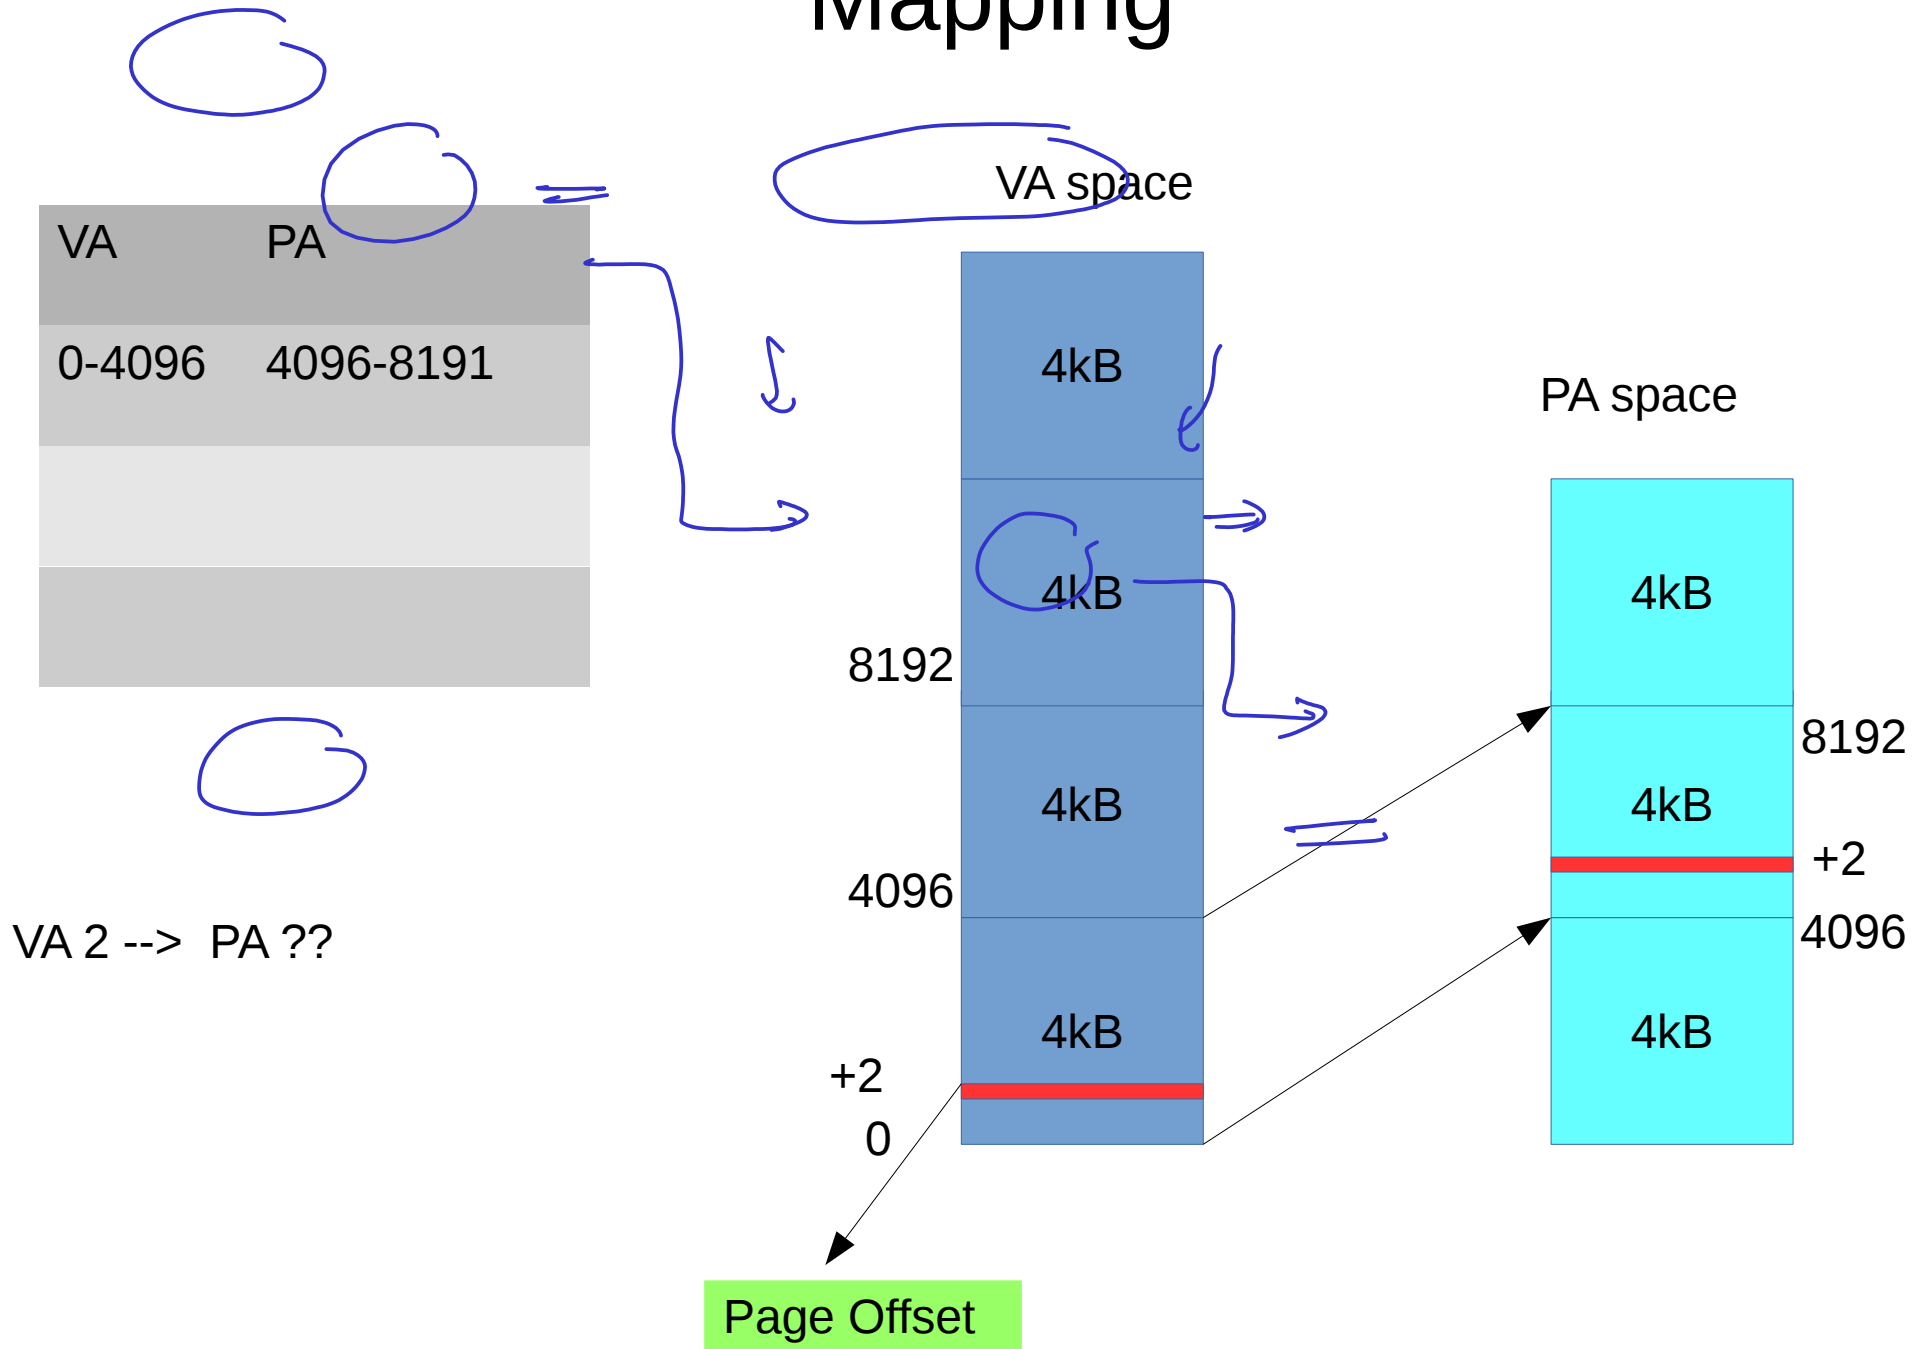
0 to 4kB	Page 0
4kB to 8kB	Page 1
To 12kB	Page 2
.	.
.	.
To 32kB	Page 7 etc

Page table: Stored in physical memory/ RAM

32-bit physical memory / RAM (4GB)

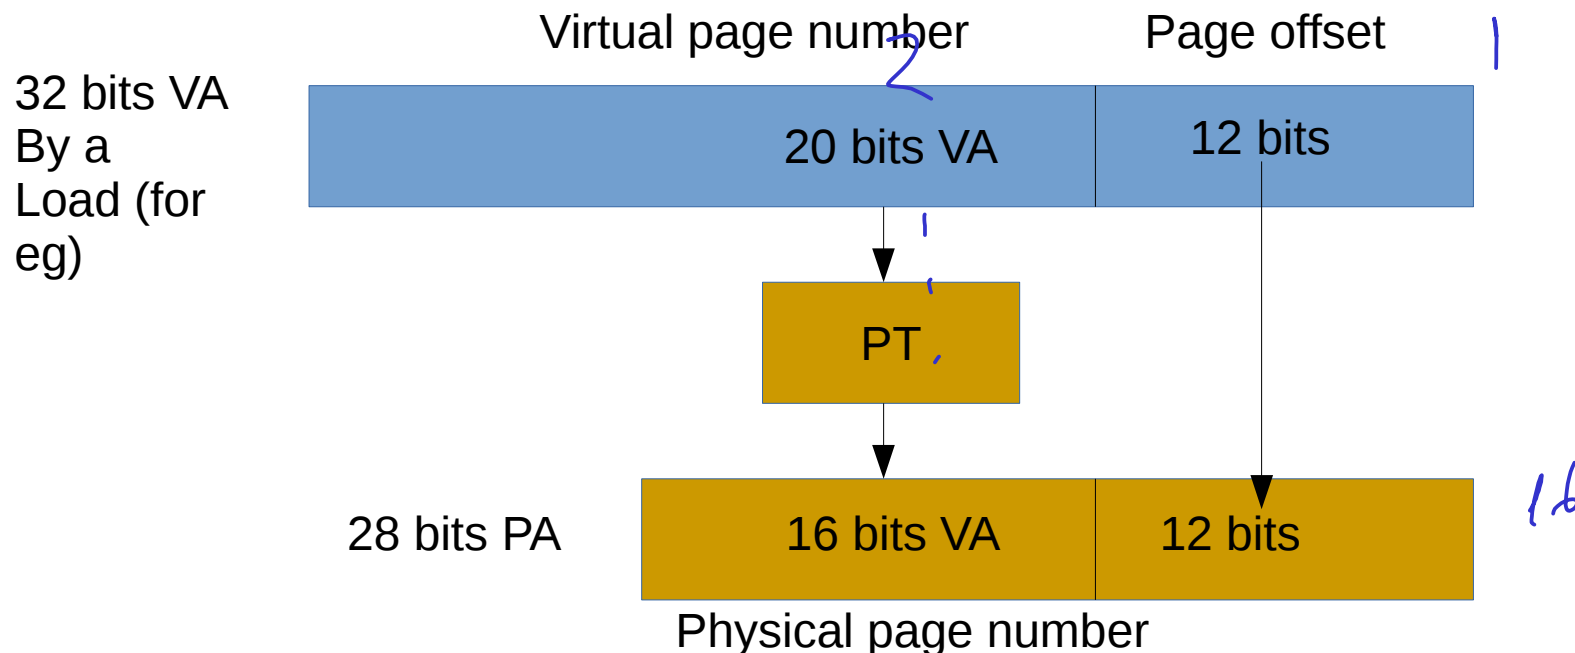


Mapping

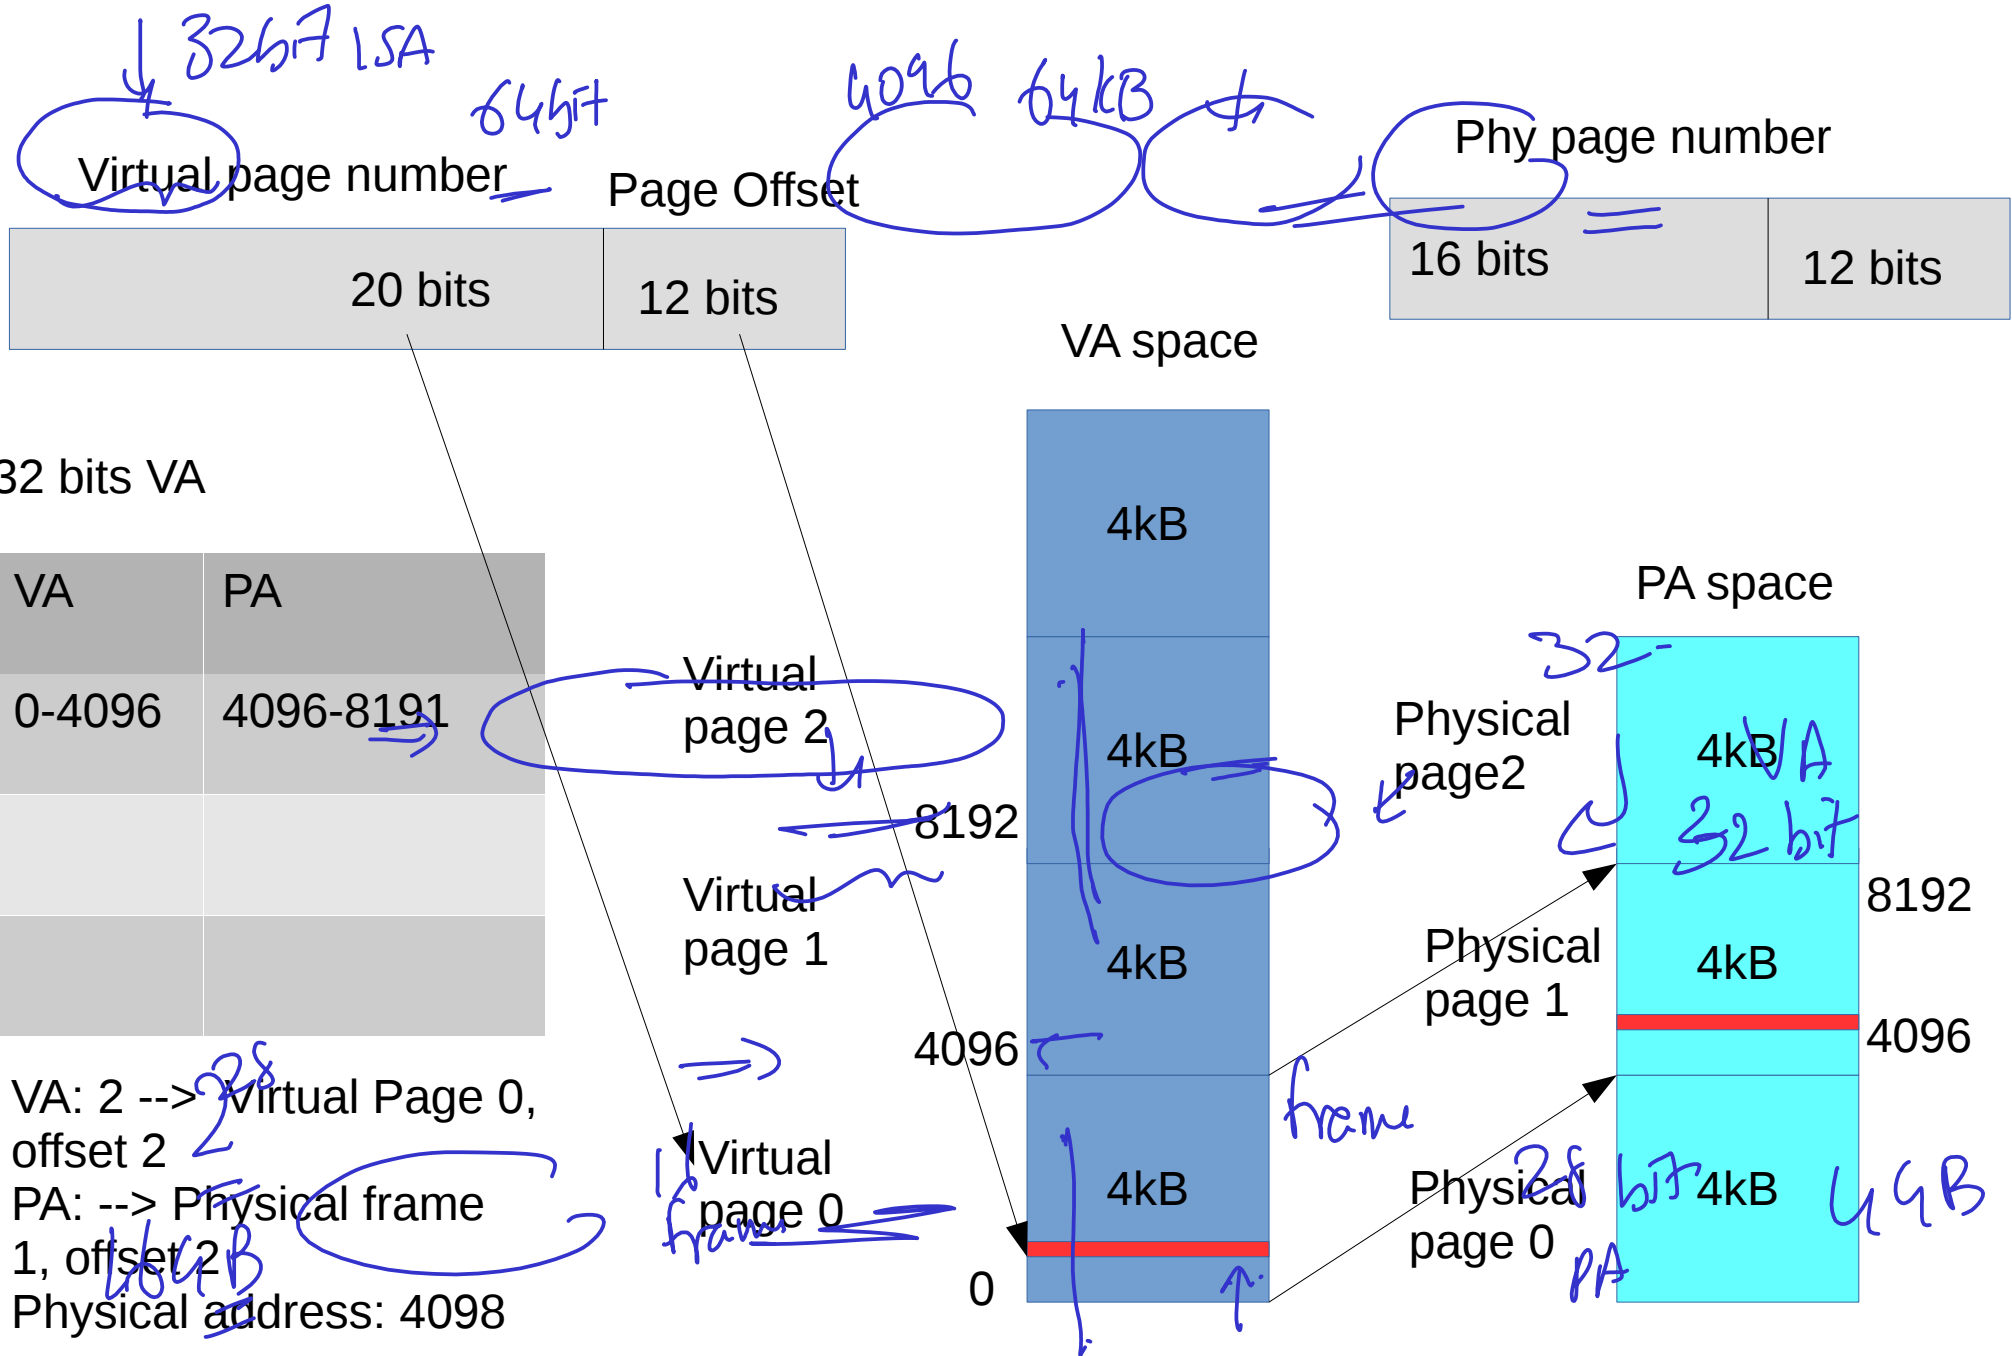


Address translation

- 32 bit virtual address space, 4kB page, 256MB (2^{28}) RAM
 - Each page has 4096 addresses – Each address do not need translation. It is part of a page offset- Form the lower 12 bits of VA
 - Each page needs a translation
 - Or in other words: $2^{32} / 4096 = 2^{20}$ pages --> These page numbers need translation to physical address. These form 20 bits of VA

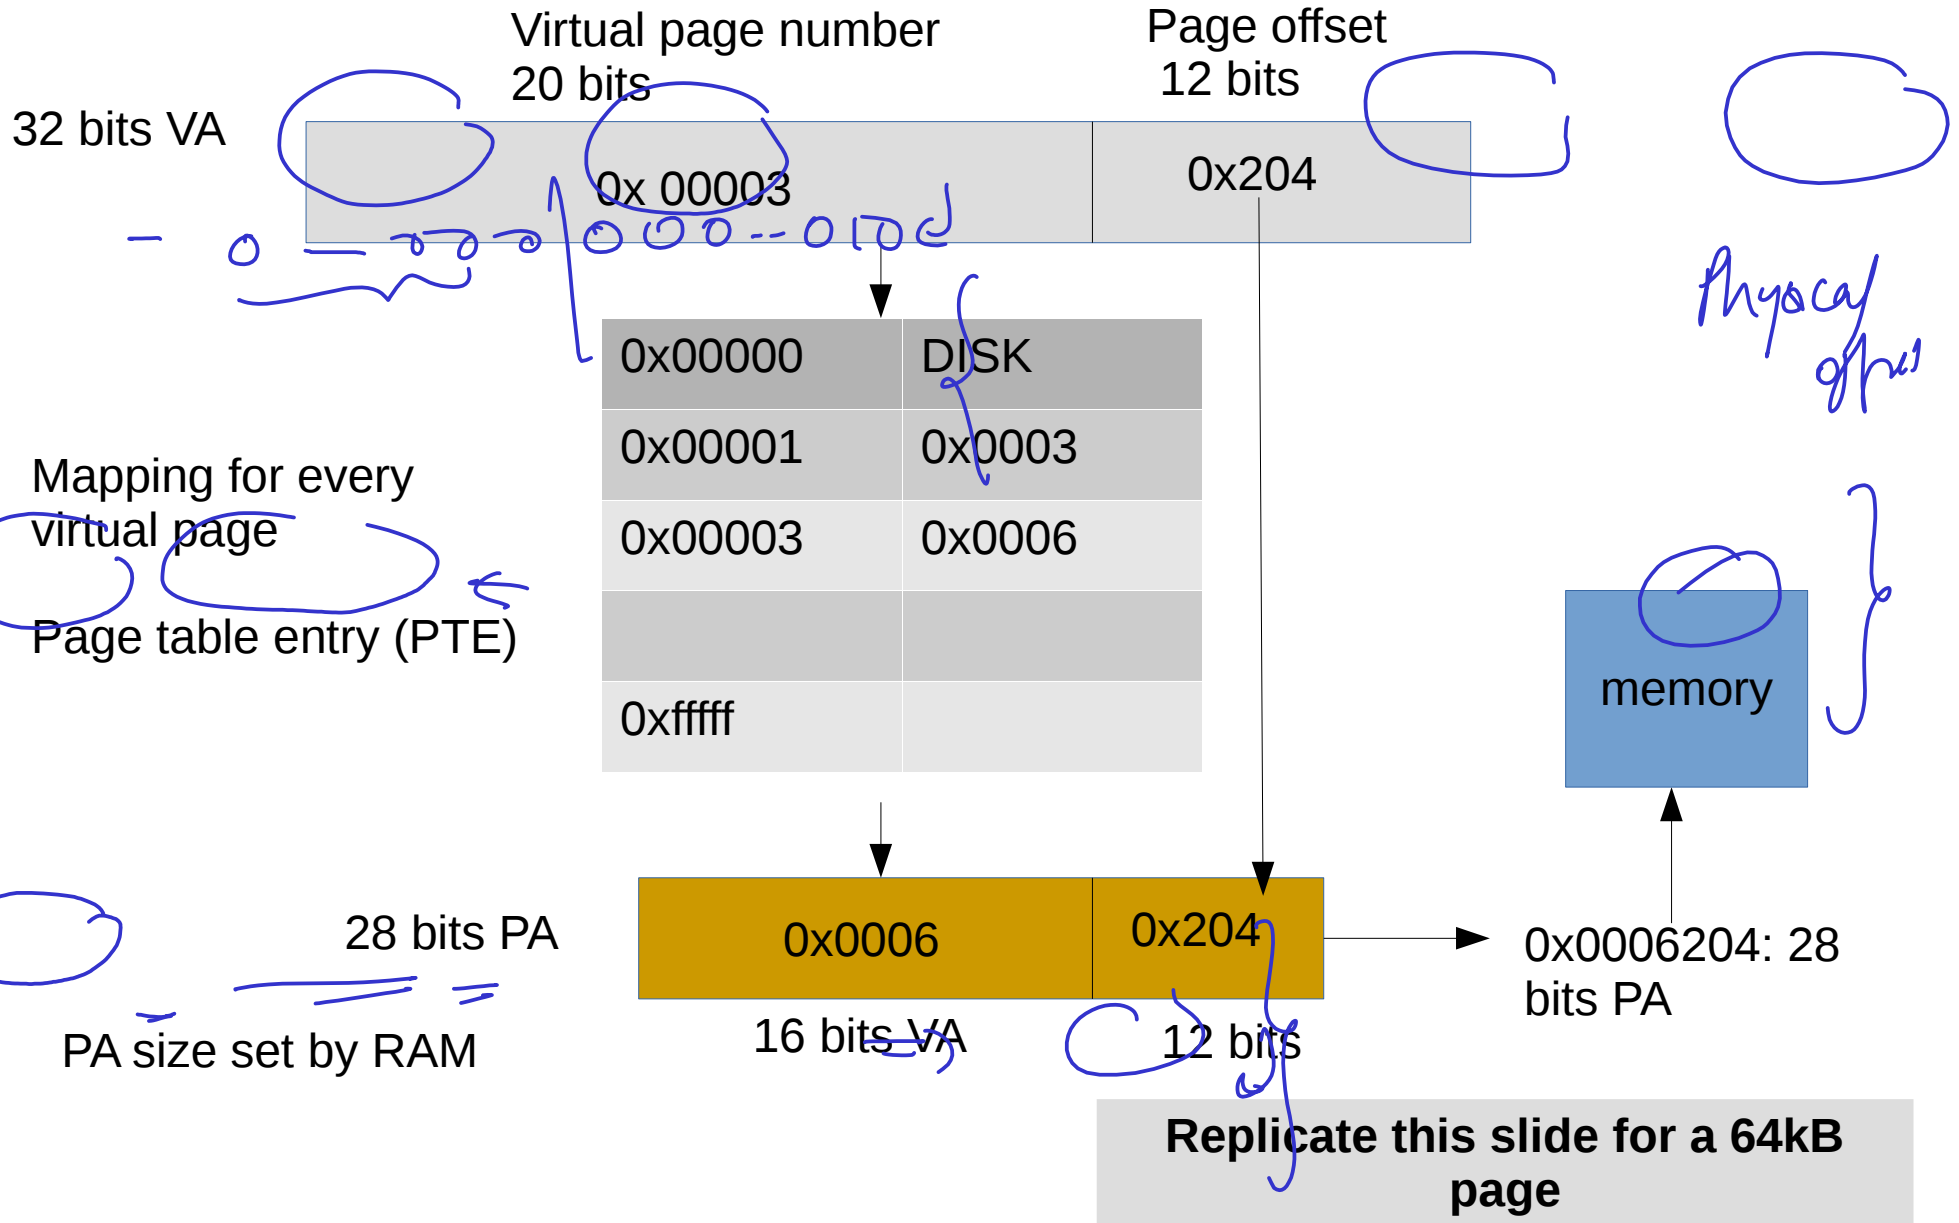


Mapping



Page table lookup

32 bit VA: 0x 00003204



Page table lookup


4kB page

256 MB RAM
 2^{28}

32 bits VA

Virtual page number

Page offset



Page fault

Mapping for every virtual page

0x00000	DISK
0x00001	0x0003
0x00003	0x0006
0xffff	

0x00000204 --> ??

Cache → RAM → HDD


2^{32}

28 bits PA

16 bits VA

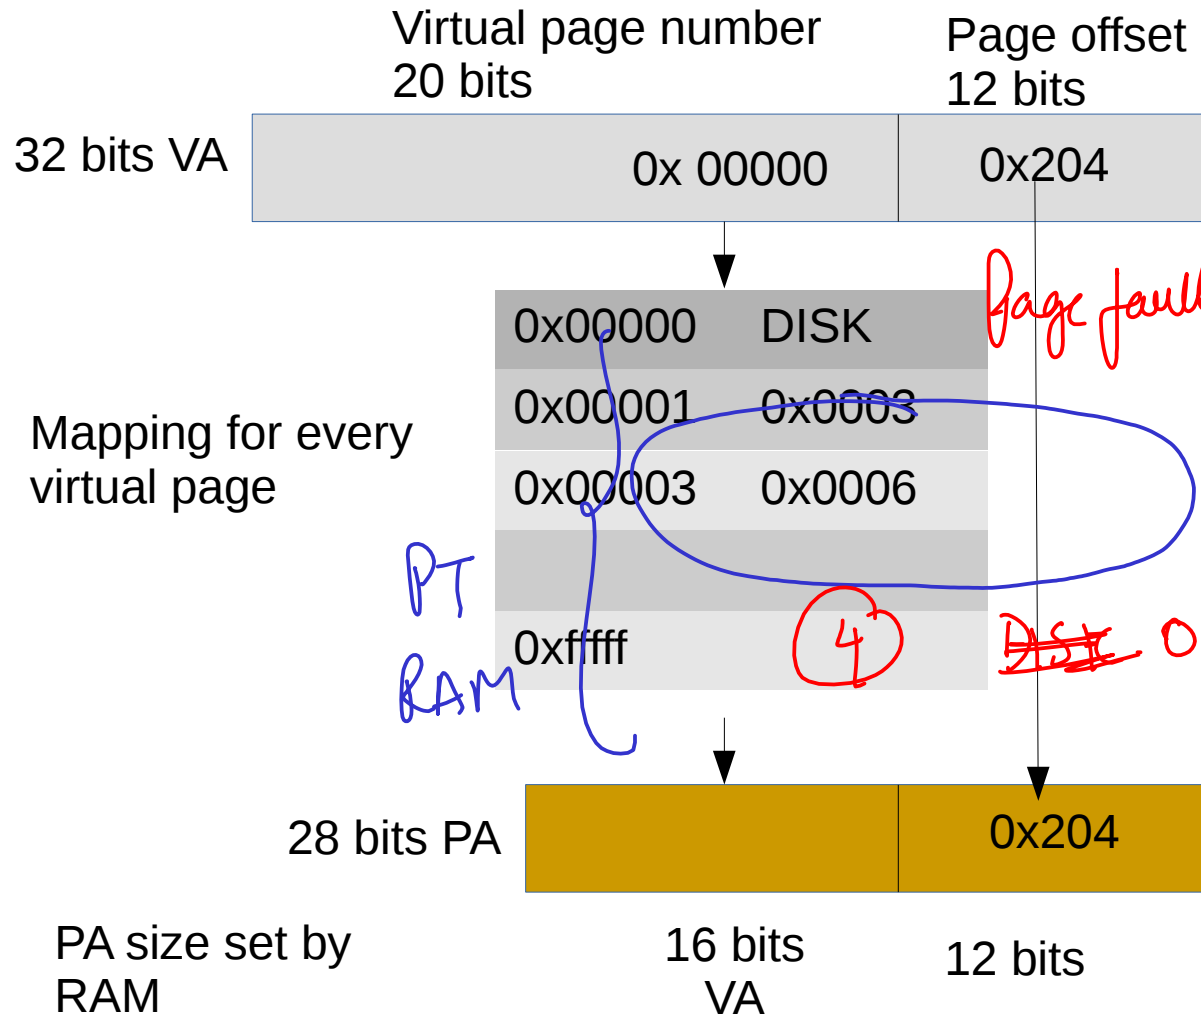
12 bits

PA size set by RAM



Page fault

VA: 0x00000204



- 1) LW R3, 40(R2) --> Address is the virtual address
- 2) PT Lookup, Compute physical address
- 3) 0x00000204 --> Data not in RAM --> CPU generates **page fault** --> exception --> OS page fault handler
- 4) OS chooses a victim to evict from RAM (based on some replacement policy), writes it to disk if the page is dirty
- 5) OS invokes a disk interrupt routine --> reads page to RAM and updates PTE

VA
↓

TLB → Miss → PT RAM.

Page fault
↓

000 → RAM
PA

OS

03
03
03

04
08


PA

RAM and updates PTE


D-cache

Page fault


VA: 0x00000204



Mapping for every
virtual page



28 bits PA




PA size set by
RAM

16 bits
VA


12 bits

- 1) Read page from the disk to RAM and OS updates PTE
- 2) Hand it over to CPU
- 3) No page fault now



Page fault – time taken

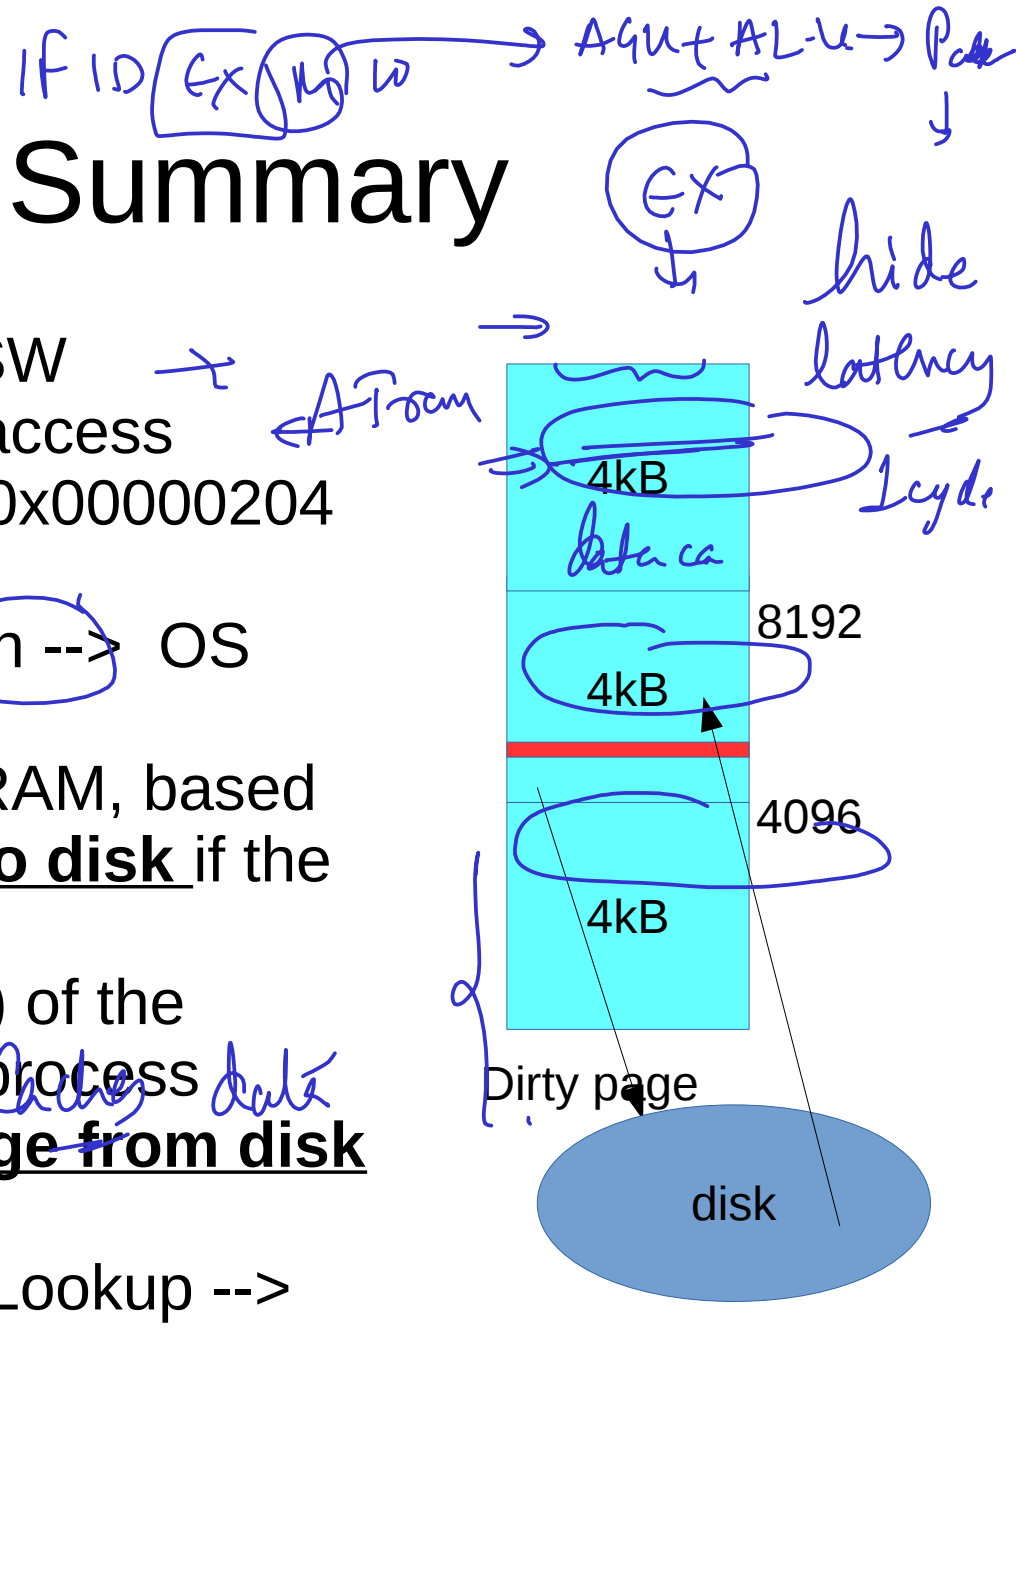
- Compute virtual address from LW/SW
- PT lookup (PT in RAM) = **memory access** and Compute physical address of 0x00000204
- Say Data not in RAM
- CPU generates page fault exception -->
- OS page fault handler
- OS chooses a victim to evict from RAM, writes it to disk if the page is dirty.
- OS read page from disk to RAM and updates PTE
- Hand it over to CPU
- PT Lookup --> No fault



TLB

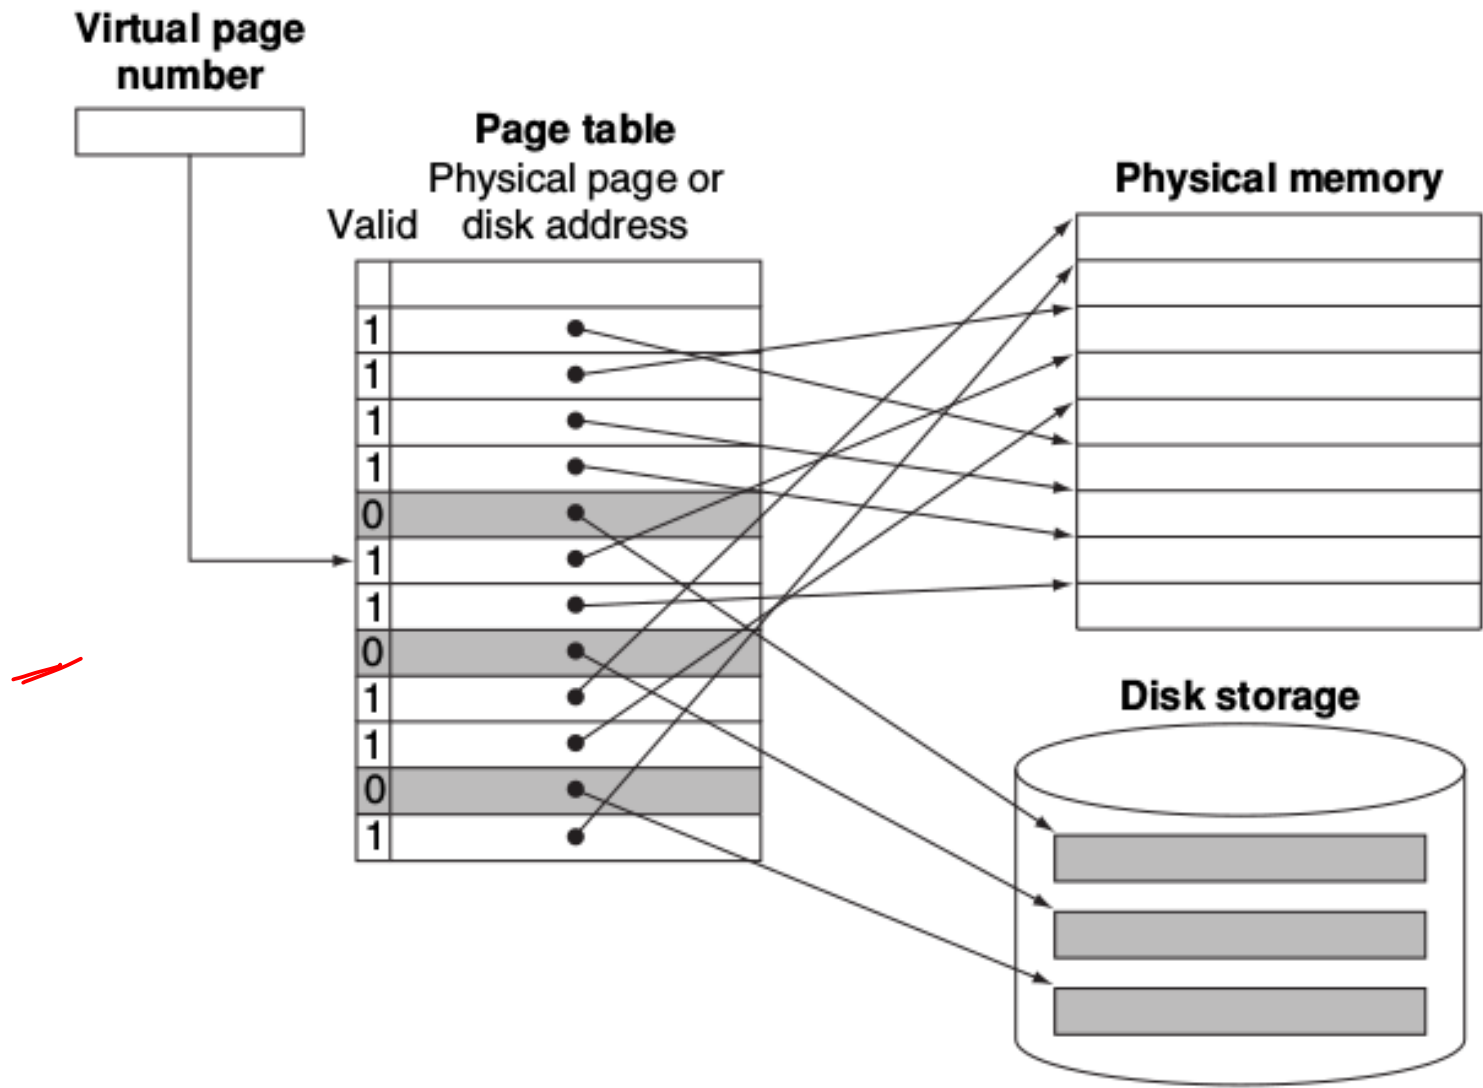
Page fault: Summary

- Compute virtual address from LW/SW
- PT lookup (PT in RAM) = memory access and Compute physical address of 0x00000204
- Say Data not in RAM
- CPU generates page fault exception --> OS page fault handler
- OS chooses a victim to evict from RAM, based on its replacement policy, writes it to disk if the page is dirty.
- Save the context (registers, PC etc) of the current process, switch to another process
- Disk interrupt routine --> reads page from disk to RAM and updates PTE
- Faulting instruction is restored, PT Lookup --> No fault



Role of OS

- Page fault:
 - How does the OS know the location of the page on disk, given the virtual address?
- For each page of a process, in the virtual address space, OS keeps track of the location on disk in a “swap space” on disk. This happens when it creates a process
 - This information can also be part of PT with the address on disk



Writes to pages in memory

- Note that physical memory or main memory acts as a cache for the ~~hard~~ disk
- Virtual memory systems must use write-back, performing writes into the page in memory, and copying the page back to disk when it is replaced in the memory.
 - “Dirty bit” concept

Page table entries




- Consider a physical memory of size 4GB
Virtual memory 32GB. Page size: 4kB
- How many page table entries does the page table has?

Eg


How many bits are needed for each of the following: Virtual address, Physical address, Virtual page number, Physical page number, Page Offset

- 32-bit ISA, 4-KB pages, 1 GB of RAM
- 32-bit ISA, 16-KB pages, 2 GB of RAM
- 64-bit ISA, 16-KB pages, 16 GB of RAM

Soln



	VA	PA	VPN	PPN	Offset
1	32	30	20	18	12
2	32	31	18	17	14
3	64	34	50	20	14



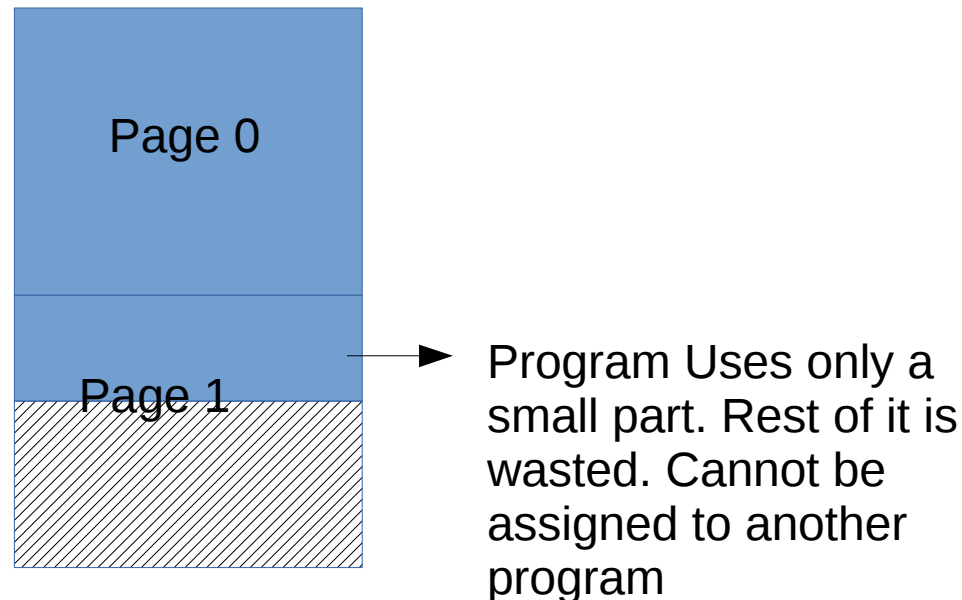
2^{14}

2^7

by 34
 $64-14$
 $34-14$

Page size

- Small pages: Large Page table
- Large pages: Small PT, but will have to replace pages in large chunks.
 - Also suffer from internal fragmentation
- Typical sizes: Few kB to a few MB (4kB, 32MB)



Problem 1: Size of a flat page table

- Page tables are stored in RAM/physical memory
- 1 entry per page in Virtual address space
- Number of page table entries (PTE) = Virtual memory size / Page size
 - 4GB/4kB
- PT Size = No of PTE * Size of each entry
 - $(4\text{GB}/4\text{kB}) * 4\text{B} = 4\text{MB}$ per program or per process
 - If Virtual address space is 64 bits, then PT size will not even fit in RAM! ($4 * 10^{15}$)
- Solution: Multilevel page tables


Problem 2: Virtual address space

64-bit VAS (4GB) VM 4GB

⇒ TLB


PT → RAM

4kB page size =



(4GB)

PT




The entire 4GB Virtual address space will not be used a program

A program uses only a limited part of this address space: Code, Data, Stack, heap

But, the Page table entries will still be available for the entire space (like the unused part, kernel space etc) and **these occupy space**


Can we omit these page table entries?



Multi-level PT

Paging the page tables

1st level PT



2nd level PT

Memory (RAM)

Translation with 2 level PTs, 4kB pages

Handwritten notes: $VPN \rightarrow PPN$, $05 \rightarrow 04$

Virtual address

0x00402204

Virtual page number
0x00402

Page offset
0x204

Outer pg no or
1st level PT = 01

Inner pg no = 10

PT 10

0000 0000 01

00 0000 0010

0x204

00 0000 0010

0001

0002

Outer page table/ 1st
level PT

2nd level PT

0x0123

PT0

0x0

0x0123

PT1

0x1

0x0003

0x1

0x0123

0x0003

0x1298

PT2

Frame
number

100

0x1298

0x204

16 bits VA

Phy page number

Overall size remains same as we still need to
have entries for 32 bits VA space!
Infact, we made it more complicated.

Program/
Virtual addr
space

RAM

PT 99

1024

Program/
Virtual addr
space

1024



Overall size remains same as we still need to
have entries for 32 bits VA space!
Infact, we made it more complicated.

Unused address space/ Reserved address space

$LW\ S1,0(S2)$
32bit

MSB
0000 0000 0000 0000 0010
10bit 10bit 10bit 10bit
2nd level PT

$4096 = 2^{12}$



X Entry not needed

VPN 20bit

Outer page table/ 1st level PT

00
01
11
100

0x0123	
0x0	0x0123
0x1	0x0003
0x1	x
	x
	x

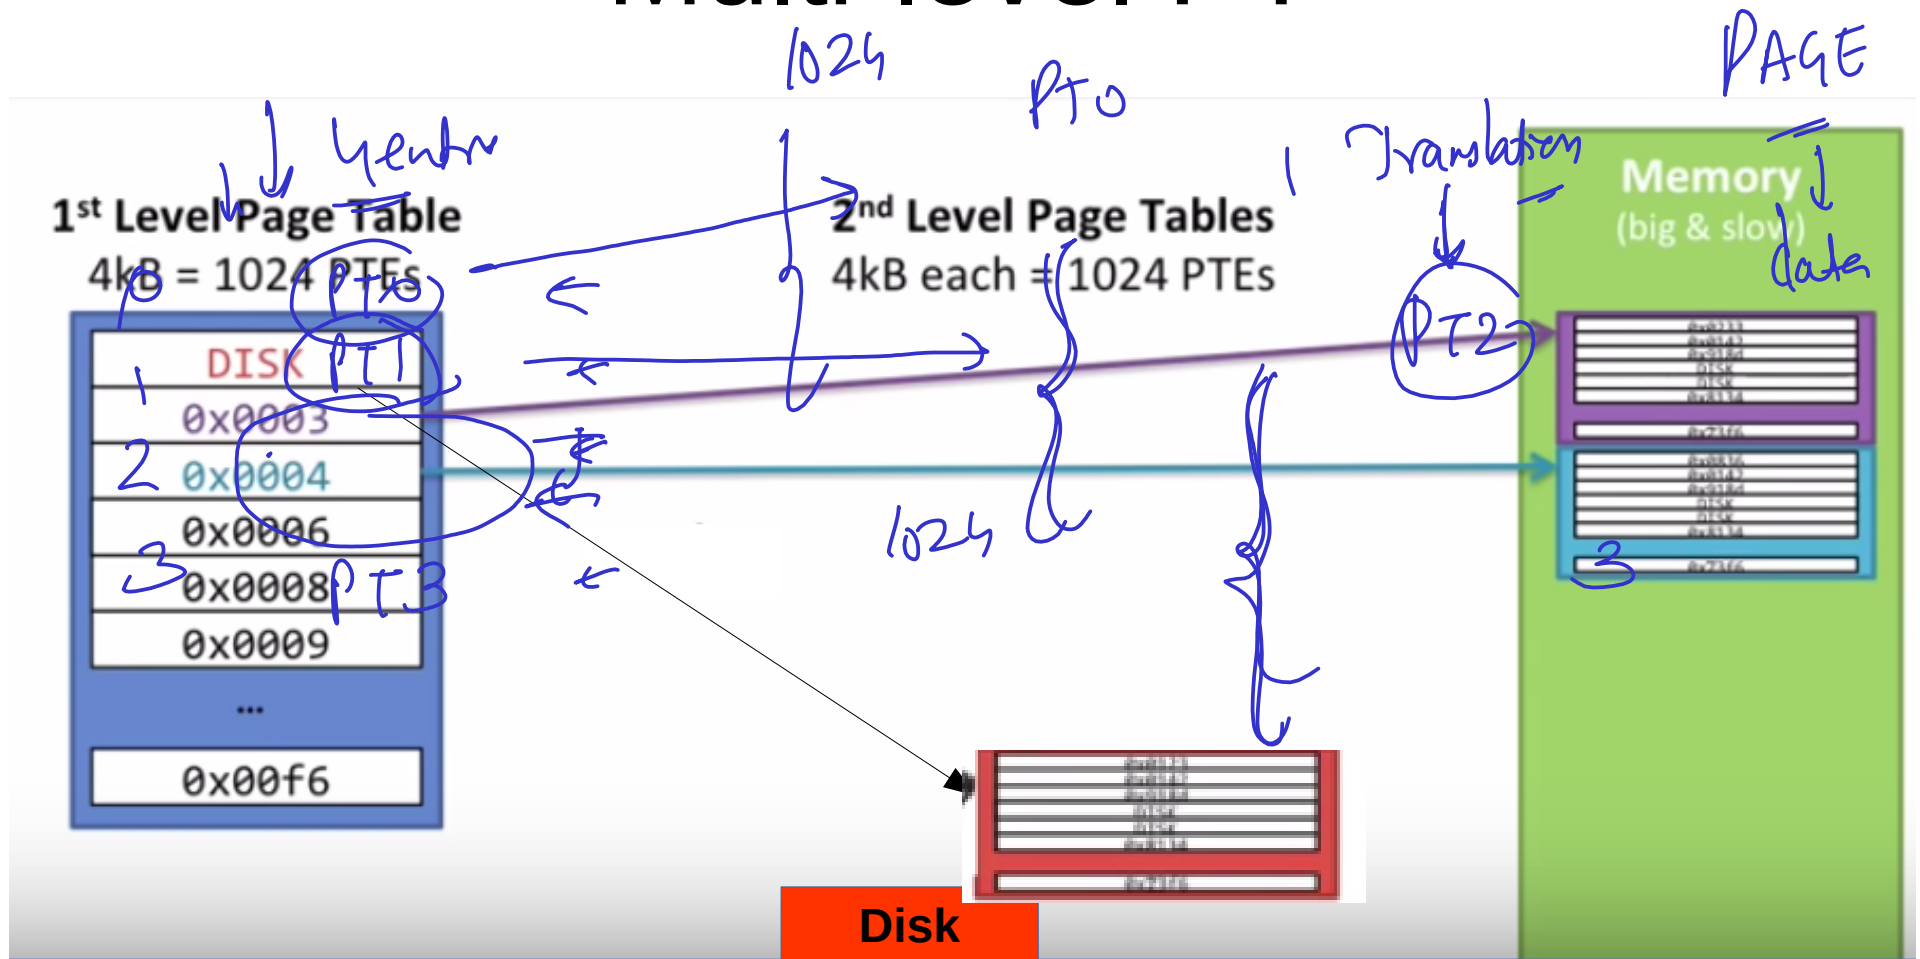
Can eliminate this entire PT

x	0x0123
x	0x0003
	0x1298

Phy page number

Lots of inner page tables may not be present.
And can be eliminated.
This will not be possible in a flat page table

Multi-level PT



1st level PT should be in RAM, 2nd level can be on disk, since we can easily find the PA through 1st level PT

Single vs 2 level PT

- 32 bit ^{PT} VA, 4kB page, 8 bytes per entry
- 1024 entry outer ^{0x00005} and inner page tables
- Program ^{PT 2nd level} uses VM from address 0 to 0x00010000

Single or flat Page Table size:

Virtual address size/Page size = No of pages

$2^{32} \text{ VA} / 2^{12} = 2^{20} \text{ pages or PT entries}$


$2^{20} * 8 \text{ bytes per entry} = 8\text{MB}$

Page fault

Memory

Single vs 2 level PT

- 32 bit VA, 4kB page, 8 bytes per entry
- 1024 entry outer and inner page tables
- Program uses VM from address 0 to 0x00010000



0 to 0x00010000

Virtual address: 32 bits

Outer pg size	Inner pg size	offset
10	10	12

0x11110000

32 bit virtual address

0000 0000 00 0000... 0000
 0000 0000 00 00 0001 000... 0000
 Outer pg no | Inner pg no | Offset

Only Outer Page no – 0 being used for this entire address range

Others can be left: x

L1
1024

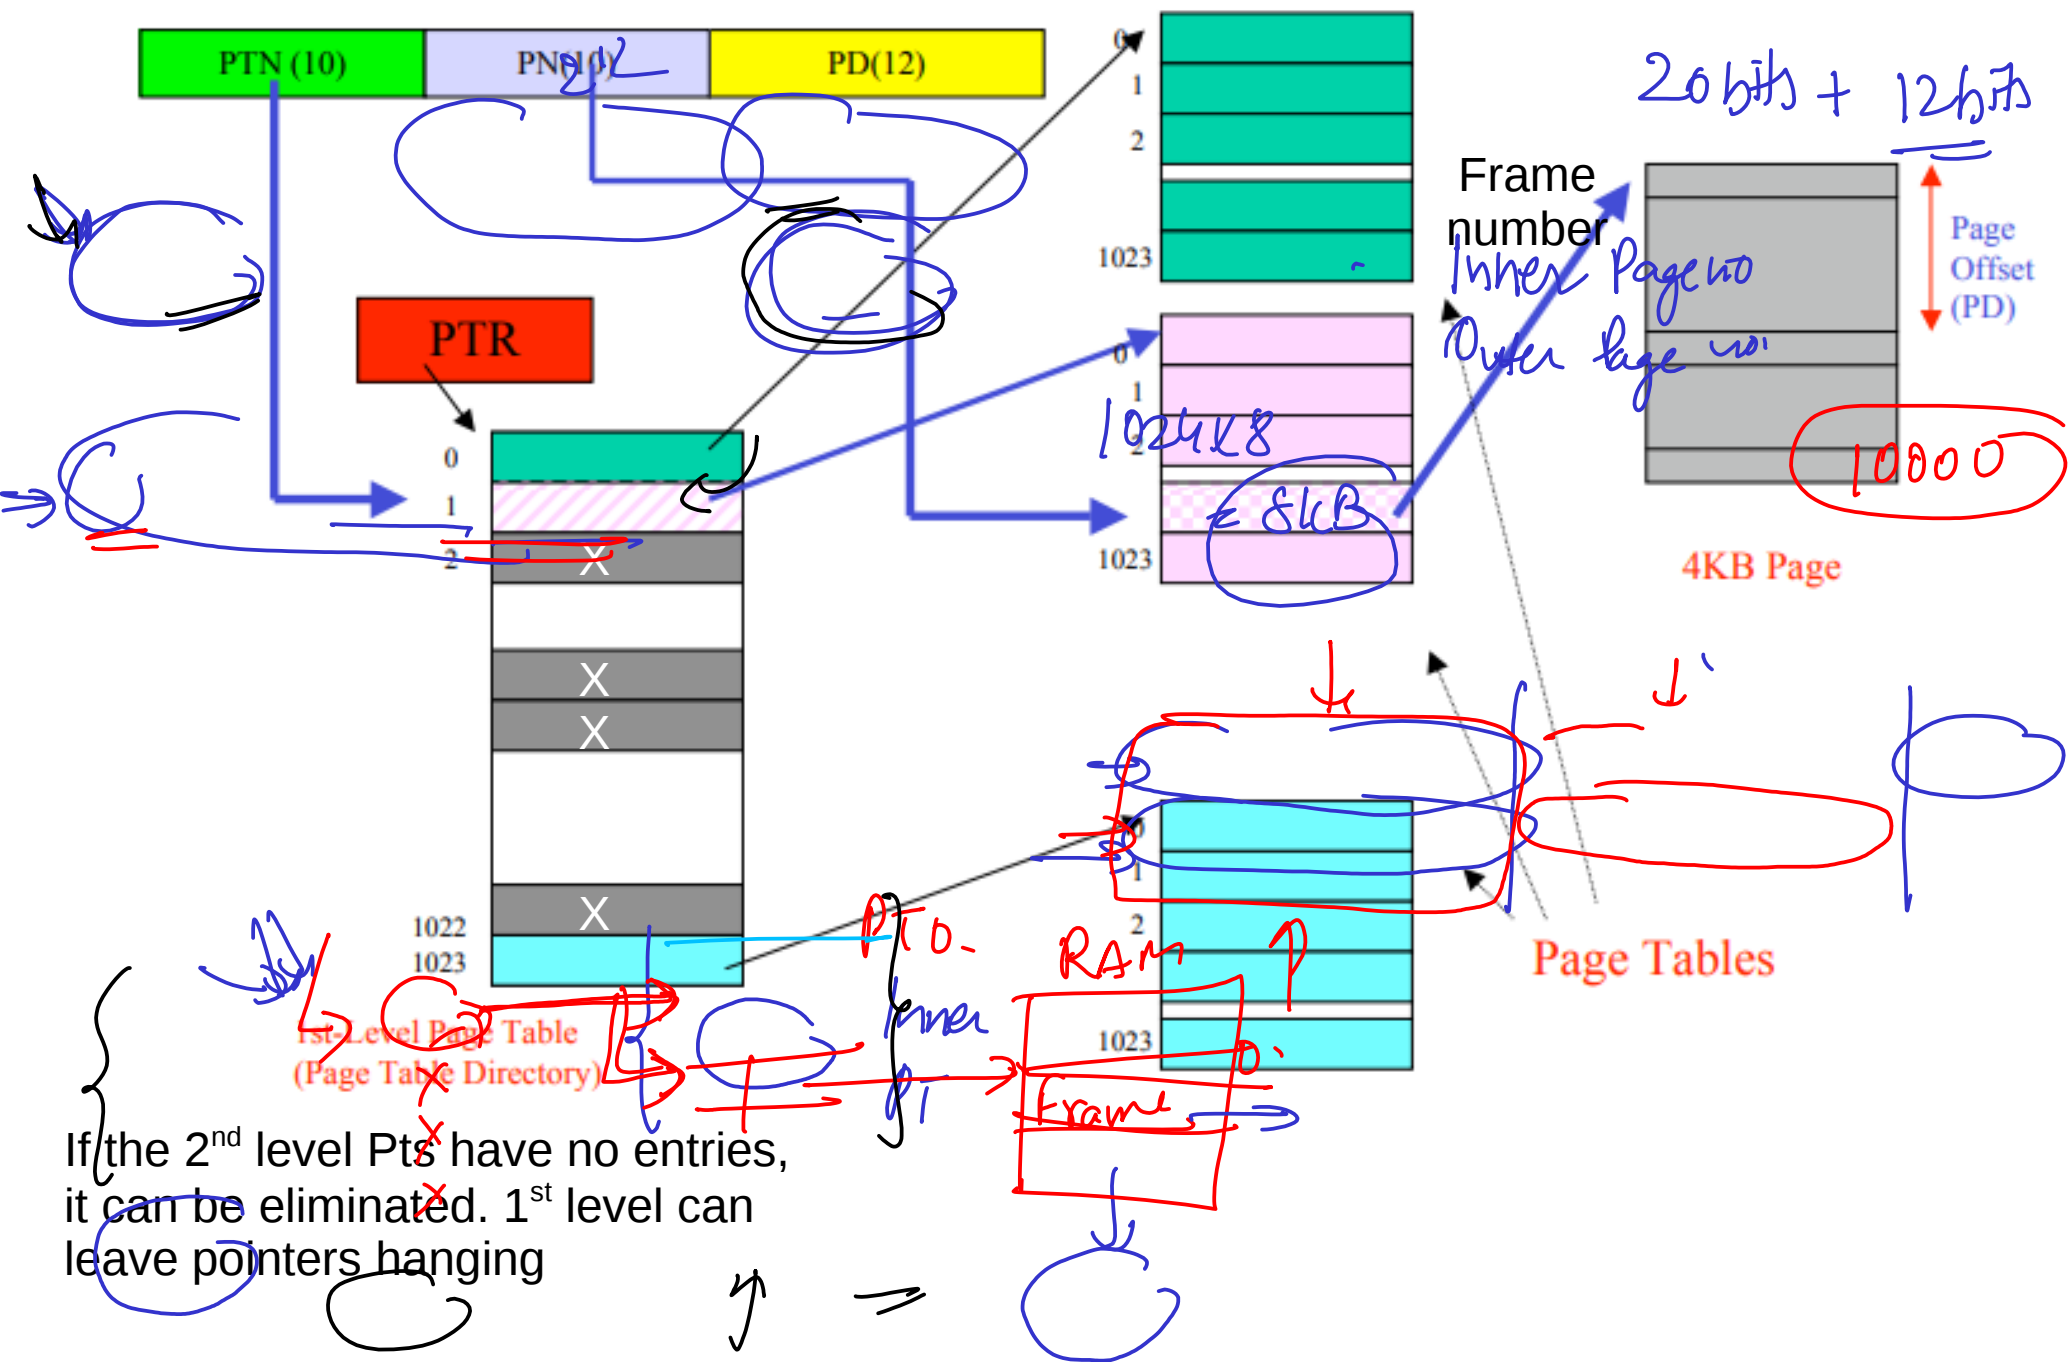
0
X
X
X

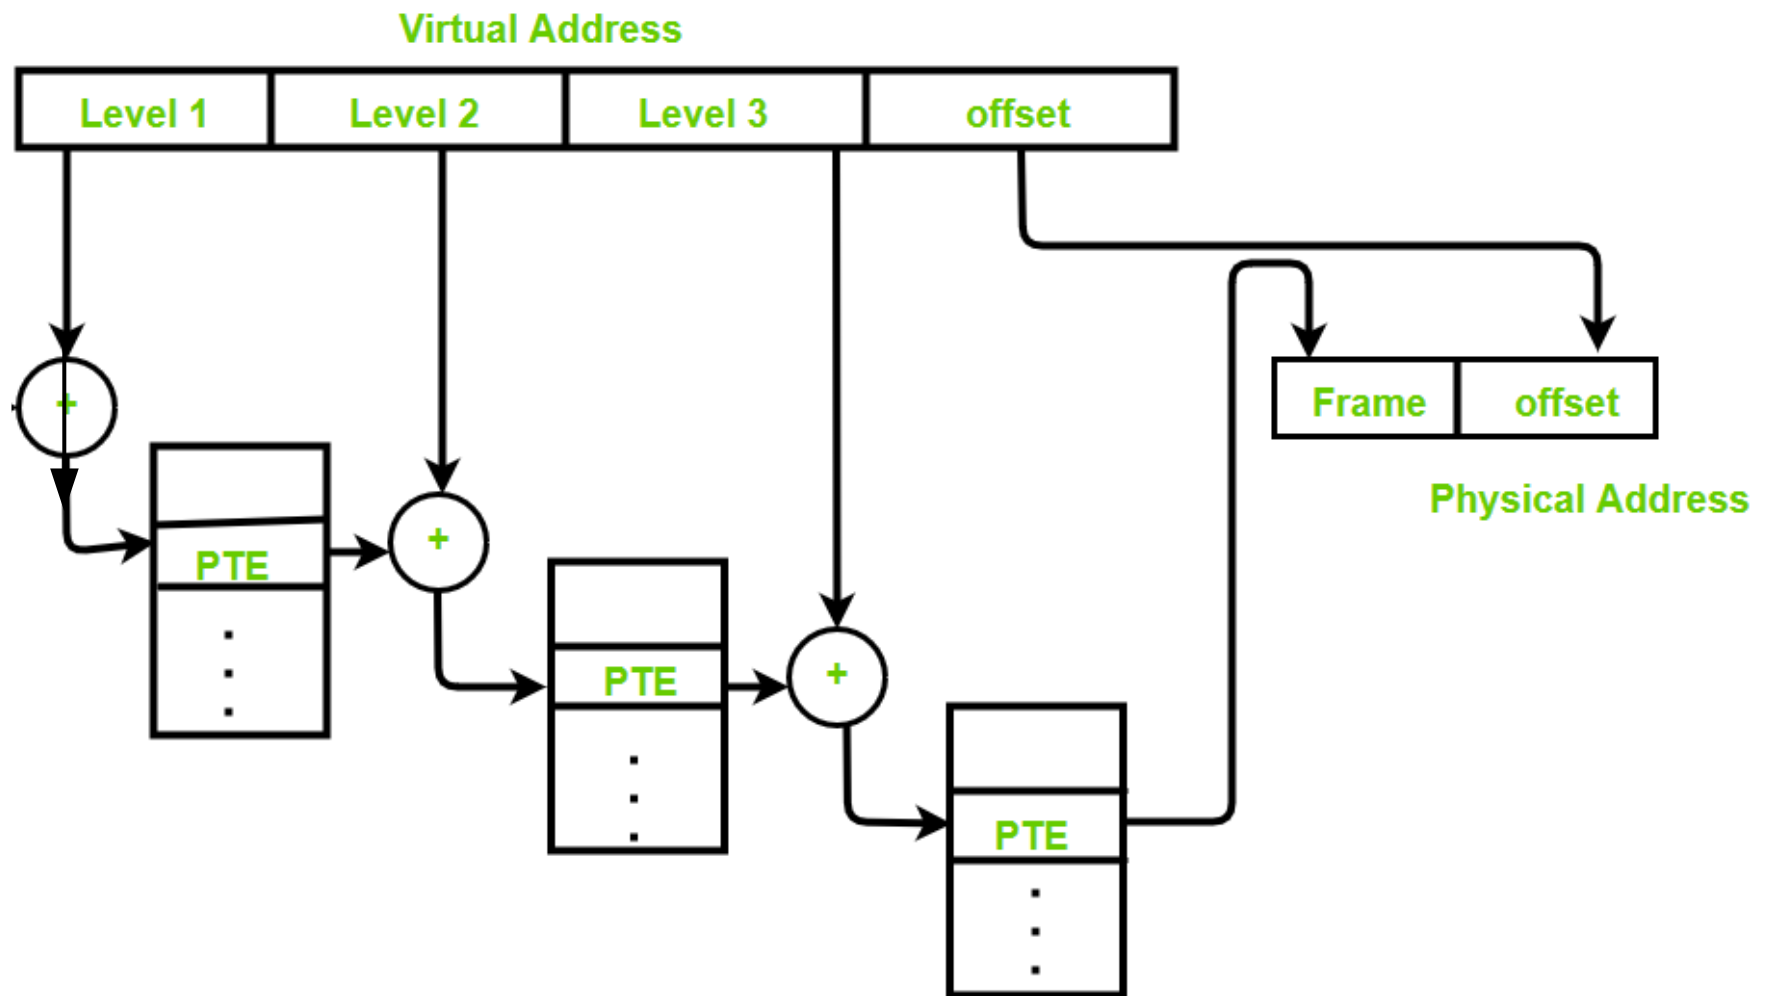
L2 1024

$2^{10} * 8 \text{ bytes} = 8\text{kB}$

$2^{10} * 8 \text{ bytes per entry} = 8\text{kB}$

16kB total size- huge memory saving compared to flat page table





3 Level paging system

