

Virtual memory – II

TLB and Caches

Acknowledgements:

CO and design- Chapter 5 - Hennessey and Patterson

CA- Quant approach- H & P

HPCA- Georgia Tech- Virtual memory

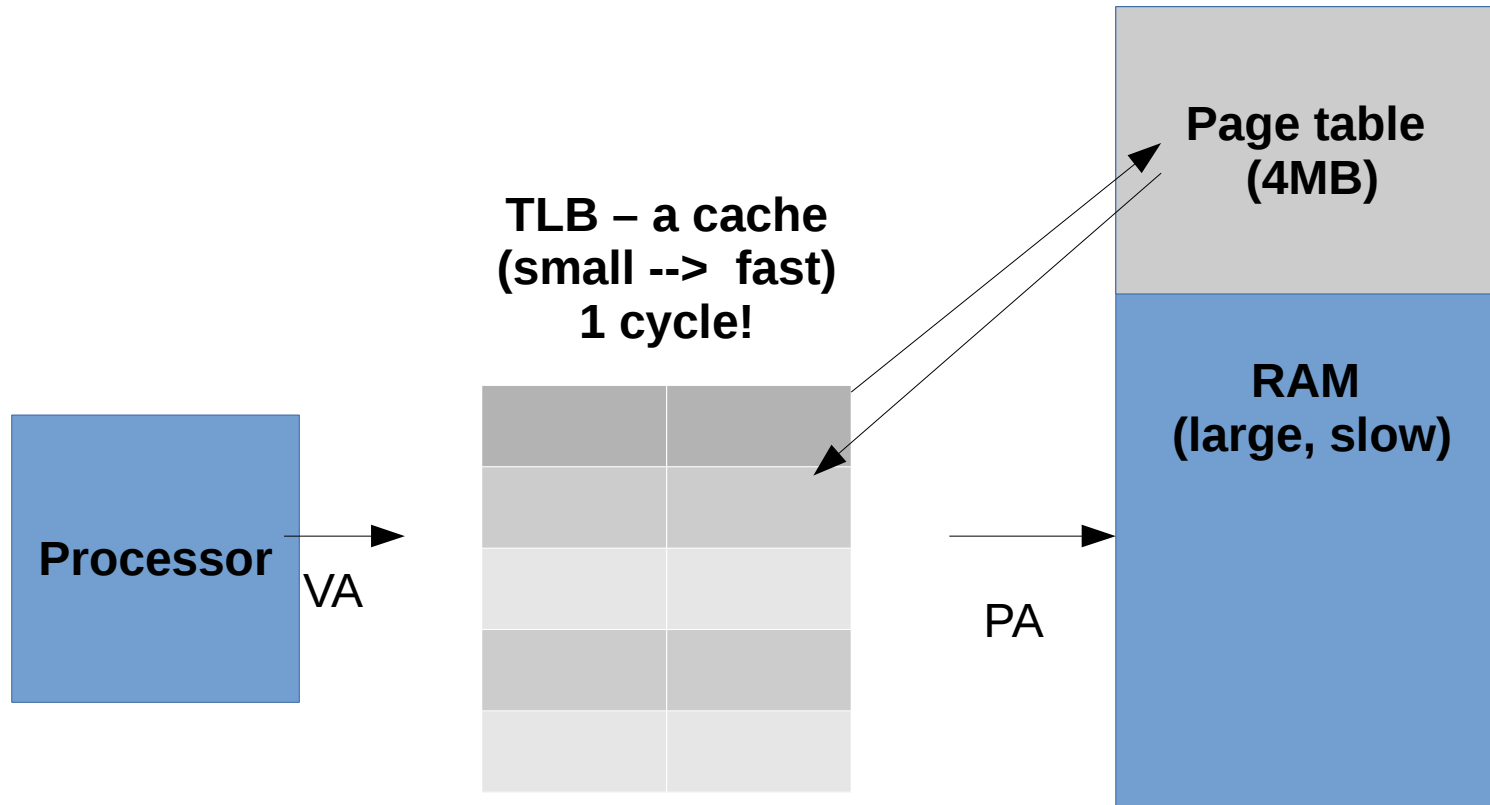
[https://www.youtube.com/watch?](https://www.youtube.com/watch?v=Dz9Hgq65iJw&list=PLAwxTw4SYaPn79fsplluZG34KwbkYSedj)

[v=Dz9Hgq65iJw&list=PLAwxTw4SYaPn79fsplluZG34KwbkYSedj](https://www.youtube.com/watch?v=Dz9Hgq65iJw&list=PLAwxTw4SYaPn79fsplluZG34KwbkYSedj)

Make VM fast

- For each memory access:
 - Access Page table in **RAM**
 - Translate
 - Access data in **RAM**
 - **2 memory accesses**
- **PT Look up fast?**

Translation look aside buffer (TLB)



- **TLB – Cache of recent virtual to physical address translations.**
- **Stores the final frame number (even in case of a multi-level PT)**
- iTLB and dTLB
- Small size, fully or highly associative
- Eg - TLB: 64 entries (for 4kB pages) or 512 entries
- Can have a 2- level TLB

Size

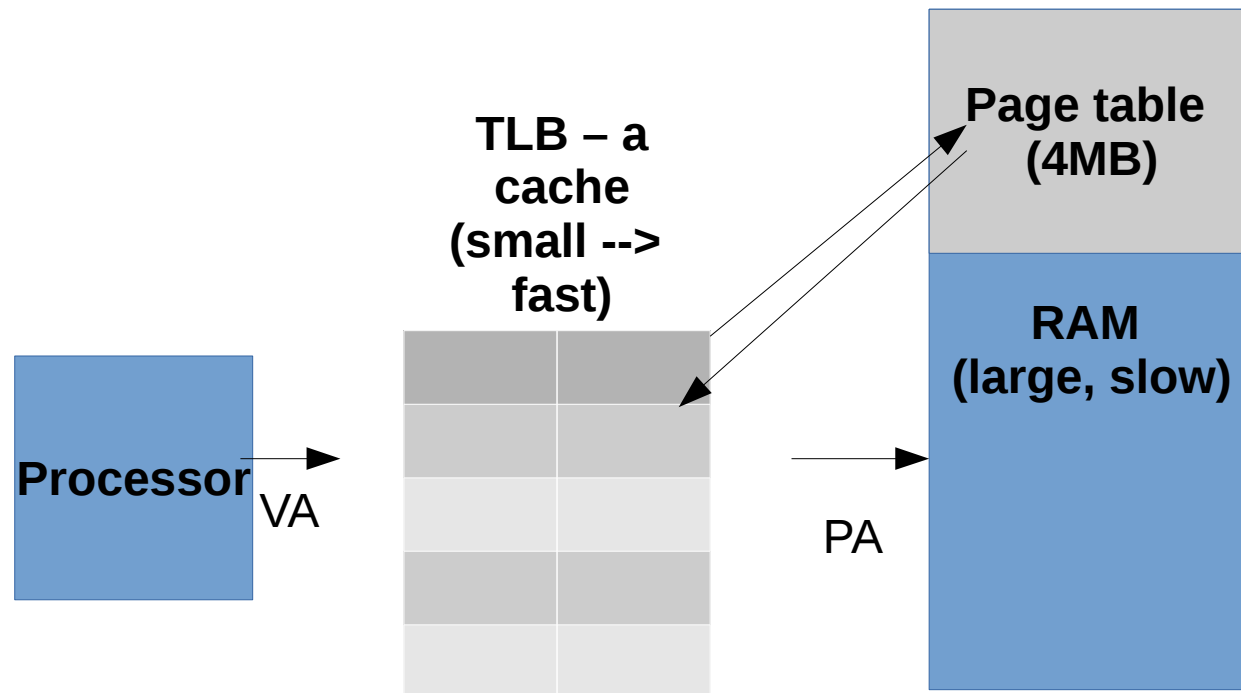
For a 32 bit program, there are 1million entries in Page Table

Assume page size of 4kB and TLB has 64 entries

--> Each TLB entry covers 4kB pages

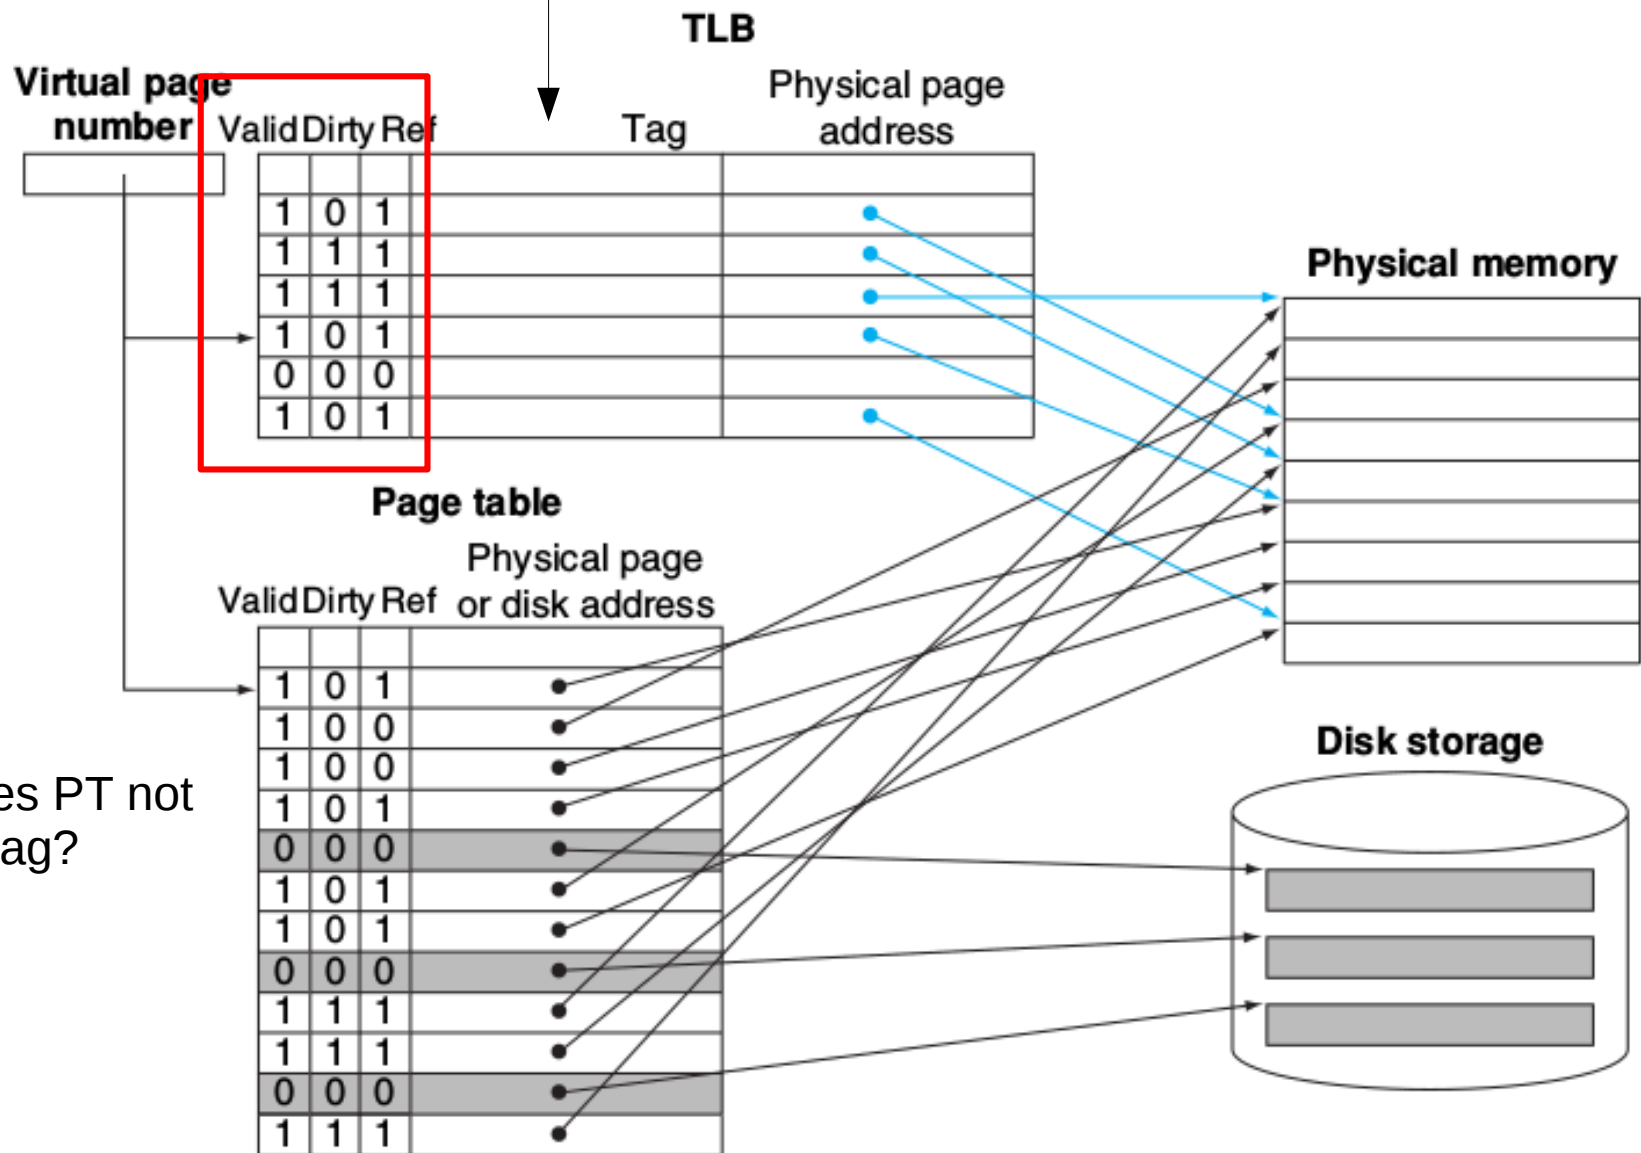

64 4kB pages --> 256kB

Translation look aside buffer (TLB)

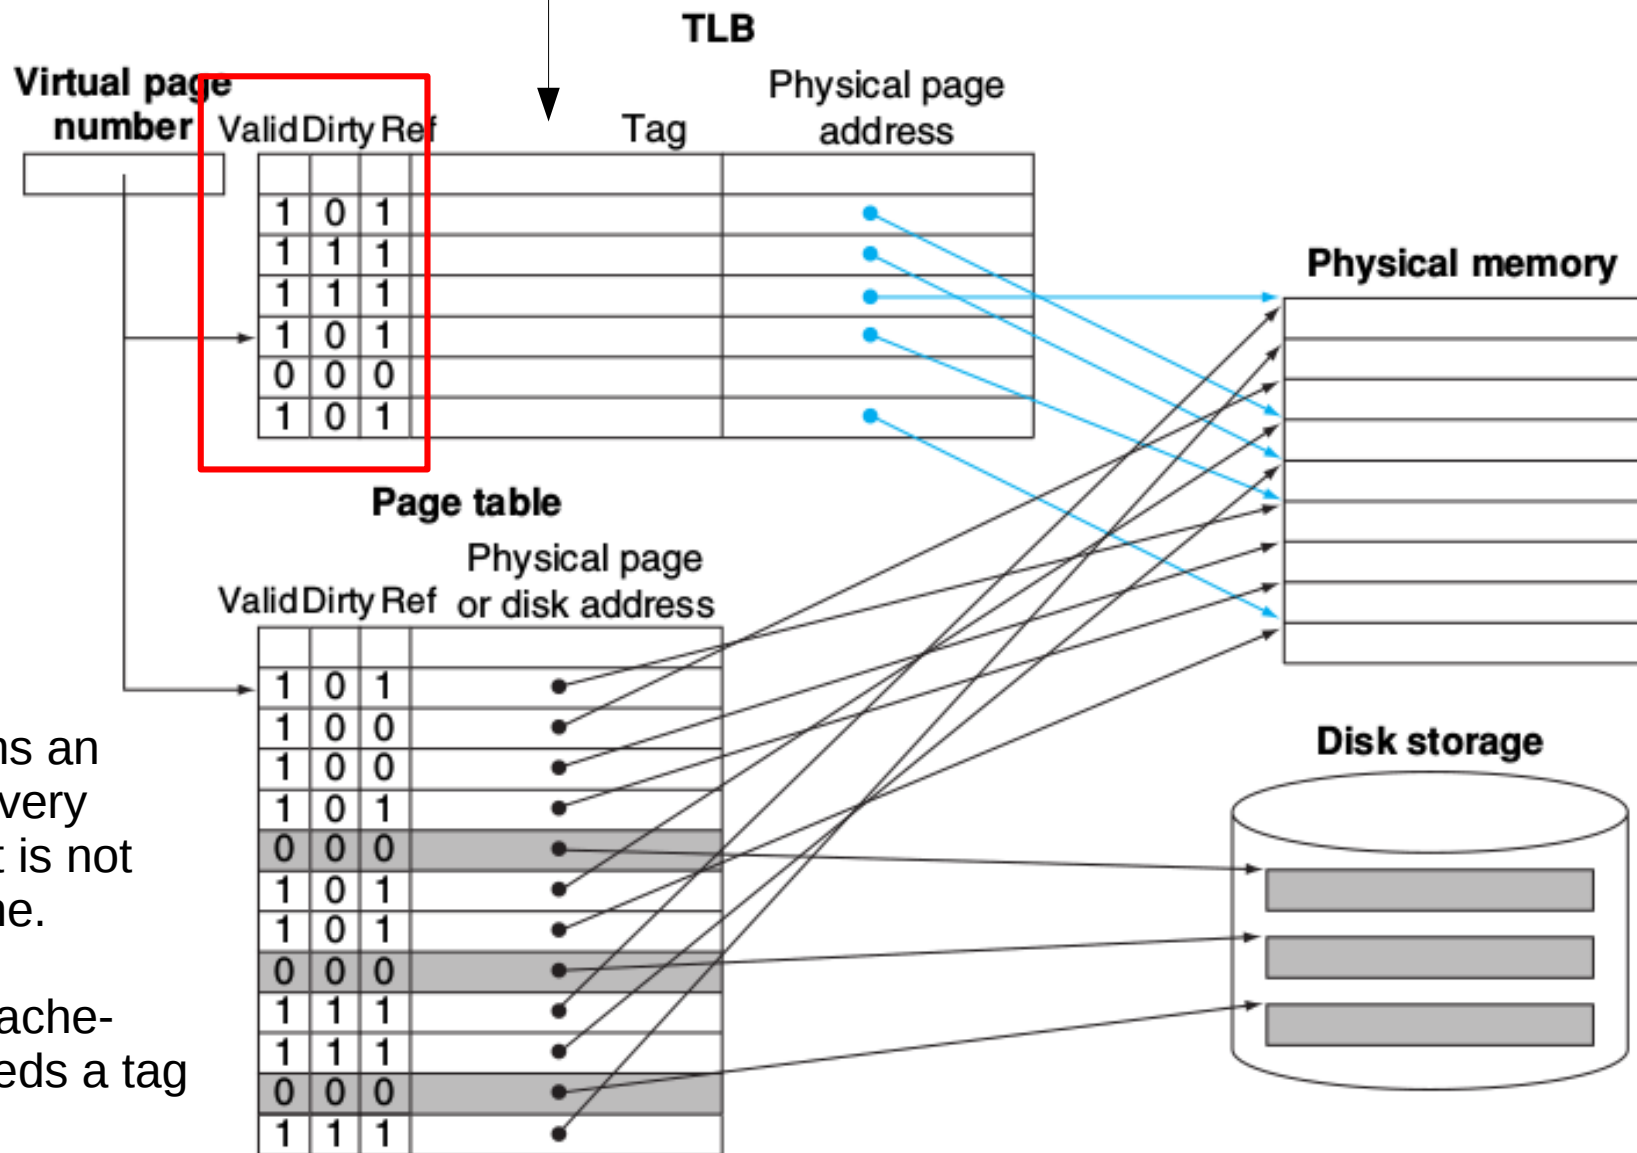



- Entry not in TLB -->
 - Go to PT and fetch the translation, go to RAM, fetch the page/data - SLOW
 - Go to PT and fetch the translation --> Page not in RAM, Go to disk and fetch, update PTE --> WORST
- Entry in TLB -->
 - Go to RAM, fetch the page/data – BEST
 - Page not in RAM? Unlikely, TLB entry will get evicted soon

How does the TLB work



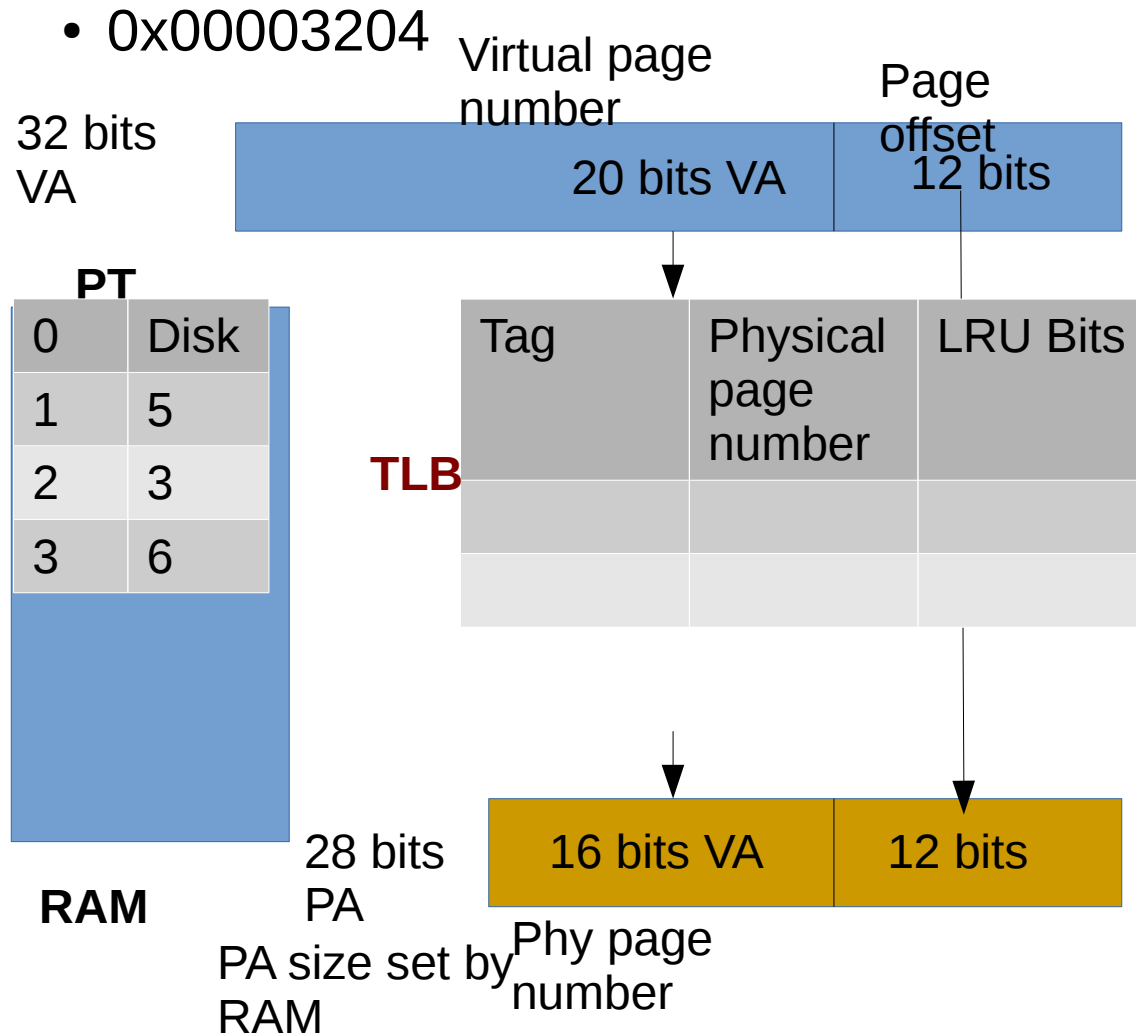
Why does PT not need a tag?



PT contains an entry for every page- so it is not like a cache.

TLB is a cache- hence, needs a tag

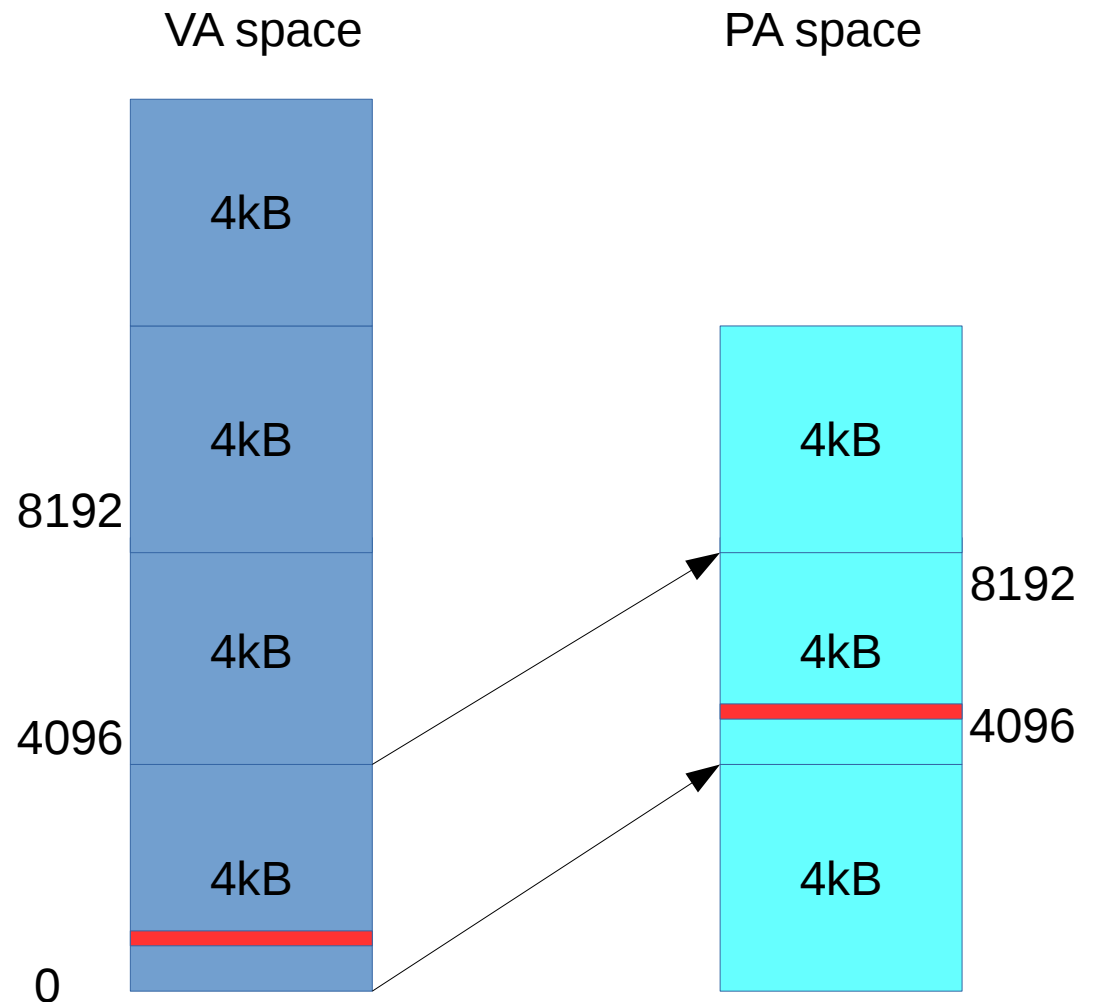
Translation using TLB- assume Fully associative TLB (no index)



- Note that the translation is for a page, not for individual bytes or words
- 0x00003204 --> TLB Miss --> Exception --> OS handles the miss routine
- **Look up the VPN as an index into the PT in RAM**
- 3 --> 6
- OS: System instructions Update the TLB
 - Tag=3
 - Phy page number=6
- Next time --> TLB hit
- Next time if you get 0x00003206 --> Page number is 3. It will be a hit.

What does the TLB store

- Note that the translation is for a page, not for individual bytes or words
- Say, Virtual address = 0: TLB Miss
- Go to the PT, fetch the translation, update the TLB
- Next, if we get a virtual address anywhere in that 4kB page (say, 2, 10, 4095), the translation will be a hit in TLB, as they share the same page number
- Remember TLB does not store data, it only stores translations
- Still need to go to cache or RAM to access data




Translation using TLB

VA size set by ISA

32 bits VA

Virtual page
number

Page offset



• 0x00001204

PT

0	Disk
1	5
2	3
3	6

Tag	Physical page number	LRU Bits
3	6	1
		0

TLB

28 bits

RAM

PA

PA size set by
RAM



Phy page number

TLB

VA size set by ISA

32 bits VA

Virtual page
number

Page offset

20 bits VA

12 bits

PT

0	Disk
1	5
2	2
3	6

Tag	Physical page number	LRU Bits
3	6	2 (LRU)
1	5	1 (MRU)

TLB

16 bits VA

12 bits

Phy page number

RAM

28 bits
PA

PA size set by
RAM

- TLB Miss
- 0x00001204 -->
Update LRU Bits if
needed

TLB Miss

VA size set by ISA

32 bits VA

Virtual page
number

Page offset

20 bits VA

12 bits

PT

0	Disk
1	5
2	2
3	6

Tag	Physical page number	LRU Bits
3	6	2
1	5	1

TLB

16 bits VA

12 bits

Phy page number

RAM

28 bits

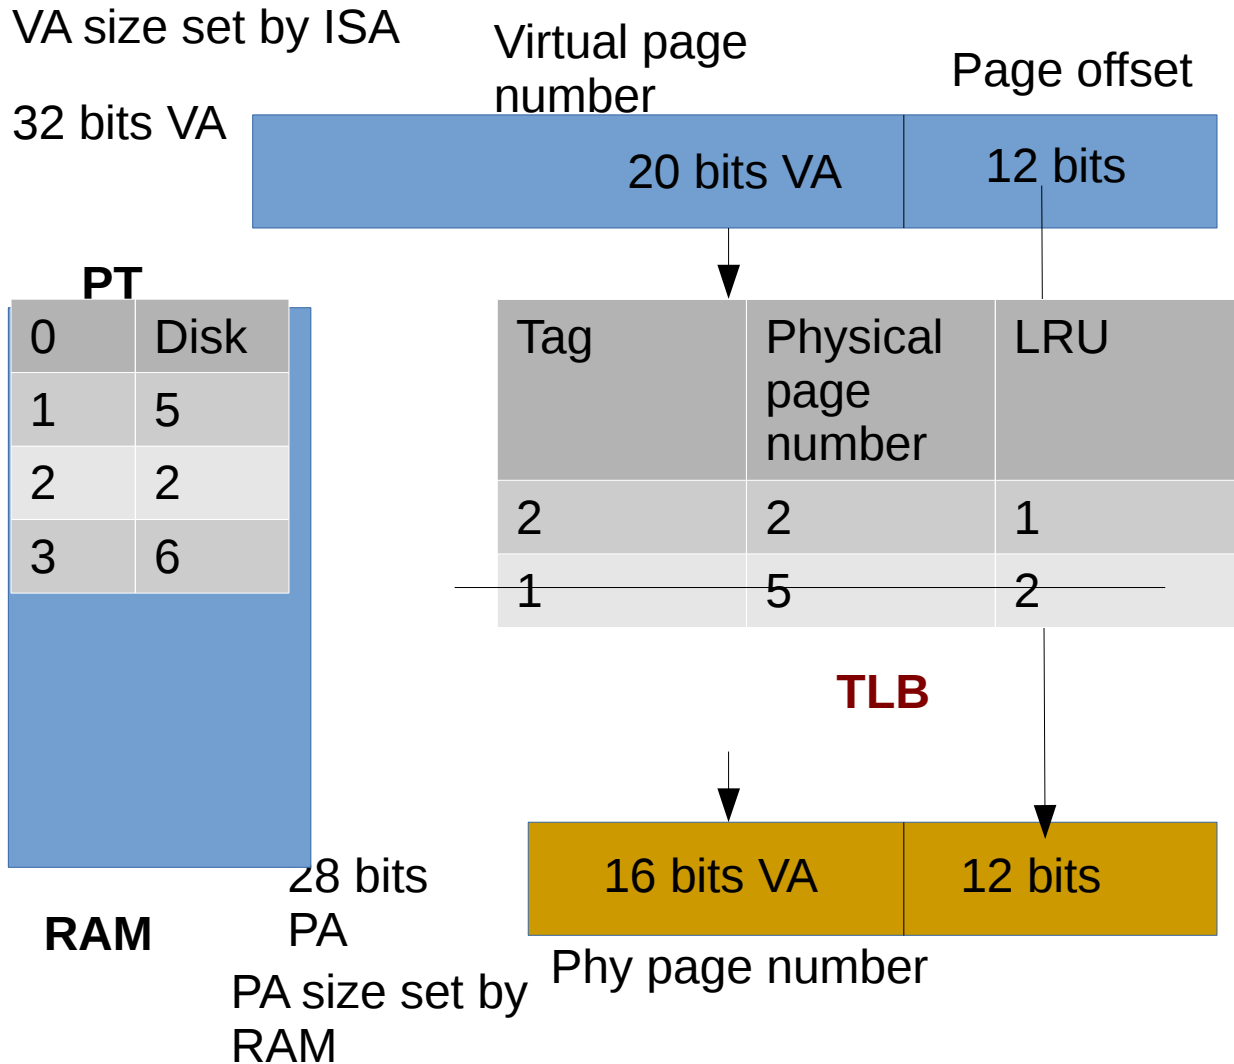
PA

PA size set by
RAM

- 0x00002204 --> TLB miss
- No place in TLB --> LRU
- VP 2 --> PP 2
- Update LRU Bits

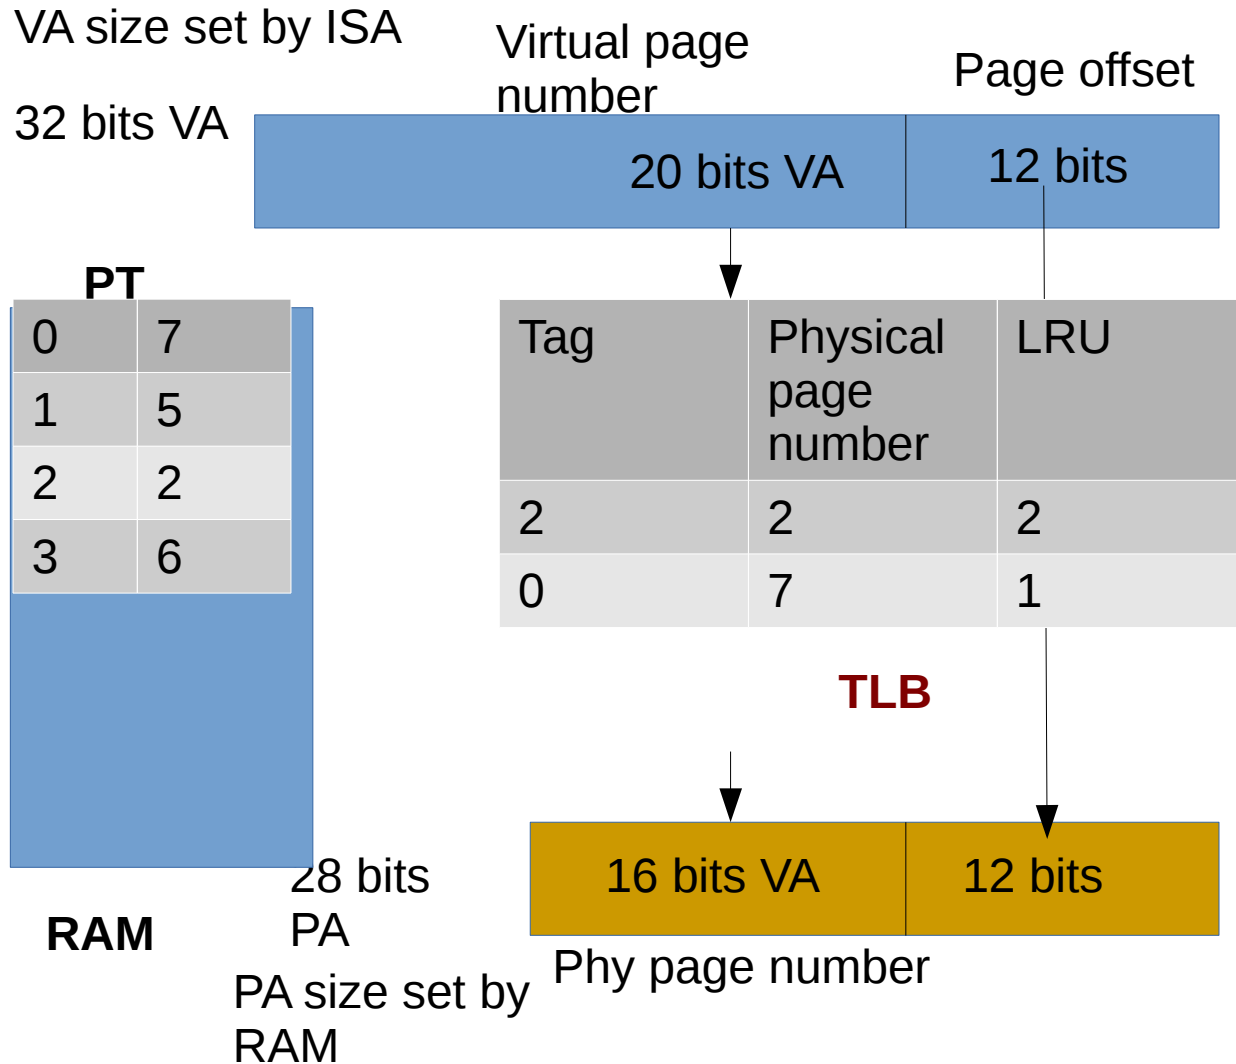
Tag	Physical page number
2	2
1	5

TLB Miss



- 0x00000204 --> TLB miss
- No place in TLB --> LRU
- Evict Tag 1
- **Translation --> Page fault**
- Load to memory from Disk, assign a new Pg number- whichever is free update PT, update TLB

TLB and PT status in the end



- No space in memory?
- Replace a page
- Check dirty bit of the page. Write back to disk before replacing
- Update page table, Update TLB entry

Handling TLB misses and page faults

- Who handles TLB misses?
 - Miss exception -- OS – Software TLB Miss handling
 - Embedded processors
 - Processor? - Hardware TLB miss handling (similar to cache miss handling)- Faster
 - Eg- x86
- Page fault exception: OS

Reduce TLB misses

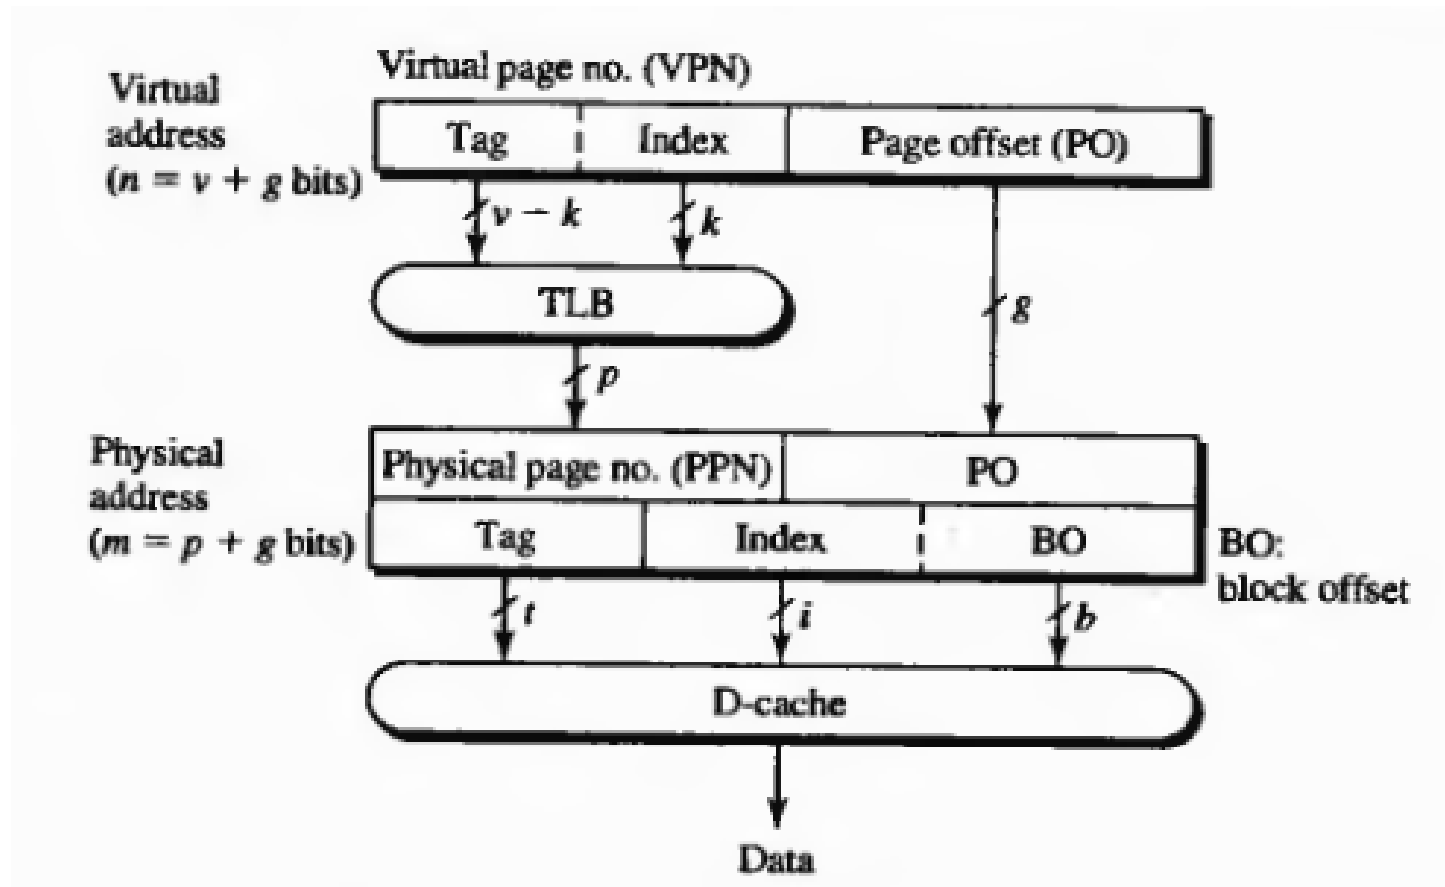
- Make pages larger
 - 64 4kB pages --> 256kB: Translate addresses to access 256kB of data
 - 32 2MB pages --> 64MB data can be covered
- Add a second TLB, like an L2 cache. Larger TLB
 - Eg: AMD Opteron

Memory security

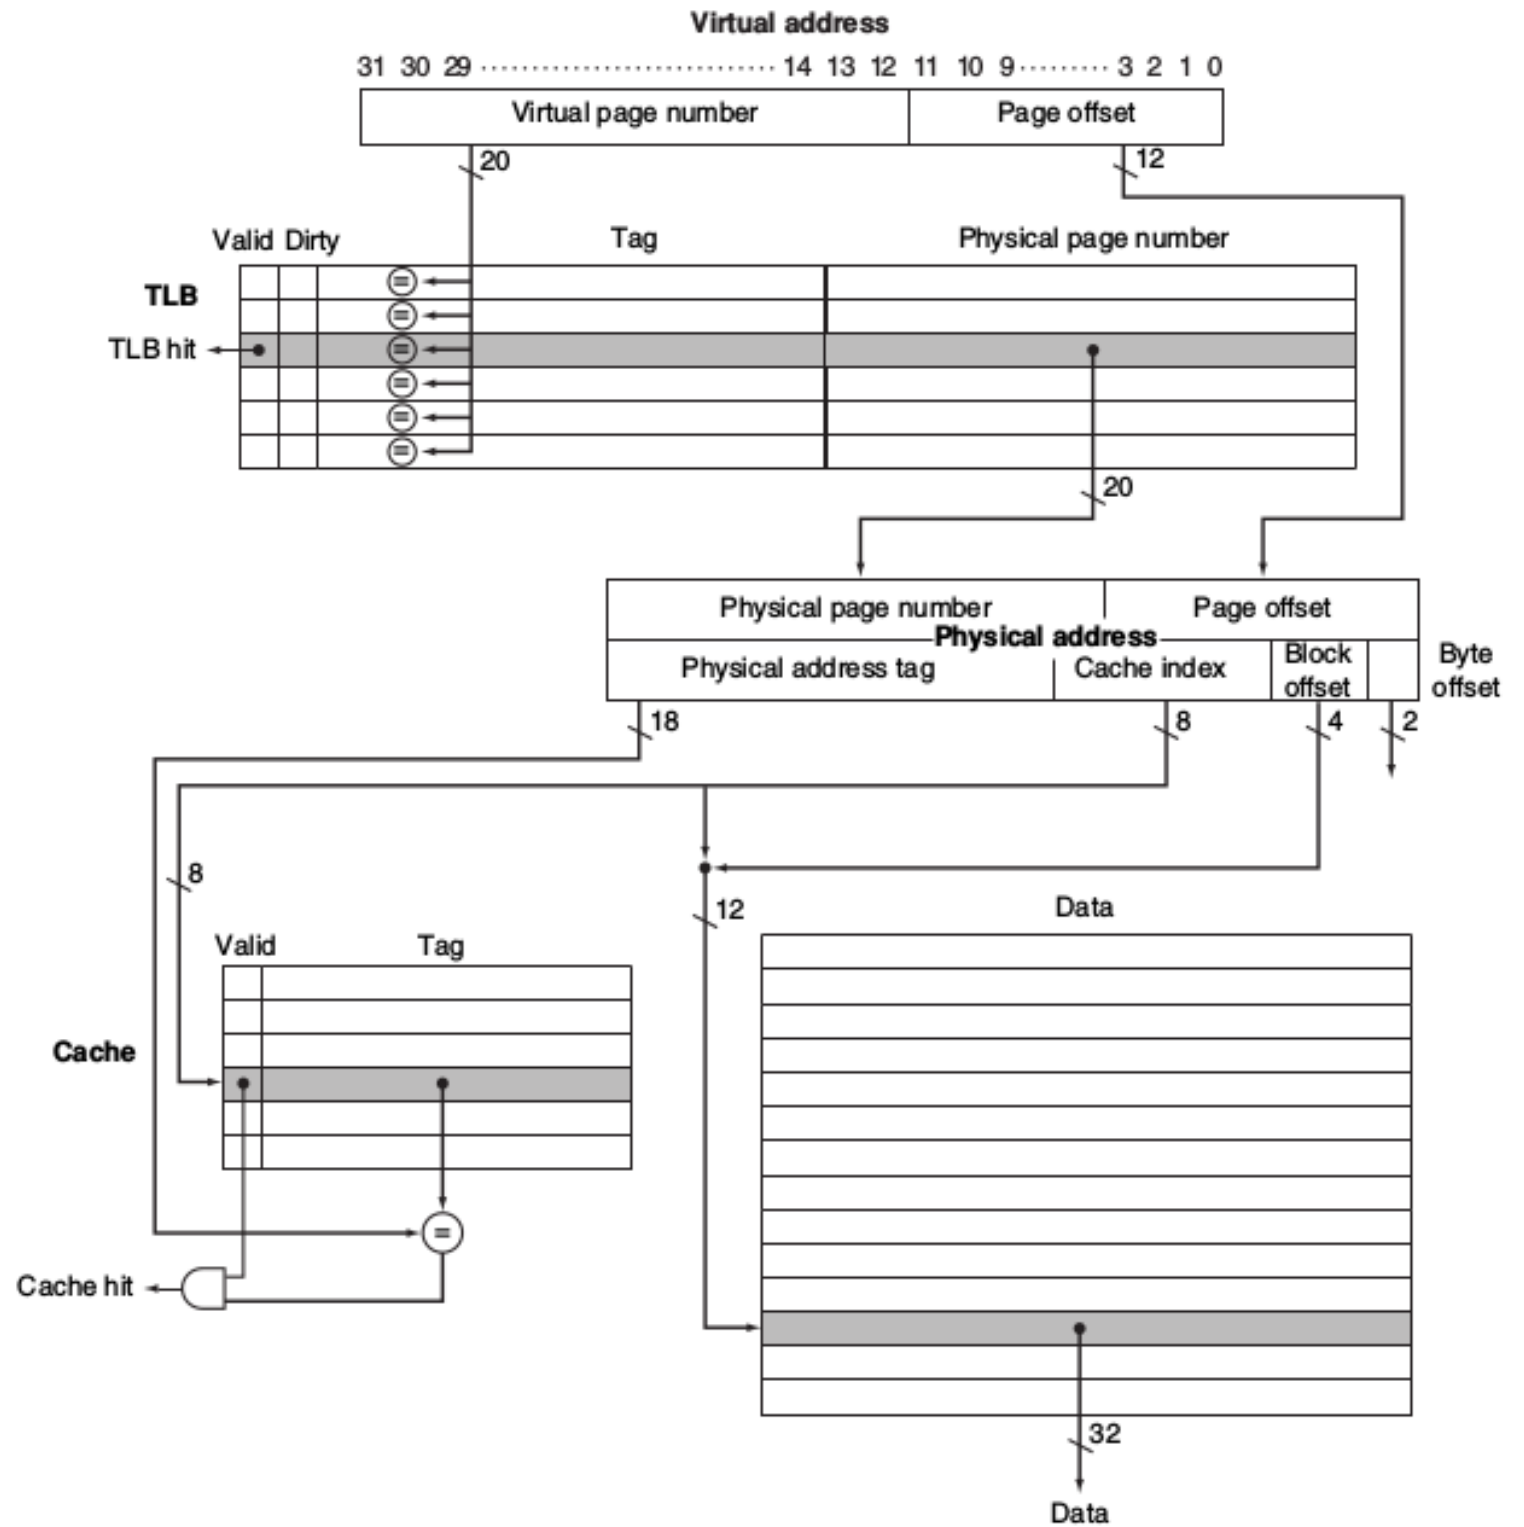
- Although multiple processes are sharing the same main memory, one process cannot write into the address space of another user process or into the operating system
- TLB: Write access bit: That provides this security

TLB and Caches

Physically indexed Physically tagged Cache



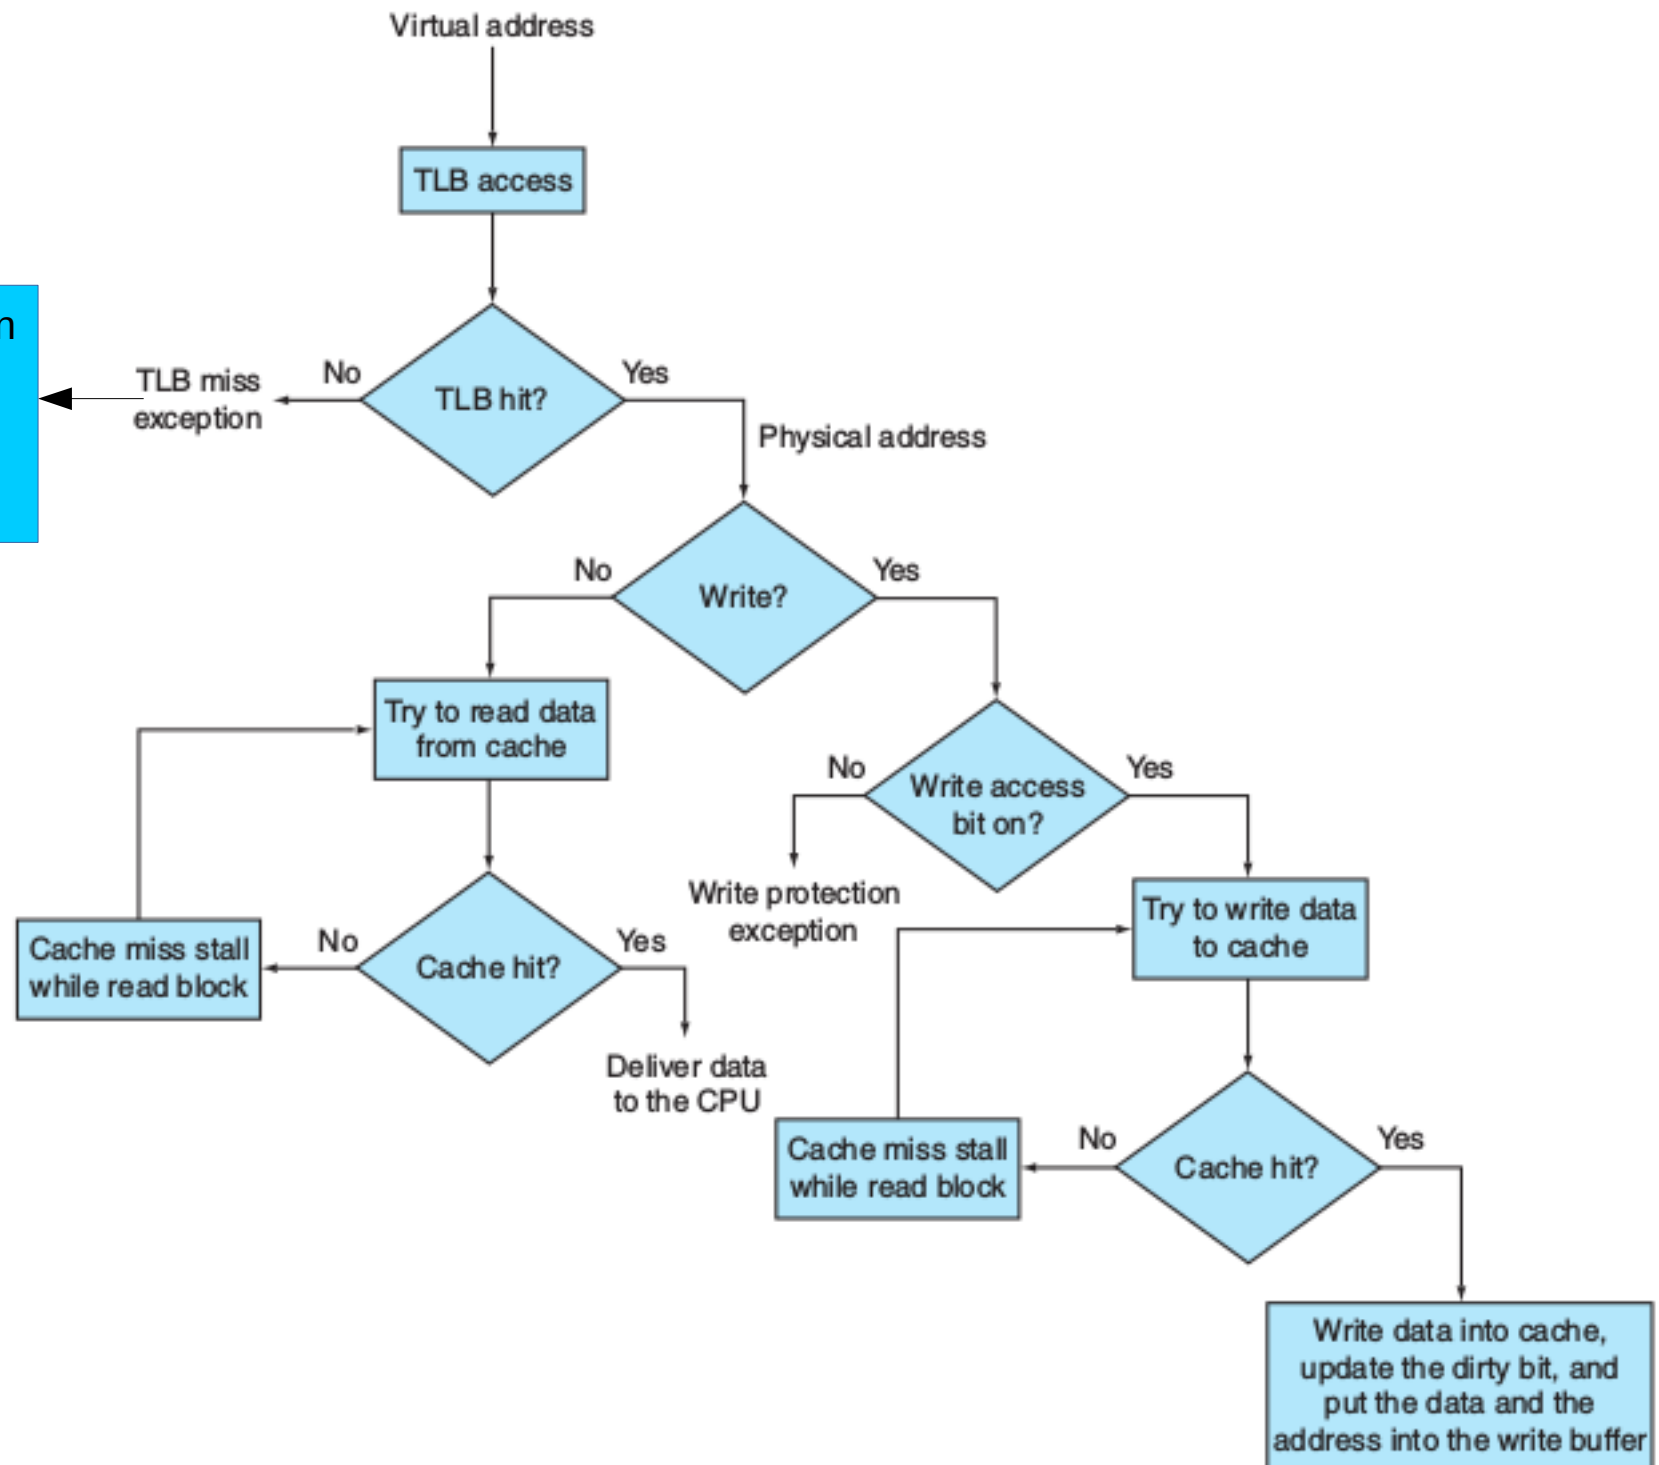
VPN = Tag.
What does it mean?



Limitations of PIPT Cache

- TLB: Say 5 cycles: Has to happen before cache access
- Cache access: 5 cycles: Overall cache access: 10 cycles
- Serialises TLB and Cache

Can result in
Page
Fault and
Cache
miss



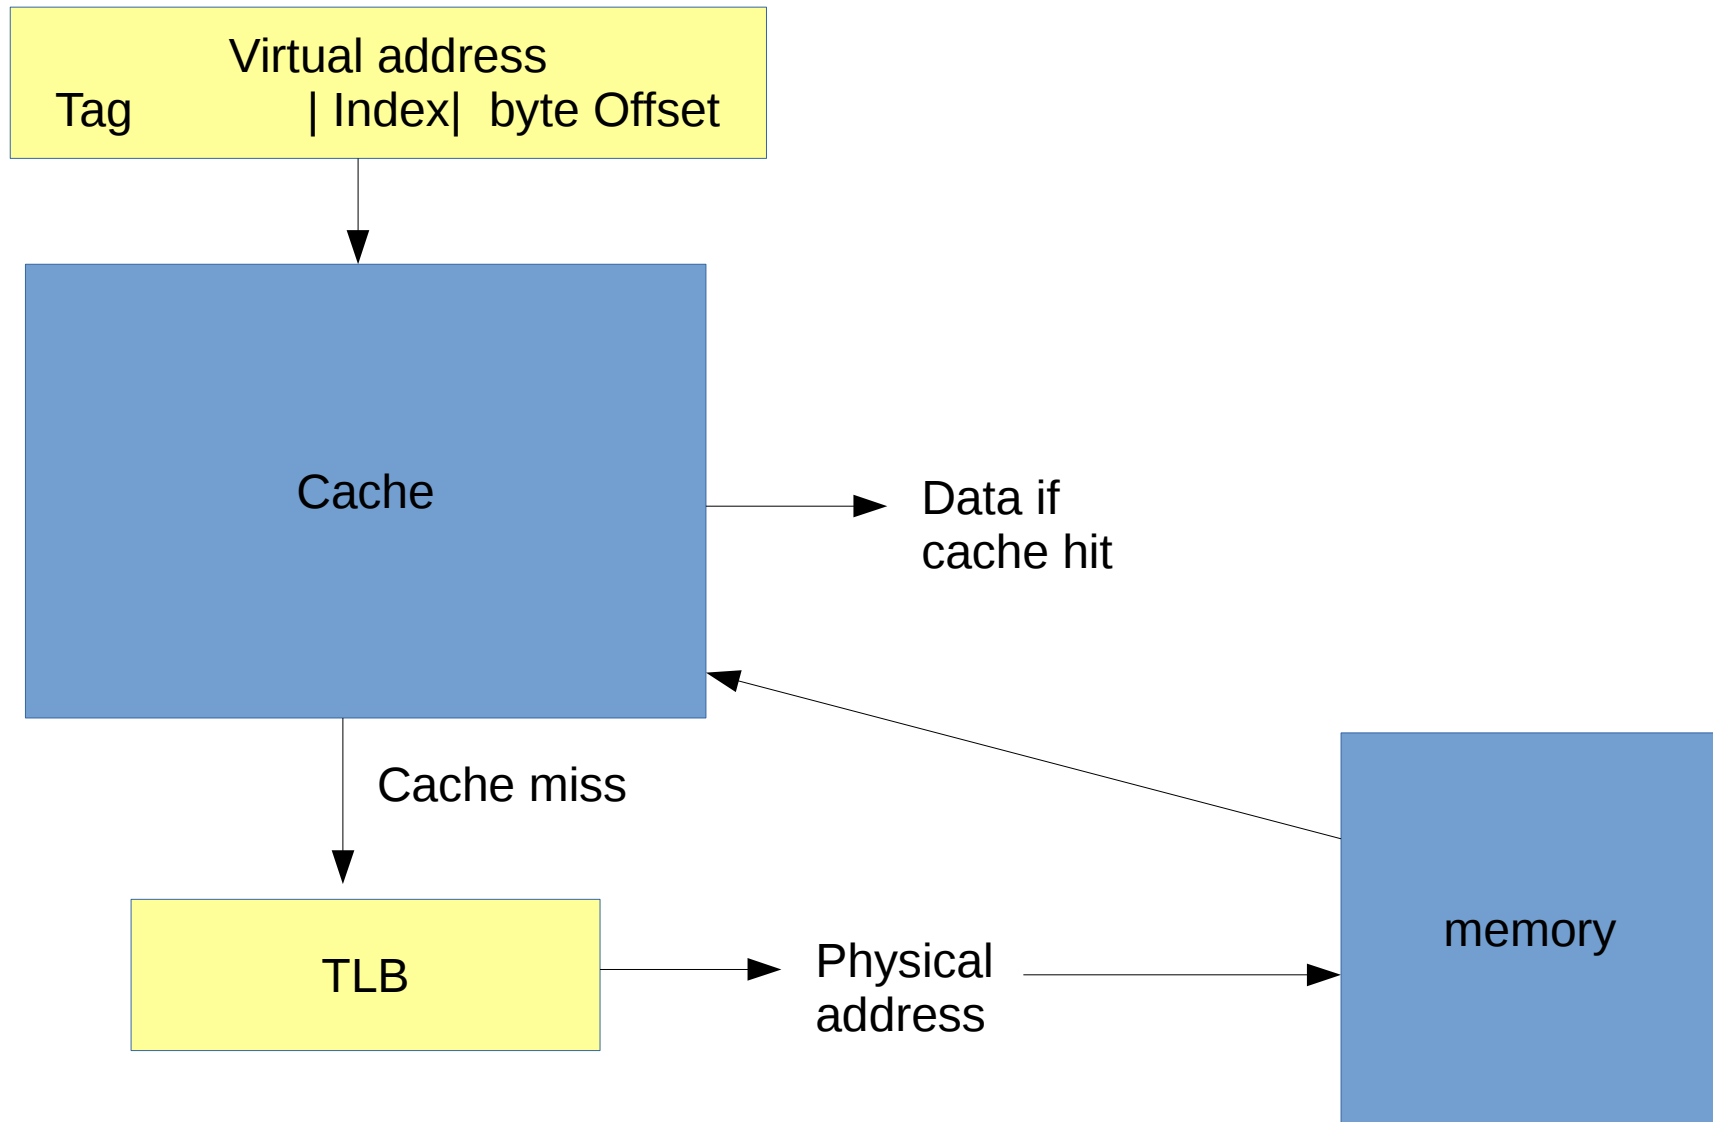
Events in the TLB, virtual memory and PIPT cache

Translation present in TLB,
Page table is 'dont care'

TLB	Page table	Cache	Possible? If so, under what circumstance?
Hit	Hit	Miss	Possible, although the page table is never really checked if TLB hits.
Miss	Hit	Hit	TLB misses, but entry found in page table; after retry, data is found in cache.
Miss	Hit	Miss	TLB misses, but entry found in page table; after retry, data misses in cache.
Miss	Miss	Miss	TLB misses and is followed by a page fault; after retry, data must miss in cache.
Hit	Miss	Miss	Impossible: cannot have a translation in TLB if page is not present in memory.
Hit	Miss	Hit	Impossible: cannot have a translation in TLB if page is not present in memory.
Miss	Miss	Hit	Impossible: data cannot be allowed in cache if the page is not in memory.
Hit	x	Hit	Best case scenario, retrieve the data from cache

Page not in memory, cannot
be definitely present in the
cache. Cache must miss

Virtually indexed Virtually tagged Cache




VIVT Cache

- No TLB latency to access cache. TLB not part of the critical path
- Limitation:
 - 2 processes can have the same Virtual address, but map to different physical addresses. TLB translations will be different

RD -- 1 : First check cache: Cache miss: TLB: 256 --> 001:

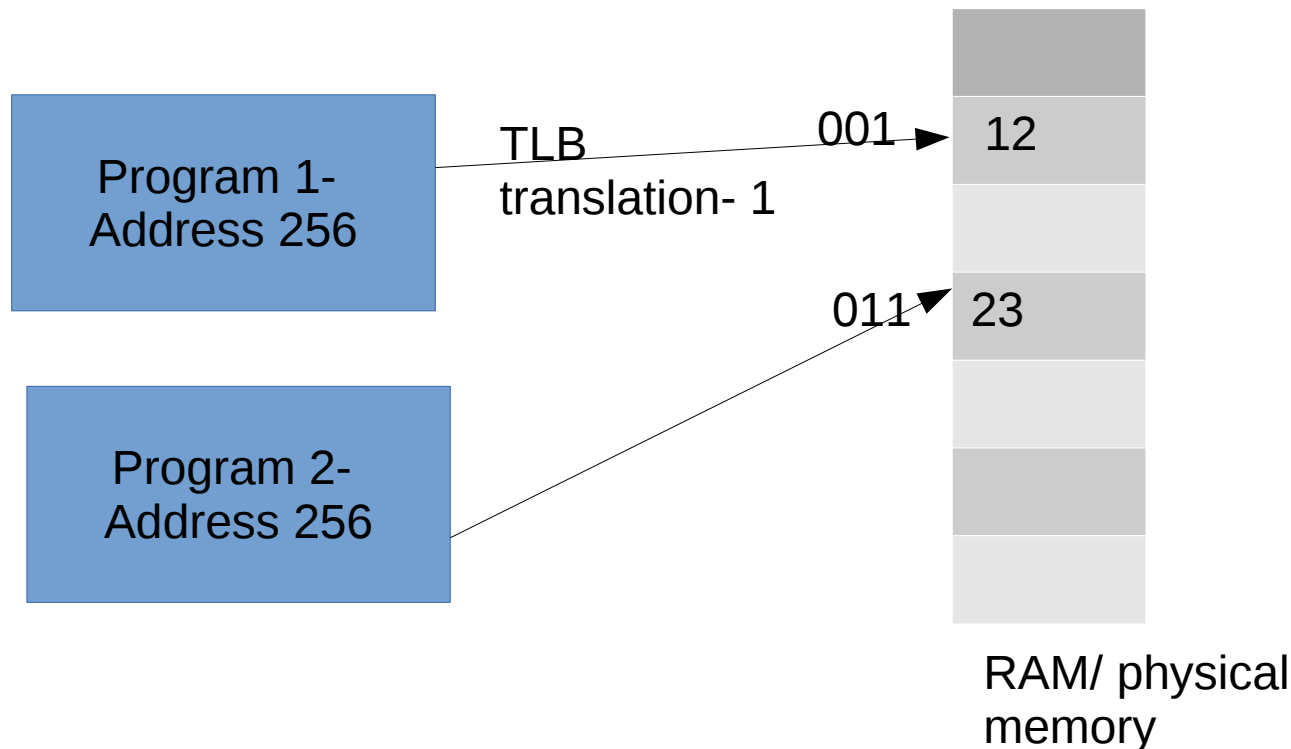
Read 12

RD -- 2: Check cache with VA: Hit: **Reads 12- Incorrect**




VIVT Cache

- When a process/context switch happens, **the entire VIVT cache and TLB will have to be flushed** so that correct data is loaded
 - Can be solved using a **process identifier** along with TLB/Cache entry



Aliasing

- Different VA, but same PA – sharing memory



Aliasing

Process 1: Virtual Address A: 0x12345000

Process 2: Virtual Address B: 0xABCDE000

4kB page: 12 bit page offset

32 bits VA	Virtual page number 20 bits	Page offset 12 bits
	0x 12345 0xABCDE	0x000 0x000

Direct mapped cache: 16 byte block: 4 bits of byte offset

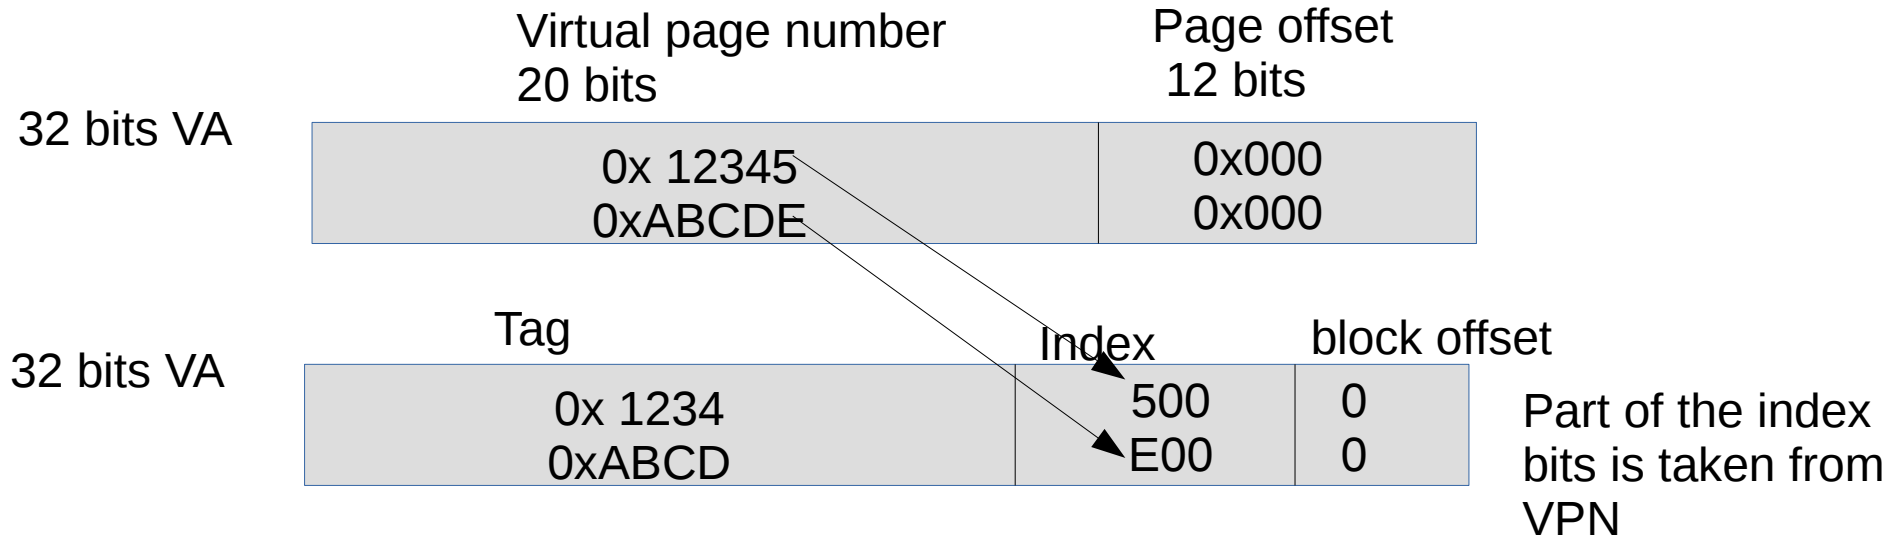
4096 cache lines: 12 bits of index

Index: Virtual index

Tag: Virtual tag

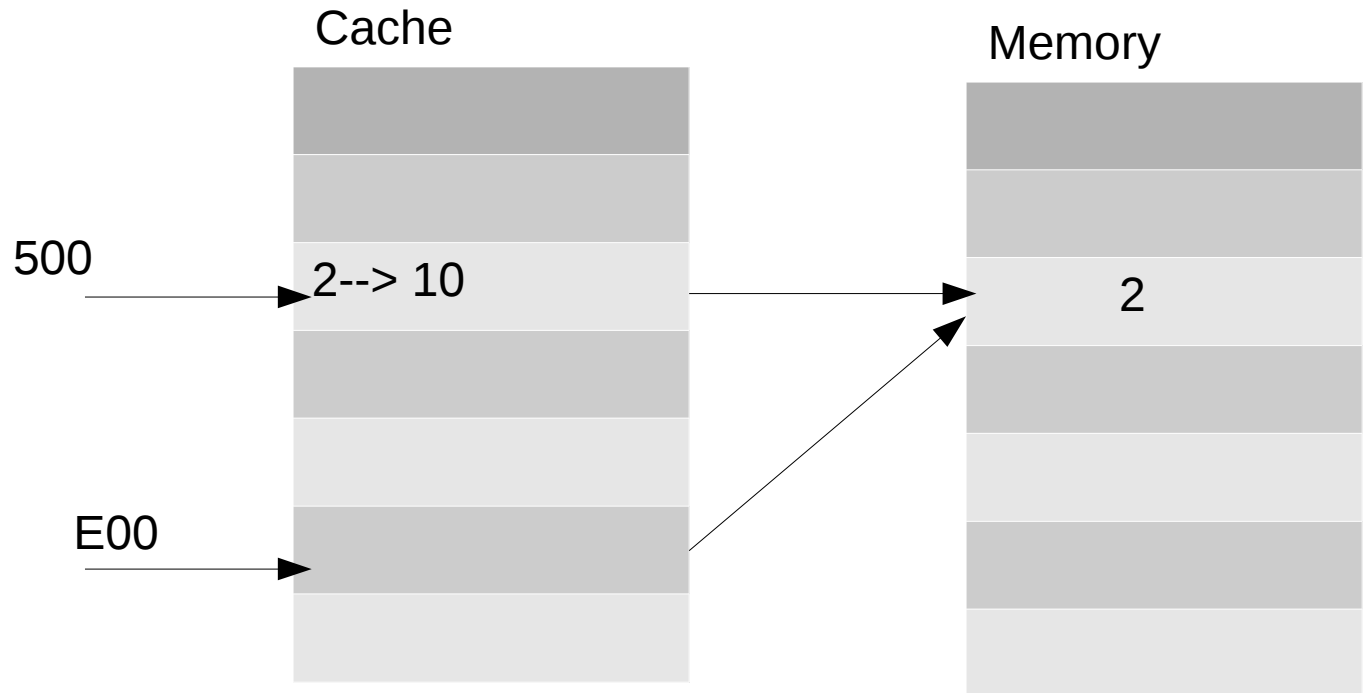
32 bits VA	Tag	Index	block offset
	0x 1234 0xABCD	500 E00	0 0

Aliasing



RD- A –index into 500:
Miss- Read 2
WR- A – Hit: Write 10. Dirty
bit=1
RD-B –Index into E00:
Miss: **Read 2 –
Inconsistent data**

**Problem: Different
indices inspite of
mapping to same
physical memory**




Aliasing solution

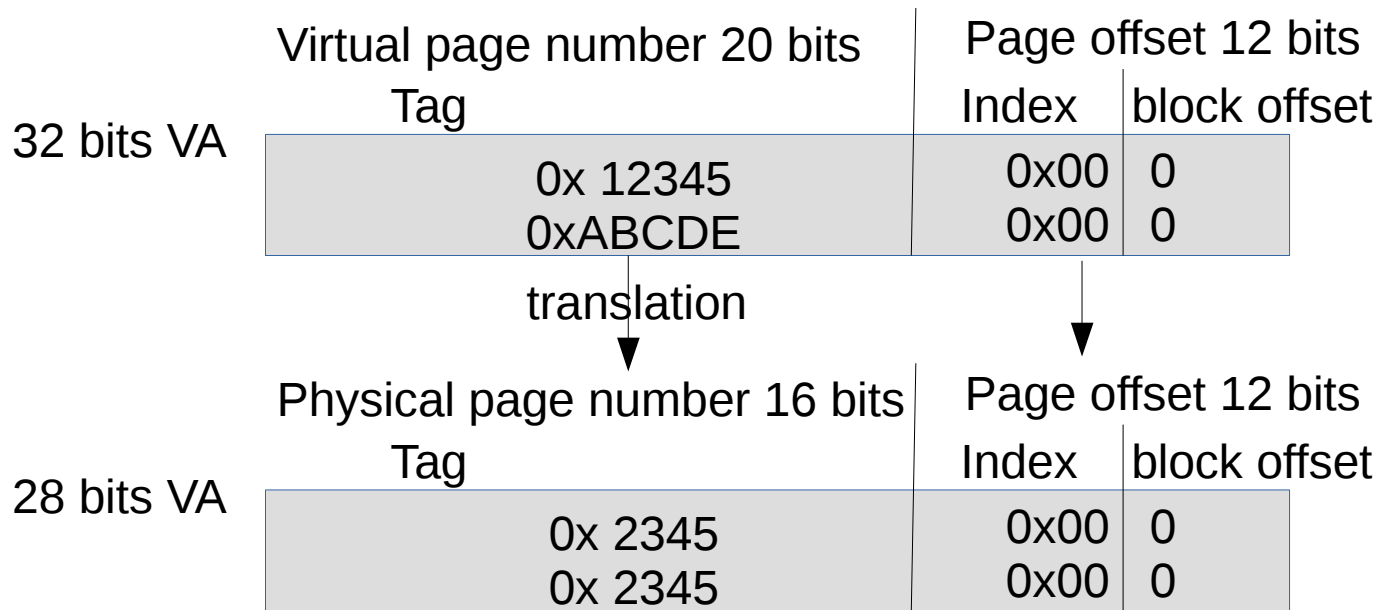
32 bits VA	Virtual page number 20 bits	Page offset 12 bits	
	Tag	Index	block offset
	0x 12345	0x00	0
	0xABCDE	0x00	0

Problem of Mapping to different indices in the cache:

- Need to have same index
- Note that **the 2 virtual addresses mapping to the same physical address will have the same page offset --> Pick the index bits from page offset**
- **What about tag?**



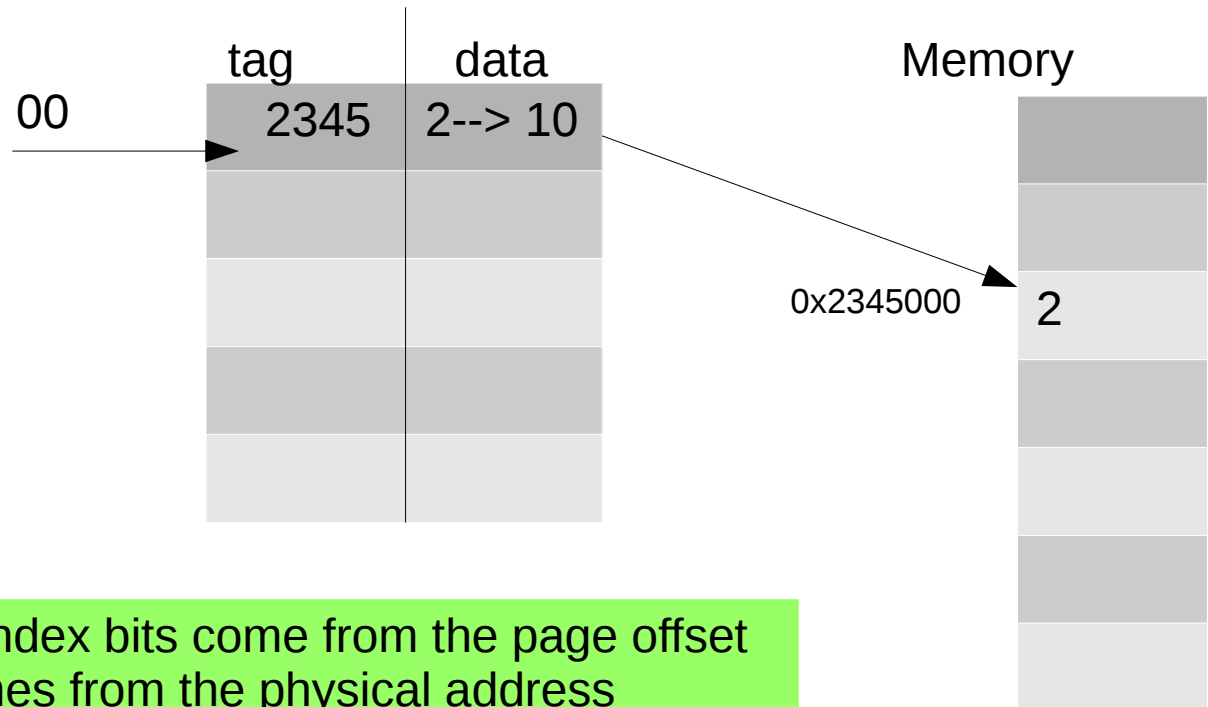
Aliasing solution



RD- A –index into 00: Miss-
Read 2
WR- A – Hit: Write 10. Dirty
bit=1
RD-B –Index into 00: Miss:
Read 10

Tag: Physical tag:
Obtained after address
translation. Both Virtual
addresses will have the
same tag

Also note that the index
bits in virtual and physical
addresses are the same

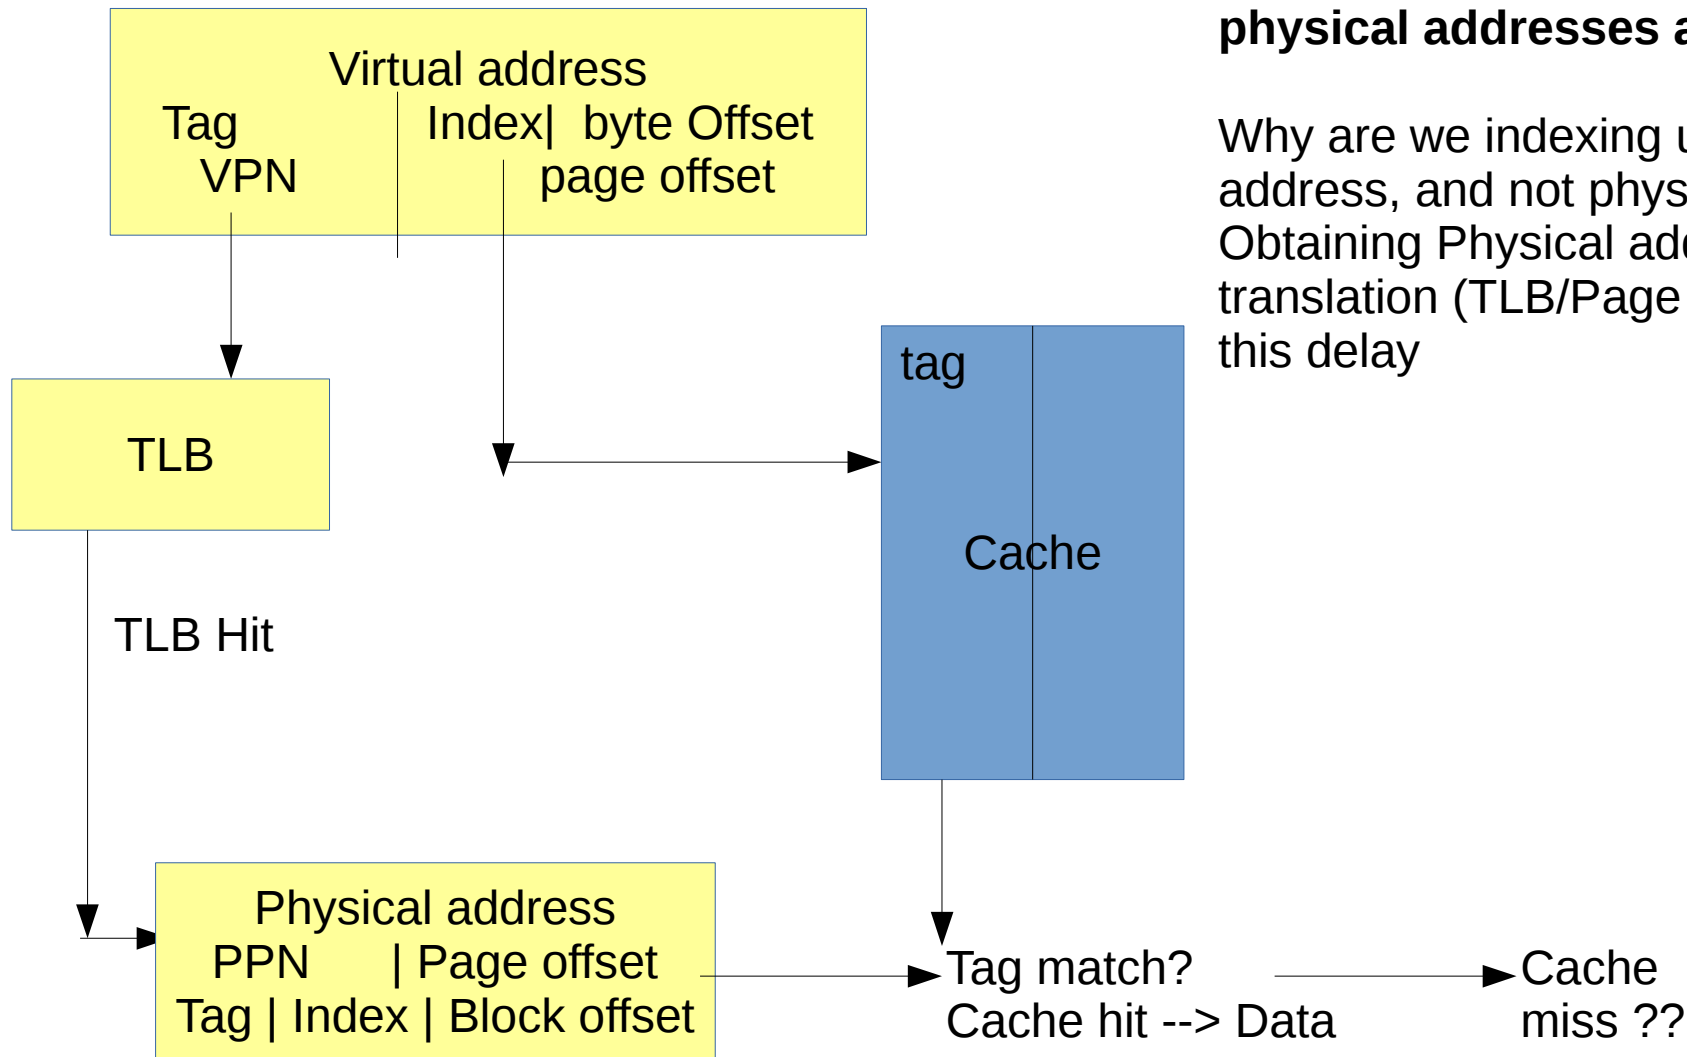


No aliasing if the index bits come from the page offset
and tag comes from the physical address

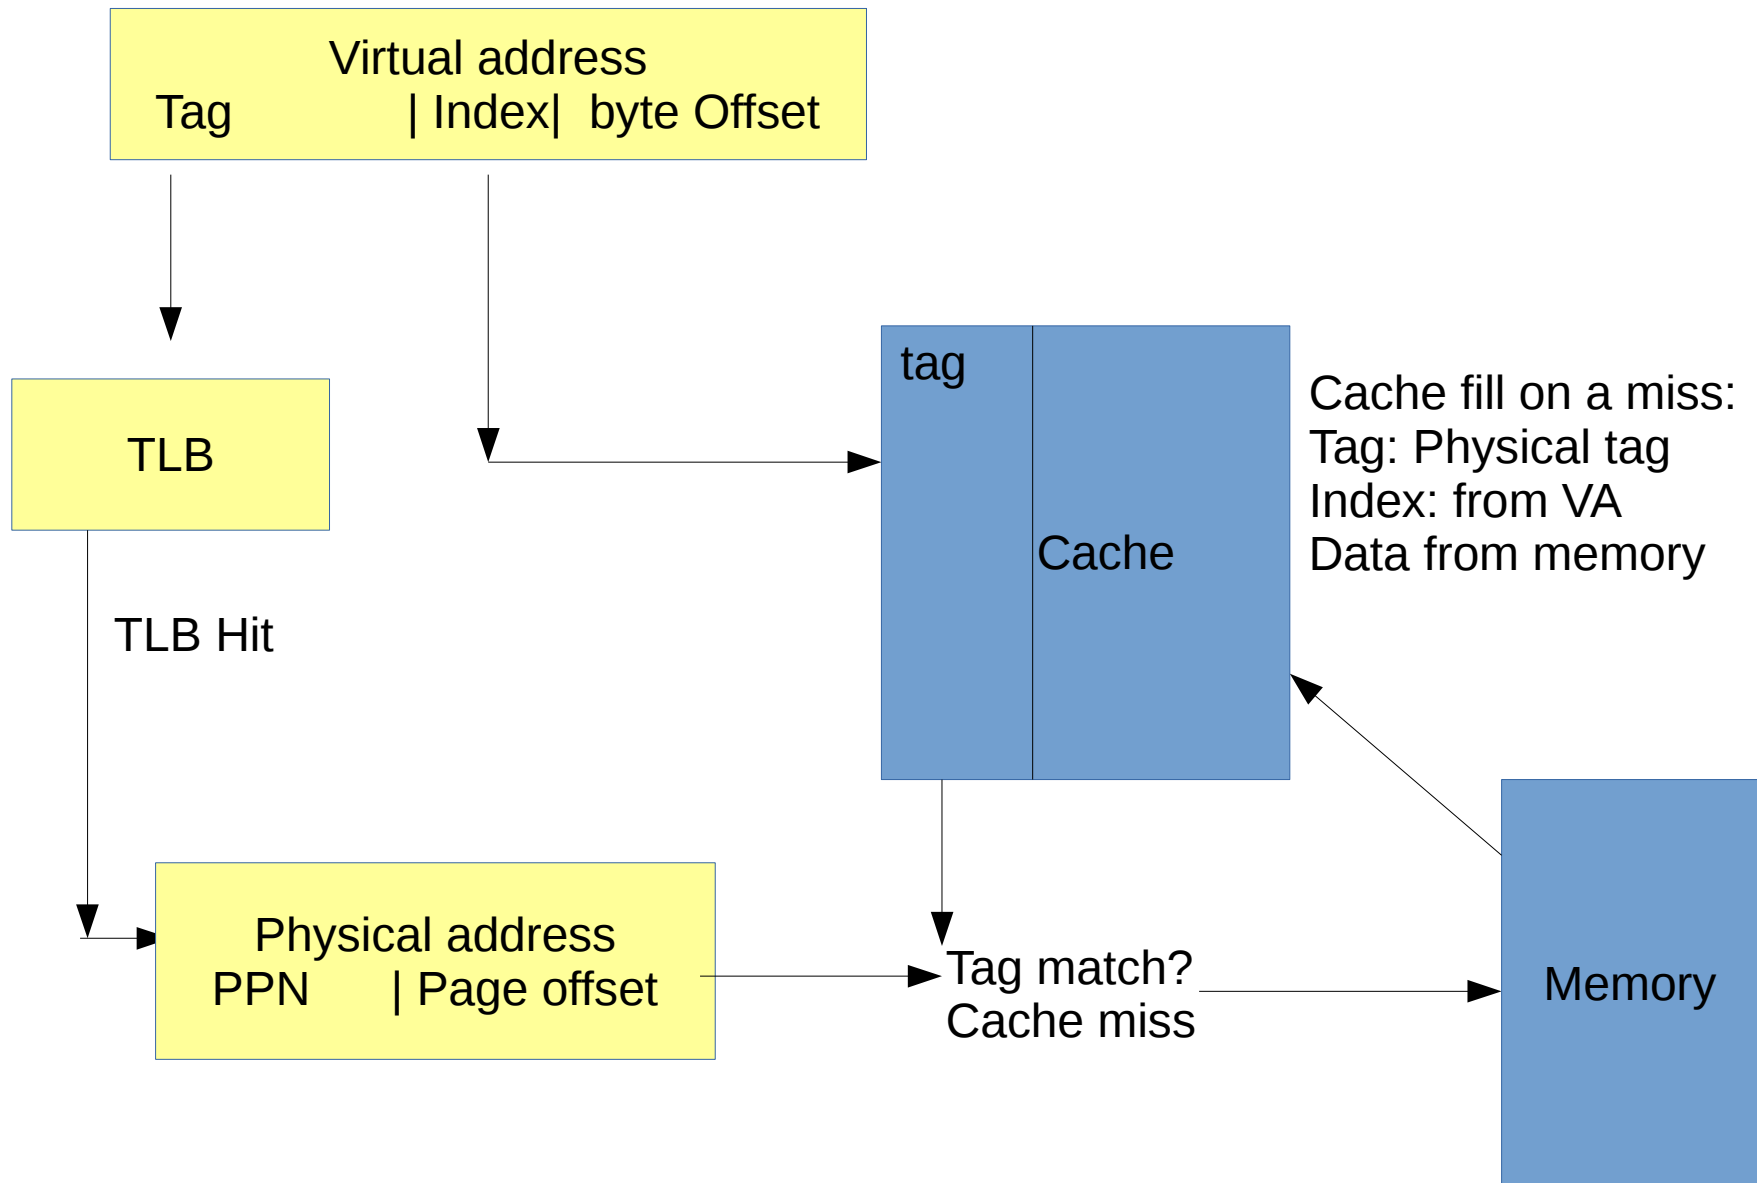
Virtually indexed Physically tagged Cache

The index bits in virtual and physical addresses are the same

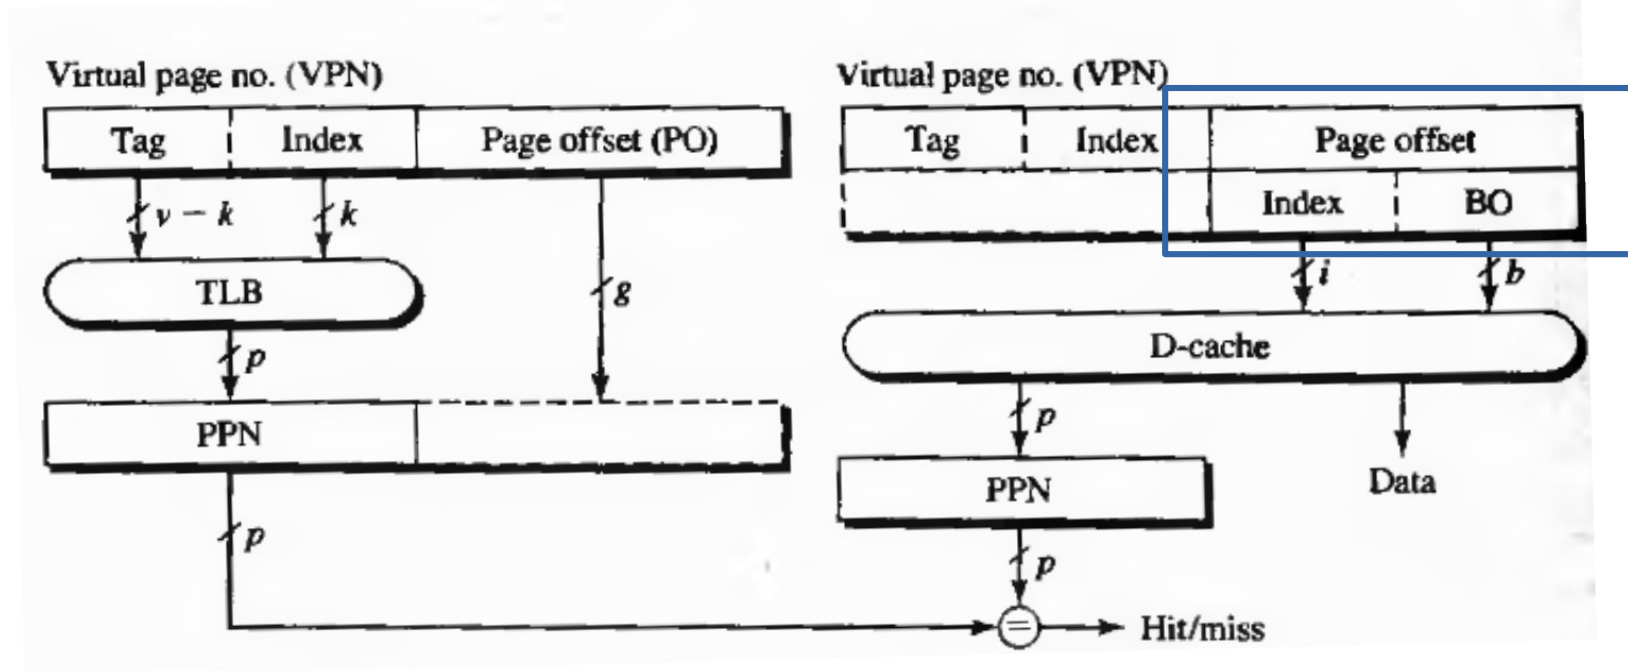
Why are we indexing using the virtual address, and not physical address? Obtaining Physical address will need translation (TLB/Page table). Avoid this delay



Virtually indexed Physically tagged Cache



VIPT Cache



- Index the cache using the virtual address. Index bits come from the Page offset- **This will limit the cache size**
- Tag will contain the physical address's tag obtained after address translation
- Cache hit check: Compare tags in the cache with the tag obtained from TLB
- Advantage: TLB and Cache look-ups are parallel.
 - No need of cache flush on context switch – ??

VIPT Cache


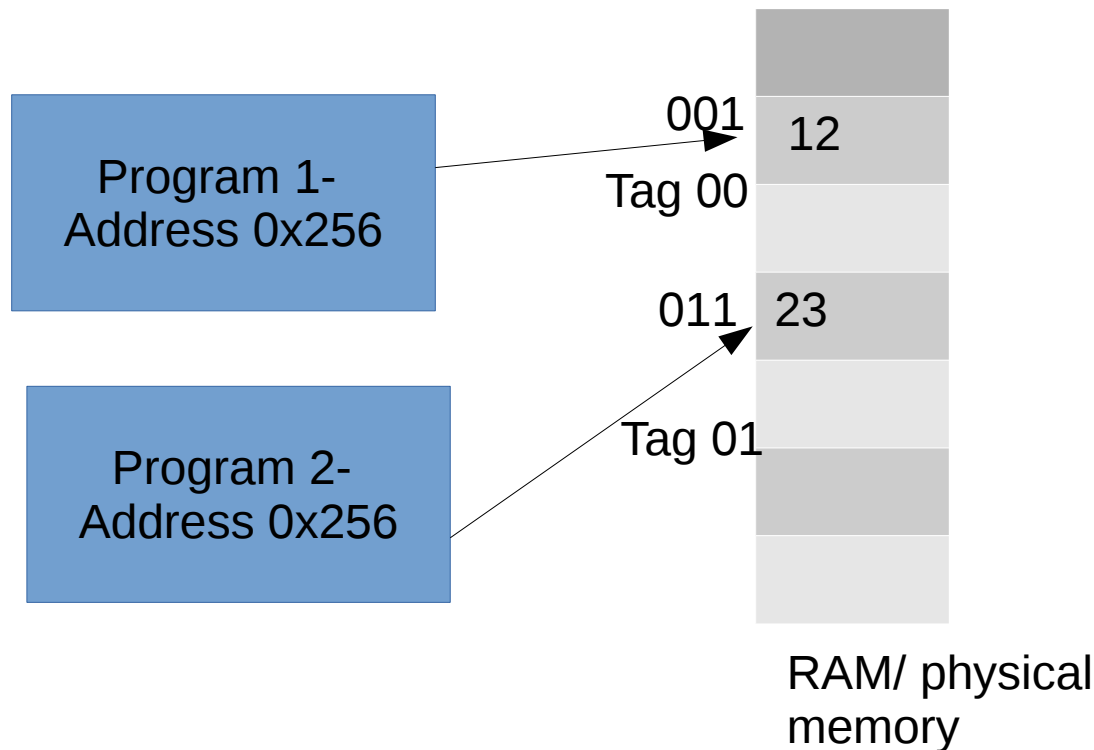
- When a process/context switch happens, the entire VIPT cache and TLB will have to be flushed so that correct data is loaded

RD -- 1 : Cache miss: TLB: 256 --> 001: Tag: 00: Read 12

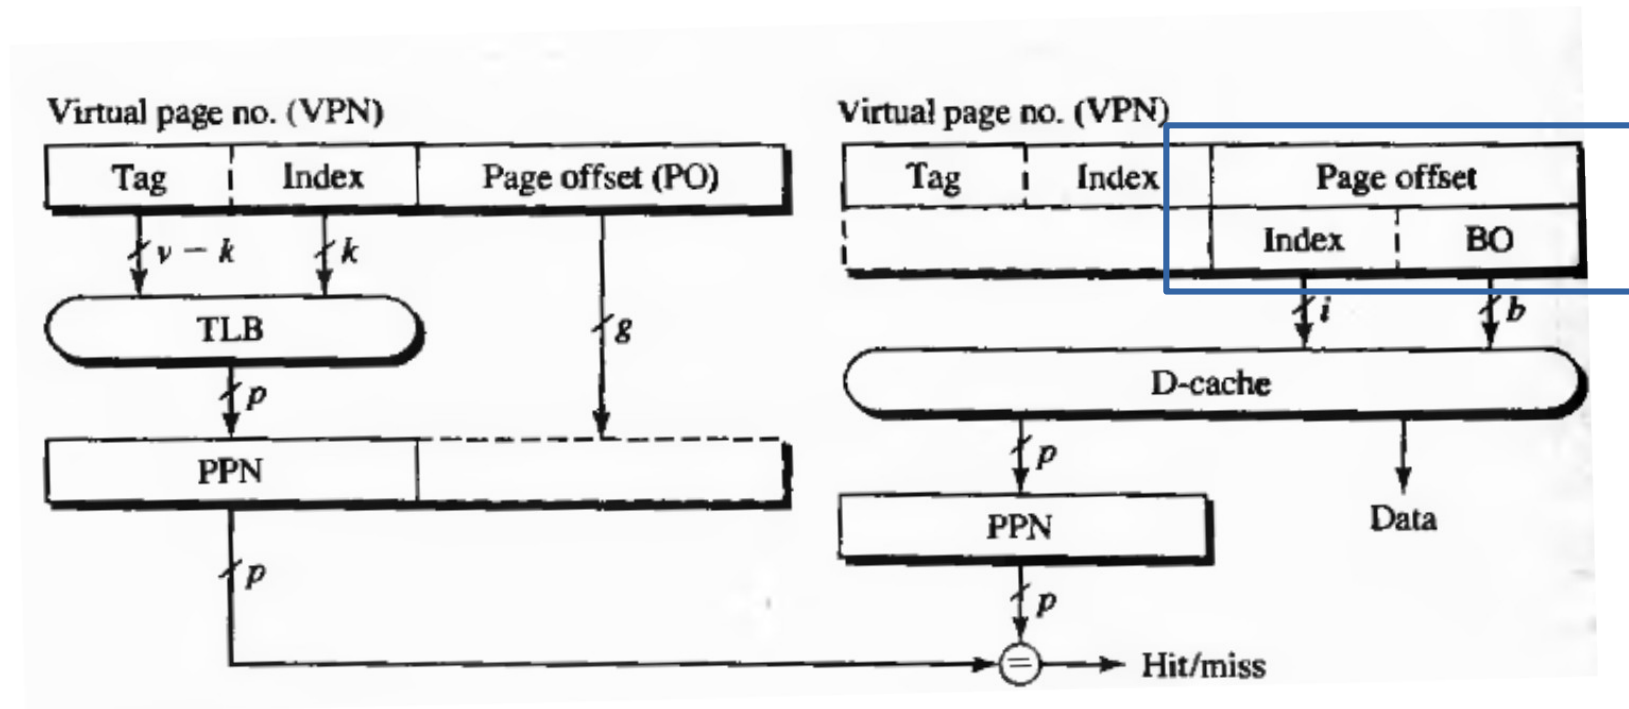
RD -- 2: Cache check: Tag mismatch- Cache miss. TLB 256

--> 011. **Reads 23**

	index	Byte offset
0010 0101	011	0



VIPT Cache size



- 4kB page: Needs 12 bits of page offset
- A 16Byte block size in the cache needs 4 bits of byte offset
 - If index has to come from the page offset to avoid aliasing, index bits = (Page offset – Byte offset) --> Max: 8 bits index
 - 2^8 lines or sets in the cache = 256 sets – maximum size

Eg


- VIPT cache: 4 way Set associative, 16 Byte block size, 4kB page. What is the maximum cache size to avoid aliasing

Eg

- VIPT: 4 way Set associative, 16 Byte block size, 4kB page. What is the maximum cache size to avoid aliasing
 - 4kB page: 12 bits of page offset
 - 16 B block size: 4 bits byte offset
 - Index: 8 bits
 - 2^8 sets * 2^4 bytes per block * 2^2 ways
 - 2^8 sets * 2^2 blocks per set * 2^4 bytes per block
 - 2^{14} bytes: Max size: 16kB

..cont

- VIPT: 4 way Set associative, 16 Byte block size, 4kB page.
What is the maximum cache size to avoid aliasing
 - 4kB page: 12 bits of page offset
 - 16 B block size: 4 bits byte offset
 - Index: 8 bits
 - 2^8 sets * 2^4 bytes per block * 4 way
 - $2^{(8+4)} * 4$ ways
 - $2^{(\text{index bits} + \text{byte offset bits})} * 4$ ways
 - $2^{\text{(page offset)}} * \text{No of ways}$
 - Page size * No of ways = Max cache size
- Eg: Pentium 4: 4 way SA cache, 4kB page--> 16kB L1 size
- Haswell: 8 way SA, 4kB page --> 32kB L1



Reading task

- Computer architecture- quantitative approach:
H&P (4th edition): 5.6 Putting It All Together:
AMD Opteron Memory Hierarchy
- Solve the homework practice problems