




Replacement policies

Prefetching



$$LMR = \frac{10}{100} \times 4-M$$

→ 10% LMR.



Replacement algorithms

No choice in a direct mapped cache

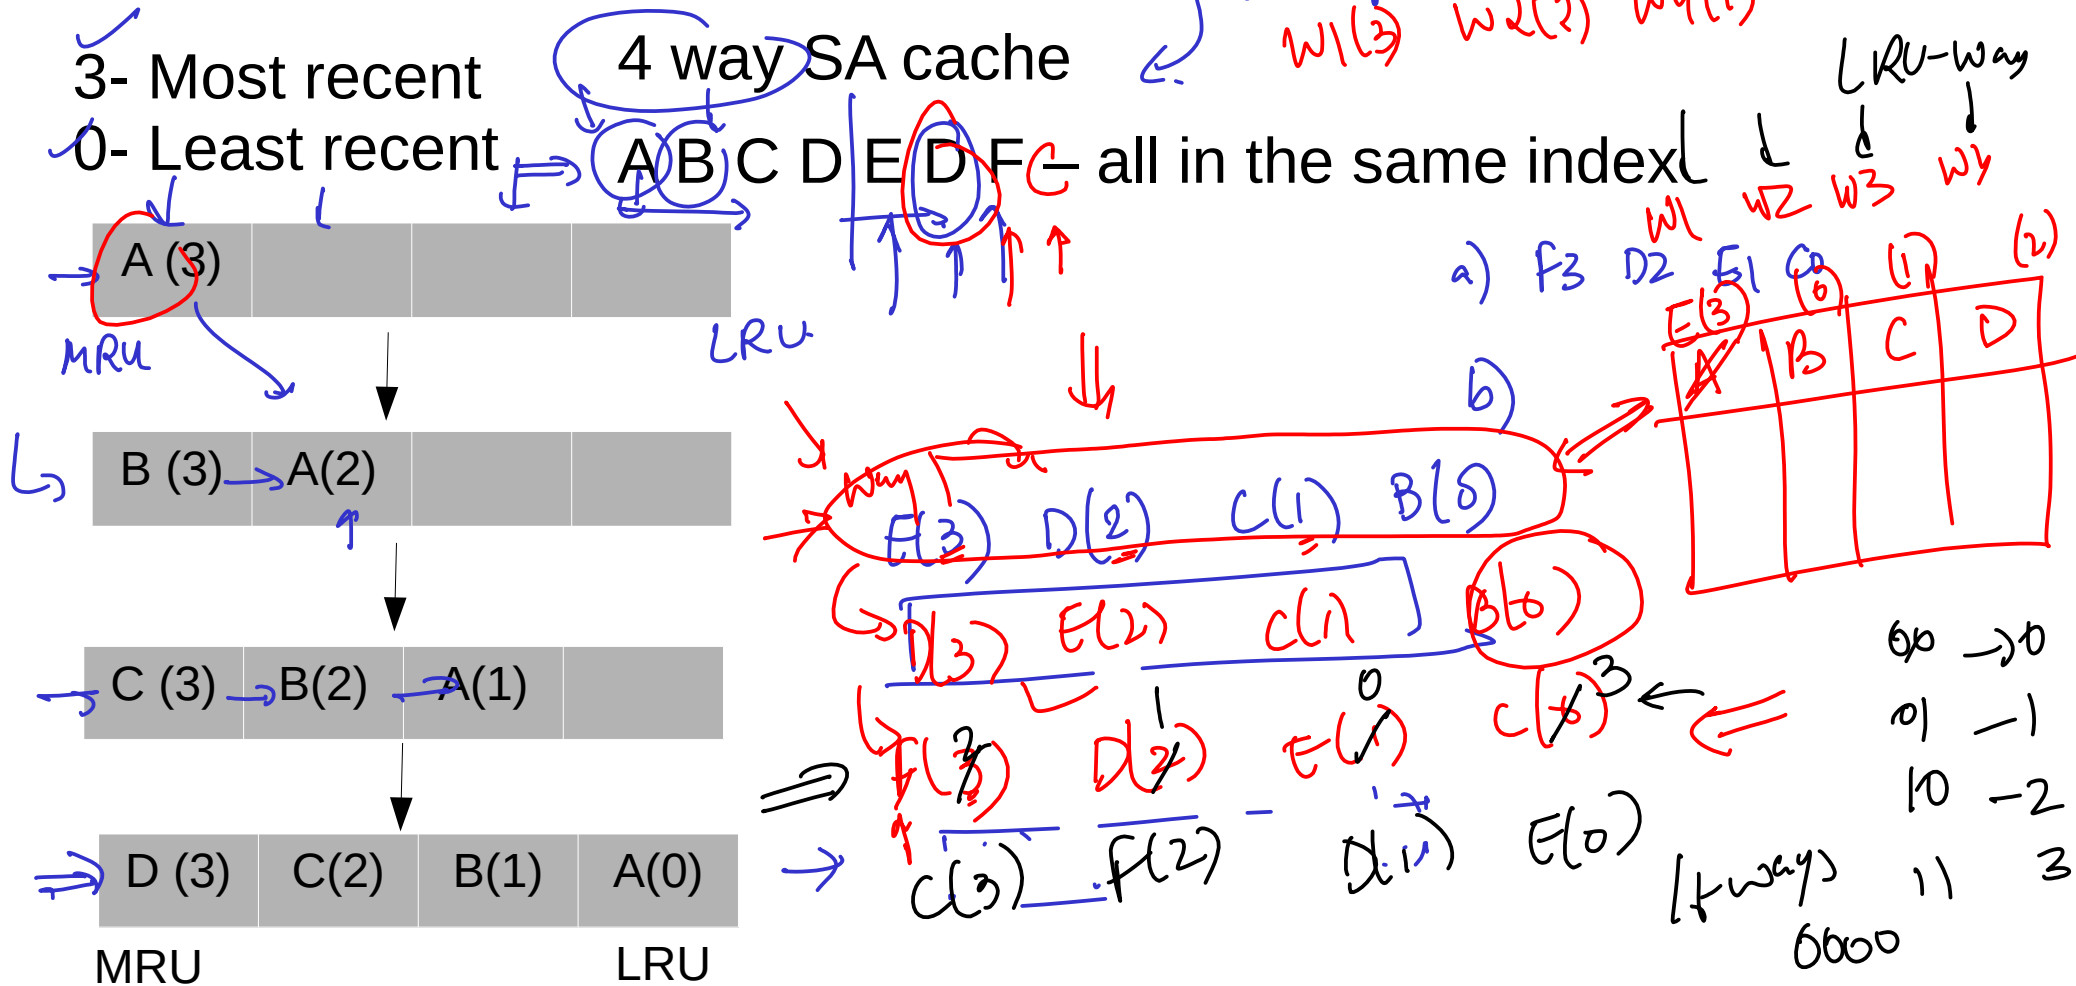
In an associative cache, which way should be evicted when the set becomes full?

- Belady's Optimal: L. A. Belady. 1966. A Study of Replacement Algorithms for a Virtual-storage Computer. IBM Syst. J. 5, 2 (June 1966), 78–101.
- FIFO (first-in-first-out)
- LRU (least recently used), pseudo-LRU
- LFU (least frequently used)

Optimal Replacement Policy/ Belady's anomaly?

- [Belady, IBM Systems Journal, 1966]
- Evict block with longest reuse distance
 - i.e. next reference to block is farthest in future
 - Requires knowledge of the future!
- Can't build it, but can model it with trace
 - Process trace in reverse
- (X,A,B,C,D,X): LRU 4-way SA. How far in the future is X going to be accessed?

LRU *shift-reg.*



Maintain "age bits" or counters for cache-lines and replace based on the smallest number

LRU



3- Most recent
0- Least recent

A B C D E D F – all in the same index



MRU

LRU




MRU

LRU



Replace LRU
Decrement
other counters



Update counters
Decrement if
more than the
current one



Update
counters


Maintain counters for each way.

Also, check all other counters if they are more than the current one

LRU with inclusive caches

To demonstrate back invalidation

A B A C A D A | E A - accesses



Fully associative 2 blocks L1

1- MRU
0 - LRU

0

1

LRU counter



Fully associative 4 blocks L1

2

3

0

1

LRU counter

3- MRU
0 - LRU

LFU

A B C D E C D B

Store the value of how many times it was accessed-frequency instead of when it was used.

Can combine with LRU --> LFRU policy

Pseudo LRU

- LRU disadvantages -->
 - Counters for each block. 4 way: 2 bit counter + Keep track of other counters.
 - Update counters on each access
- Pseudo LRU: Single bit
 - 1 – when accessed
 - 0 – is the least recently accessed
 - Replace the block which has a 0 bit (randomly if more than one block which has '0')

PLRU

64 ways + 6-bit

Way 1



MRU

Set a block bit to '1' if accessed.



↓ MRU




Random

On a hit, no need to check bits of other blocks.

Just set the block to '1'









When the last one is set to '1', zero out the others



Replace a block which has '0' bit, and set the bit to 1. This is done randomly, so not true LRU

4-way superscalar



Cache Prefetching

- Steven P. Vanderwiel, "Data Prefetch Mechanisms" ACM 2000
- H&P – Chapter "Memory Hierarchy design"


Motivation

- Cache: “on demand” memory fetch policy
 - processor requests a word --> cache miss --> fetch
 - Store only previously accessed data
 - Larger block size --> fetches consecutive words, but on-demand
 - Disadvantage: Evict all blocks (evict useful data)
- Anticipates cache misses and issues a fetch to the memory system--> placed in a prefetch buffer
- Proceeds in parallel to processor computation
- Significantly improve overall program execution

Introduction

- Speculate: Fetch instructions or data from memory to cache before they are needed
- Data or instruction prefetching
 - Easier to guess instructions
- How?
 - Software based: Programmer/compiler inserts "prefetch" instructions in the program
 - Hardware based: In processor: watches the stream of instructions or data

Introduction




- What to fetch:
 - Involves predicting which address will be needed in the future
 - Misprediction: Is ok. Prefetched data will not be used
- Predict based on past patterns
 - Prefetching algorithm
- PreFetch data up in the memory hierarchy before they are actually needed by the processor

Metrics

Prefetches

- ✓ Accuracy = No of useful prefetches / Total prefetches
- ✓ Timeliness = When to prefetch?
 - Too early
 - Might replace useful data
 - Might get replaced by the time it is used (Stall)
 - Too late
 - Processor has to wait (Stall)

Hide \rightarrow Mem access latency



8

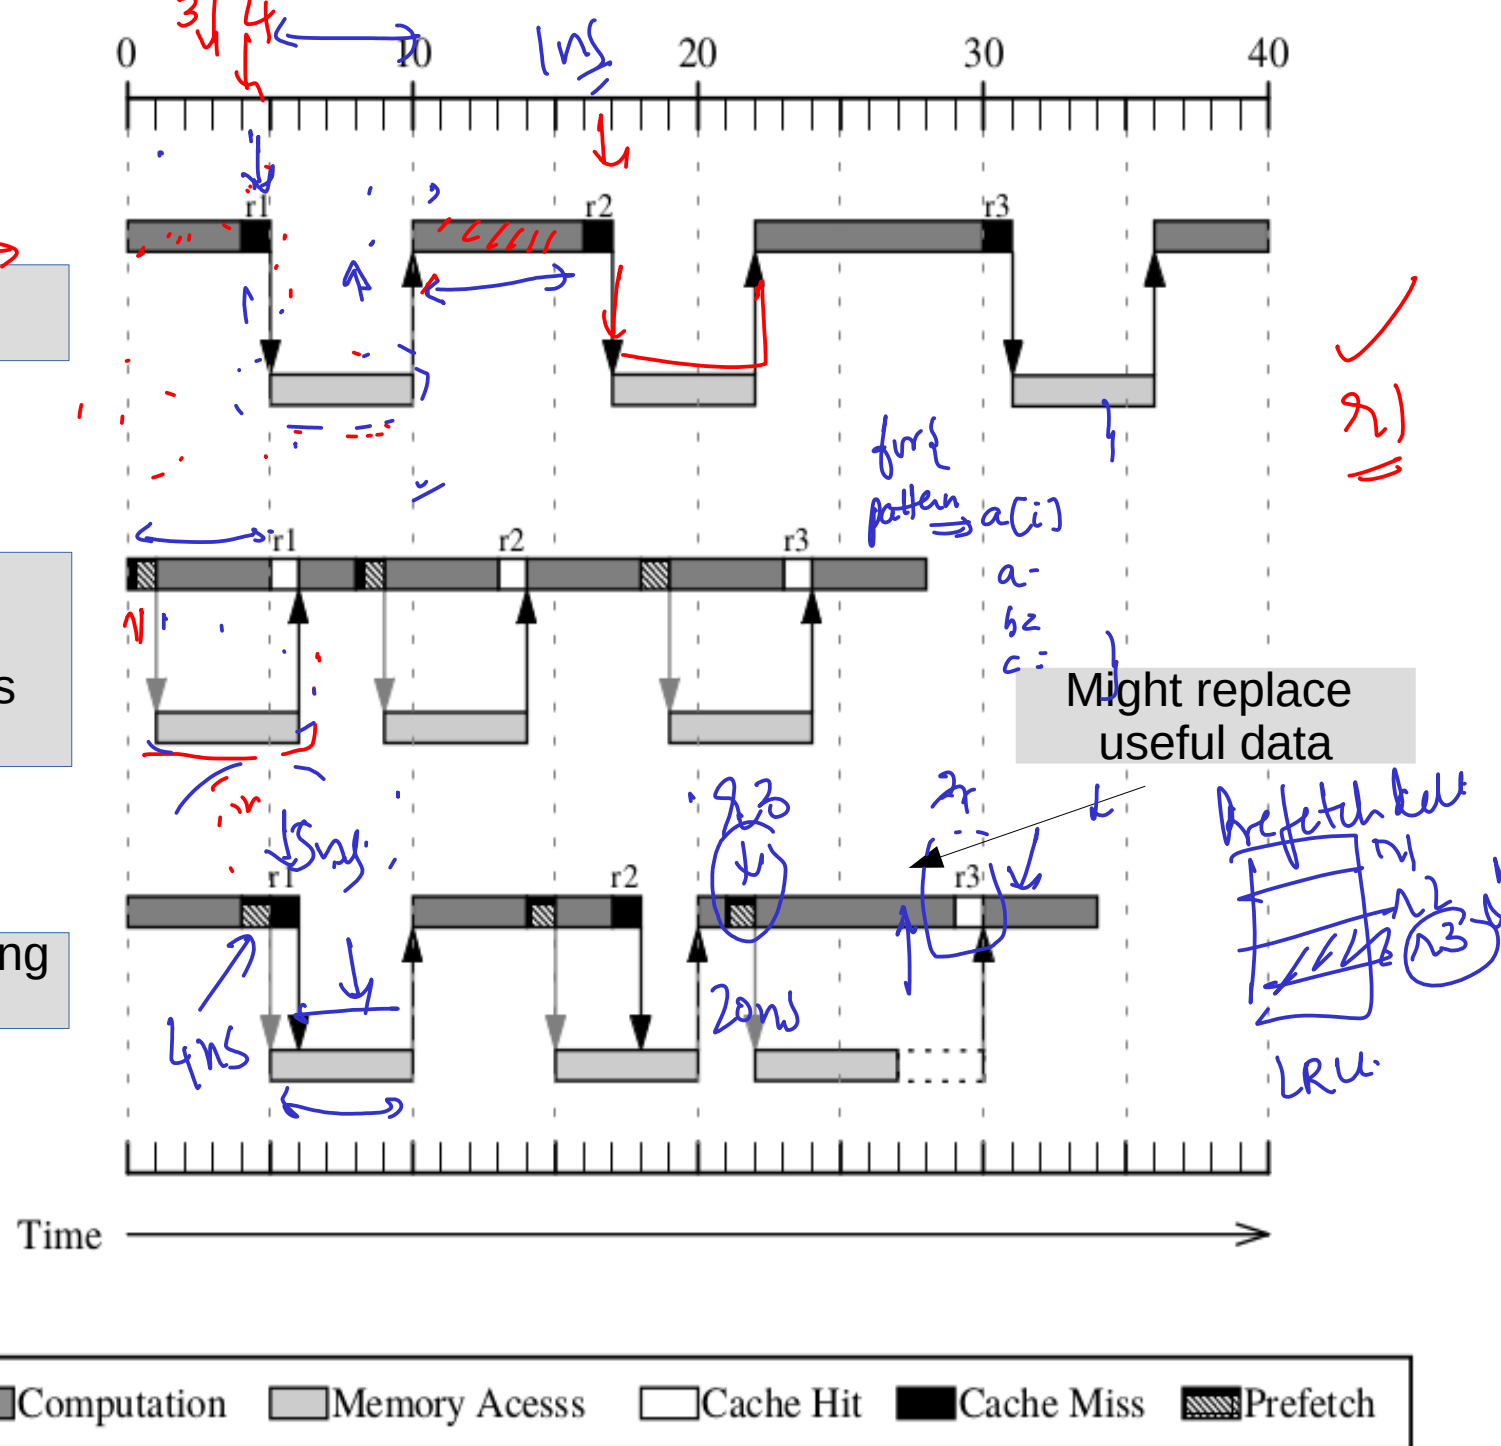
Timeliness

no prefetching

$P \rightarrow \text{Computer}$ Prefd

perfect prefetching-
No misses
Fetch "memory cycles
Latency- earlier"

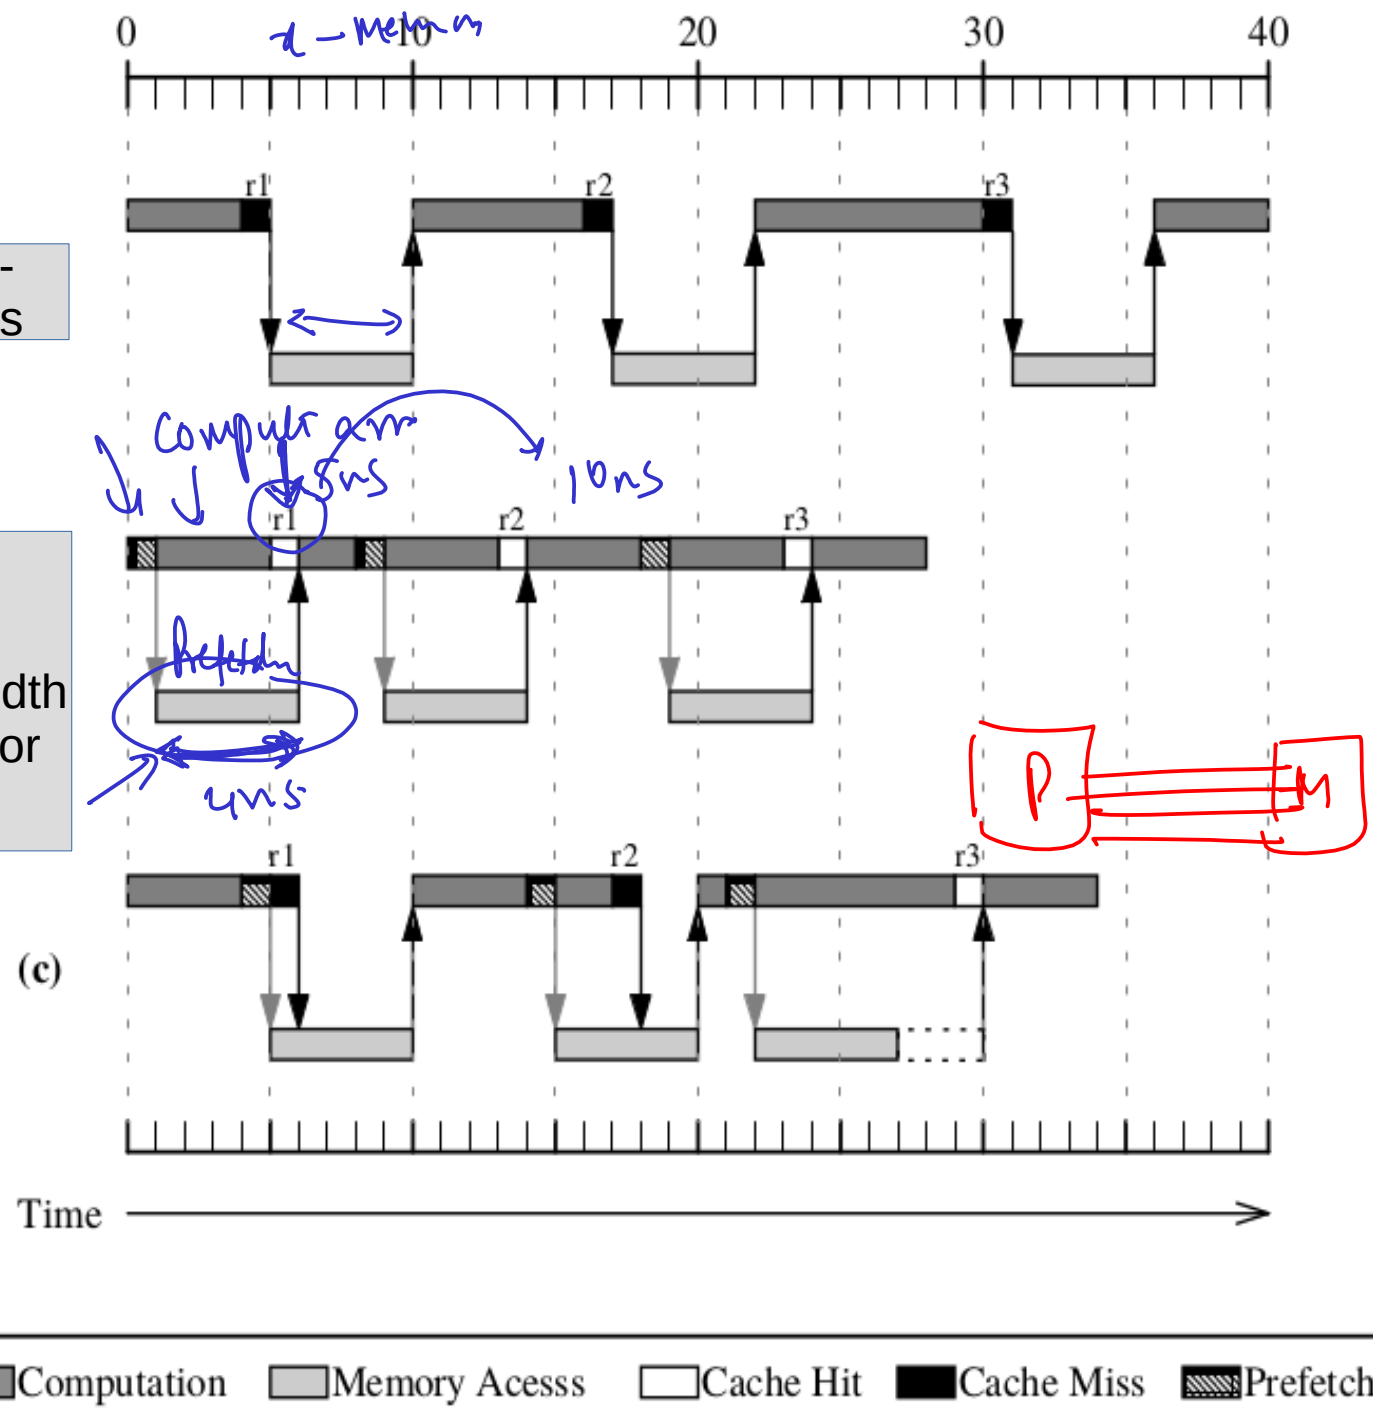
late
degraded prefetching
Prefetch late



Memory accesses


31 units of time--
3 mem accesses

19 units of time--
3 mem accesses
Higher Memory bandwidth
Work done by processor
is more



Software prefetching

Assume a Cache with 4-word blocks




$\Rightarrow a[i]$
 $\Rightarrow a[i+4]$

```
for (i = 0; i < N; i++)  
    ip = ip + a[i]*b[i];
```

$a[i+1] * b[i+1]$

no prefetching
Will miss every 4th iteration




```
for (i = 0; i < N; i++) {  
    fetch( &a[i+1]);  
    fetch( &b[i+1]);  
    ip = ip + a[i]*b[i];  
}
```

one =

$a[i]$ $a[i+1]$

$a[i]$ prefetch buffer

simple prefetching
Prefetch every iteration --
-> fetches 4 blocks



loop unrolling
Parallel

```
for (i = 0; i < N; i+=4) {  
    fetch( &a[i+4]);  
    fetch( &b[i+4]);  
    ip = ip + a[i]*b[i];  
    ip = ip + a[i+1]*b[i+1];  
    ip = ip + a[i+2]*b[i+2];  
    ip = ip + a[i+3]*b[i+3];  
}
```

$a[i] = i+4$ $fetch(ba[i+1])$
 $(i+6)$
 $(i+8)$
 $(i+10)$

prefetching with loop unrolling
Unroll y times, where y = number of words in each block

loop \rightarrow LW $a[i]$
 \rightarrow LW $b[i]$
 \rightarrow MAC SW
(LW $a[i+1]$)

Cache misses will occur during the first iteration of the loop

Unnecessary fetches during the last iteration
This code does not consider memory latency

Software prefetching for loops

{ LW 4(5)
 MAE
 SW-
 LW i+2
 LW i+3
 Add i+4
 Beg loop

⇒ for (int i=0; i<1024; i++) {
 array1[i] = 2 * array1[i];
 }

Works well with loops and regular access patterns

- Prefetch instructions added either by the programmer or compiler
- Prefetch for the next iteration
- Prefetch directives improve performance


i=7
 for (int i=0; i<1024; i++) {
 prefetch (array1 [i + k]);
 array1[i] = 2 * array1[i];
 }
 k=7

Prefetch k elements ahead

K= Prefetch distance

⇒ What is k?


⇒ If k=7, Compulsory misses: 0 to 6 will be misses
 0 to 6



Next-line $i+1$
 0+1
 0+k
 Strided prefetching
 offset prefetch.

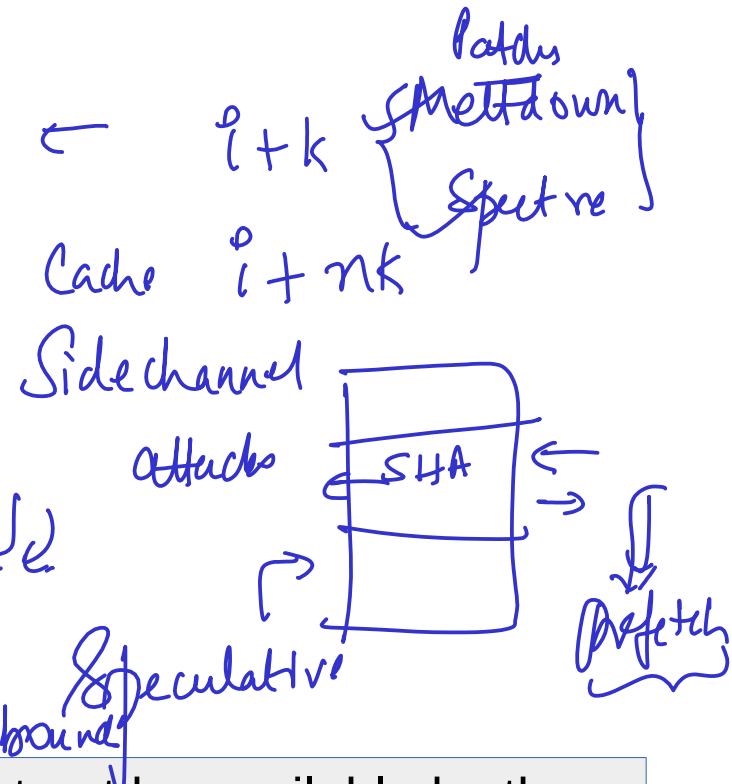
Software prefetching for loops

```
for (int i=0; i<1024; i++) {
    array1[i] = 2 * array1[i];
}
```



Works well with
loops and regular
access patterns

```
for (int i=0; i<1024; i++) {
    prefetch (array1 [i + k]);
    array1[i] = 2 * array1[i];
}
```



- Currently $i=0 \rightarrow k=1$, Prefetched data might not be available by the next iteration. (Too late)

- $i=0, k=20$, Prefetched data might come in too early and replace useful data (Too early)

Not easy to predict k at compile time

- Can combine loop unrolling with software prefetching

Software prefetching for loops

↓
Prefetch how many iterations ahead?

- L/S iterations ahead

L = average memory latency in processor cycles *hide.*

S = shortest execution time of 1 iteration

L = 5 cycles of memory latency

S = 1 cycle

Prefetch how many iterations ahead?

→ 5 iterations ahead

$\frac{5 \text{ cycle.}}{1}$

Limitations

- + • Works for Regular and predictable array accesses
- • Code expansion, more instructions
- • Statically done by compiler, does not exploit runtime information
- • Programmer has to do this manually

exception
page fault }

Hardware prefetching

- Sequential prefetching
- Strided (distance) prefetching

→ Best offset prefetcher (BOP)


$a+k$ ←

$a+k$

$a+s$

→ $lw\ a[0] \leftarrow$

$a[i+1]$



Sequential prefetching

- Prefetch on miss
 - Prefetch next block $y+1$ if y is a miss
 - If $y+1$ is already in cache, do not prefetch

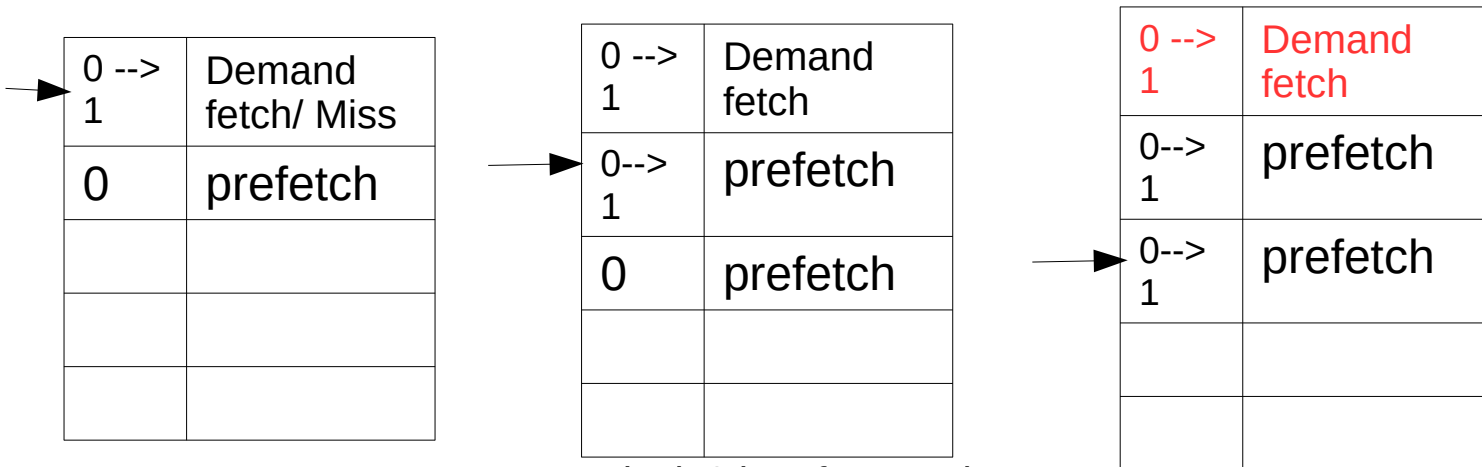


Demand fetch means miss
No prefetch generated here since
it was a hit

Overall 2 misses

Tagged prefetching

- A tag bit with each cache line
- 0:
 - Initially
 - Reset to 0 on replacement
 - Prefetch
- 1:
 - When the line is referenced
 - Brought into cache on demand (after a miss)
- **Prefetch is initiated when tag changes from 0 --> 1, that is, if the line is referenced (indicated by arrow)**
 - When a data is prefetched and accessed, more confidence in prediction



Block 1--> Initially 0, changed to 1
Initiate prefetch for next block

Block 2 is referenced and is a hit. 0--> 1
Initiate prefetch for next block

Overall 1 miss


Strided prefetching

- Strided (distance) prefetching: Consecutive blocks that are fetched are “y” addresses apart
 - Block b , $b + y$, $b + 2y$?
- Check 2 consecutive loads, and the distance between their memory accesses. Set this as the stride
 - Predict their access pattern


Where to place the prefetched data - is a design choice

H/W based stream buffers

- Placing in cache is easiest --> Pollutes the cache. Replaces useful data
- Buffers: k subsequent addresses are fetched into a buffer of depth k:
 - Eg- k=4--> Miss on A fetches A+1, till A+4 – **Sequential/ Next-line prefetching**
- Multiple streams for multiple misses



Multiple Stream buffers for data


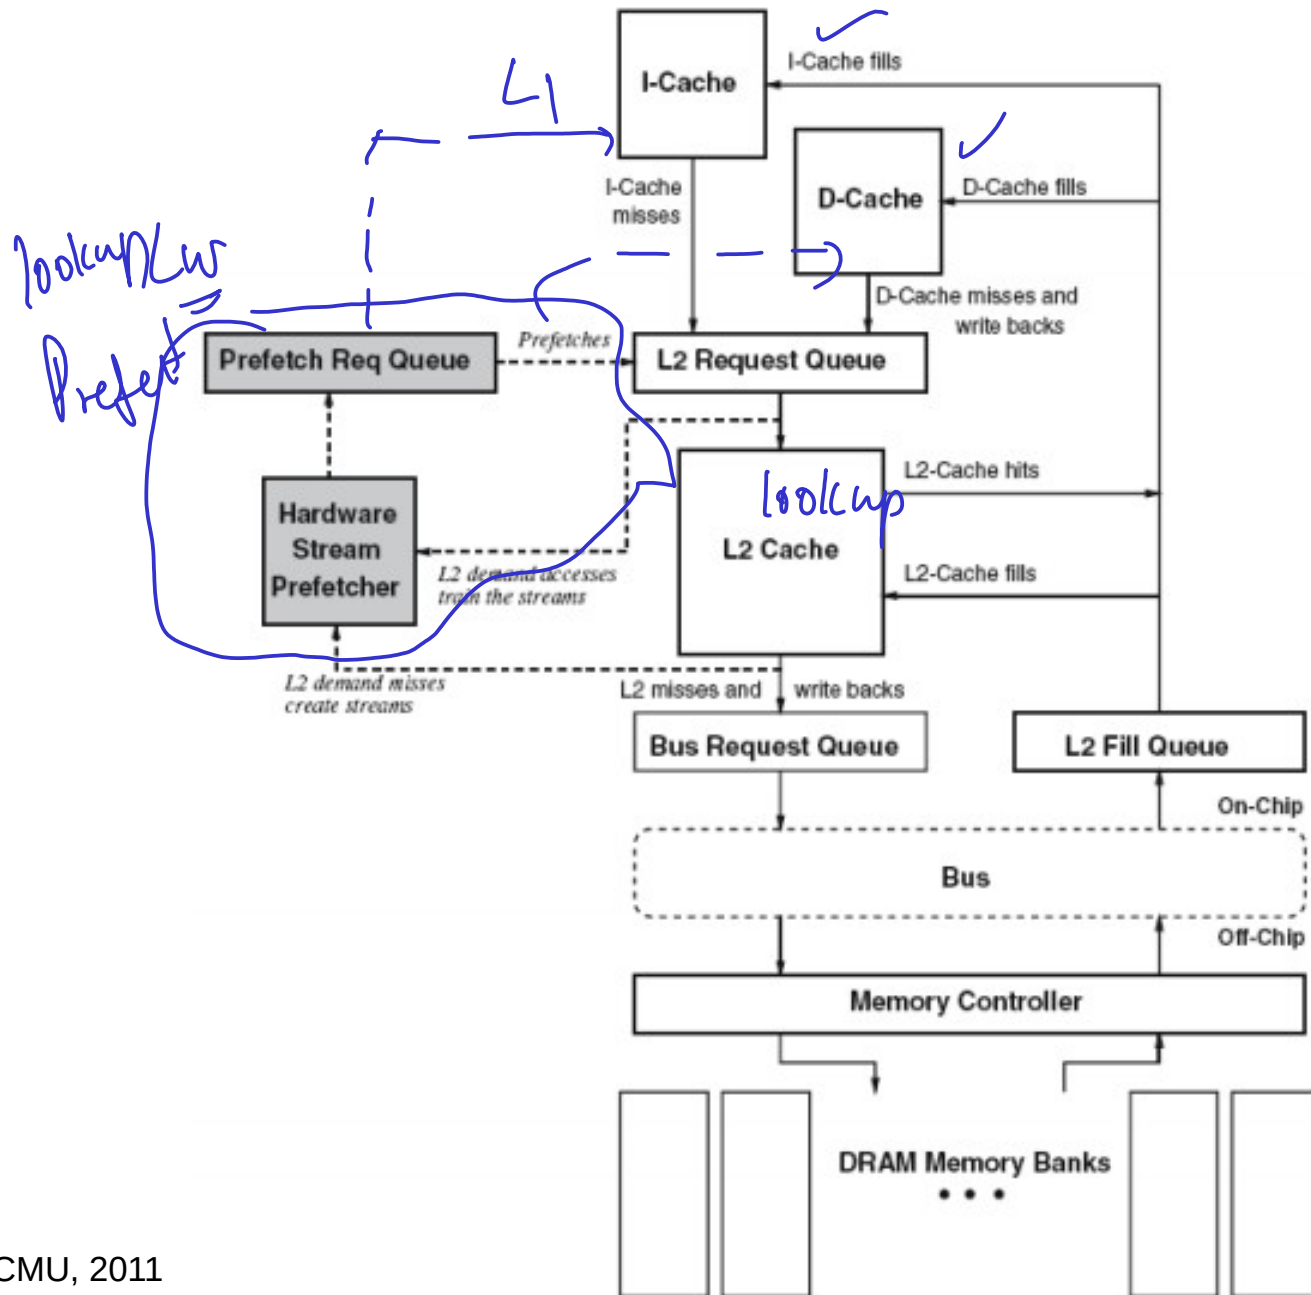


- Each stream buffer holds one stream of sequentially prefetched cache lines
- On a cache/load miss, check the head of all stream buffers for an address match
 - Hit, pop the entry from FIFO, update the cache with data. Fetch the next in line data
 - Miss-- allocate a new stream buffer to the new miss address
 - If all streams are full, replace a stream buffer following LRU policy
- L2 and stream buffer can be checked parallelly


Buffers

- Where to place the buffer: btn L1 and L2, memory to L2?
- When to access the buffer (parallel vs. serial with cache)
- When to move the data from the prefetch buffer to cache
- Size the prefetch buffer

Overall system



Performance



Speedup due to hardware prefetching on Intel Pentium 4 with hardware prefetching turned on

Acknowledgements

- TU Berlin, Software and Hardware Prefetching
- UCB- CS 152
<https://inst.eecs.berkeley.edu/~cs152/sp16/lectures/L06-Memory.pdf>